

BST Application

Range Query: 1D

e> - A>

邓俊辉

deng@tsinghua.edu.cn

你这个人太敏感了。这个社会什么都需要，唯独不需要敏感。

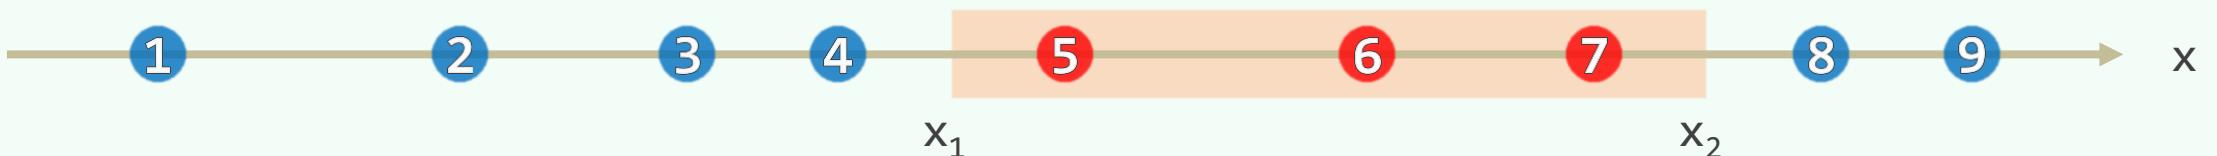
1D Range Query

- ❖ Let $P = \{ p_1, p_2, p_3, \dots, p_n \}$ be a set of n points on the x-axis
- ❖ For any given interval $I = (x_1, x_2]$
 - COUNTING: how many points of P lies in the interval?
 - REPORTING: enumerate all points in $I \cap P$ (if not empty)
- ❖ [Online] P is fixed while I is randomly and repeatedly given
- ❖ How to PREPROCESS P into a certain data structure s.t.
the queries can be answered efficiently?



Brute-Force

- ❖ For each point p of P , test if $p \in (x_1, x_2]$
- ❖ Thus each query can be answered in LINEAR time
- ❖ Can we do it faster? It seems we can't, for ...
- ❖ In the worst case,
the interval contains up to $\Theta(n)$ points, which need $\Theta(n)$ time to enumerate
- ❖ However, how if we
ignore the time for **enumerating** and count only the **searching** time?



Binary Search

- ❖ Sort all points into a sorted vector and add an extra sentinel $p[0] = -\infty$
- ❖ For any interval $I = (x_1, x_2]$
 - Find $t = \text{search}(x_2) = \max\{ i \mid p[i] \leq x_2 \} // O(\log n)$
 - Traverse the vector BACKWARD from $p[t]$ and report each point $// O(r)$ until escaping from I at point $p[s]$
 - return $r = t - s // \text{output size}$



Output-Sensitivity

- ❖ An **enumerating** query can be answered in $O(r + \log n)$ time
- ❖ $p[s]$ can also be found by binary search in $O(\log n)$ time
- ❖ Hence for COUNTING query, $O(\log n)$ time is enough //independent to r
- ❖ Can this simple strategy be extended to PLANAR range query?

TTBOMK, unfortunately, no!



BST Application

Range Query: 2D

θ> -A₂

昔者明王必尽知天下良士之名；既知其名，又知其数；
既知其数，又知其所在。

邓俊辉

deng@tsinghua.edu.cn

Planar Range Query

❖ Let $P = \{ p_1, p_2, p_3, \dots, p_n \}$ be a planar set

❖ Given $R = (x_1, x_2] \times (y_1, y_2]$

- COUNTING: $|R \cap P| = ?$
- REPORTING: $R \cap P = ?$

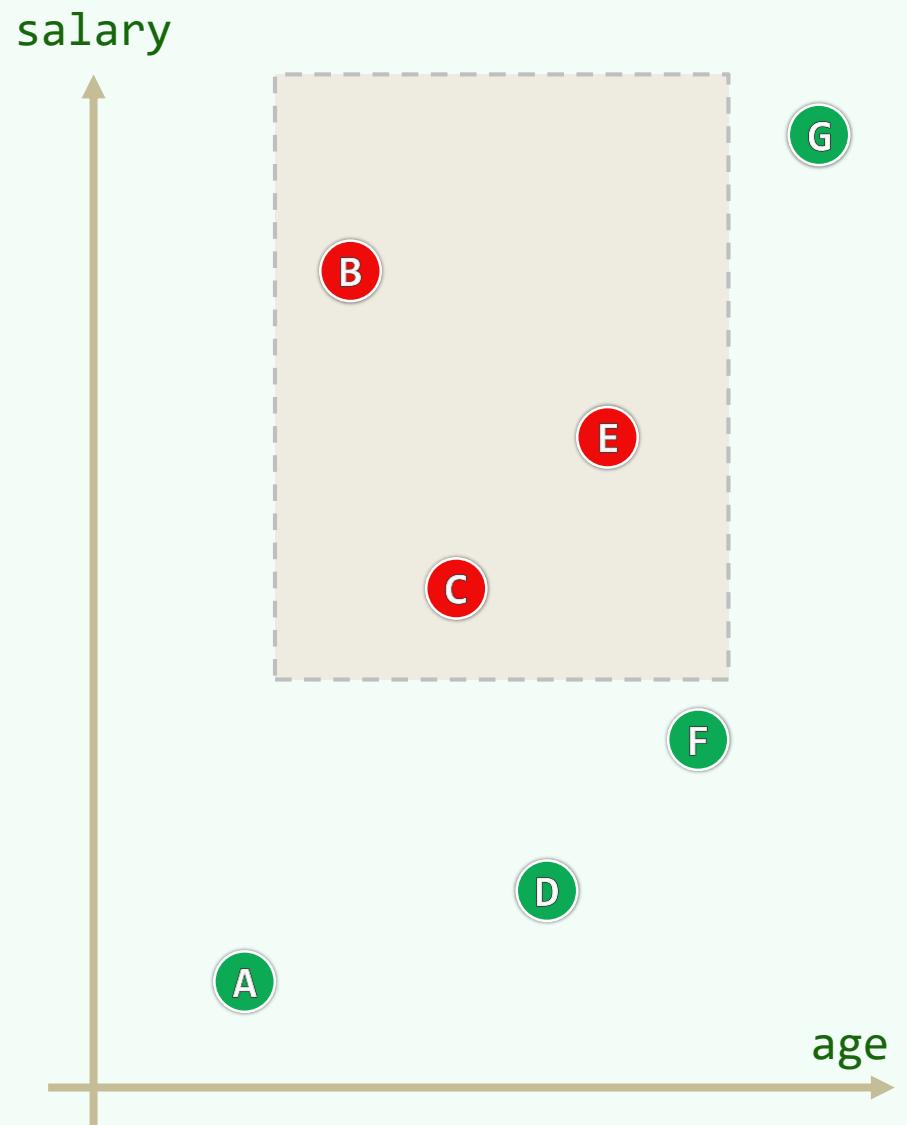
❖ Binary search

doesn't help this kind of query

❖ You might consider to

expand the counting method using

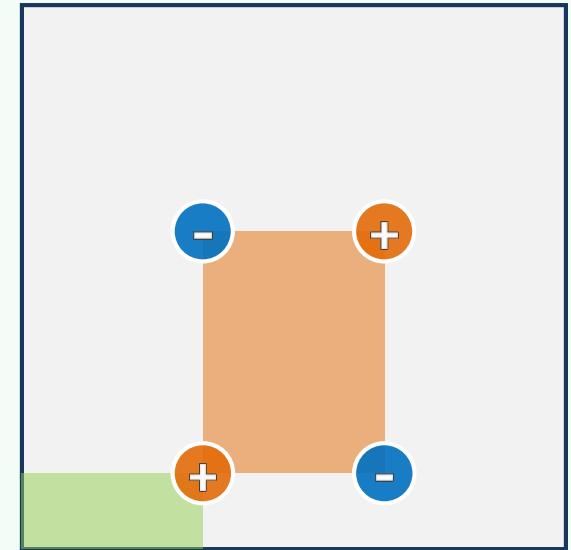
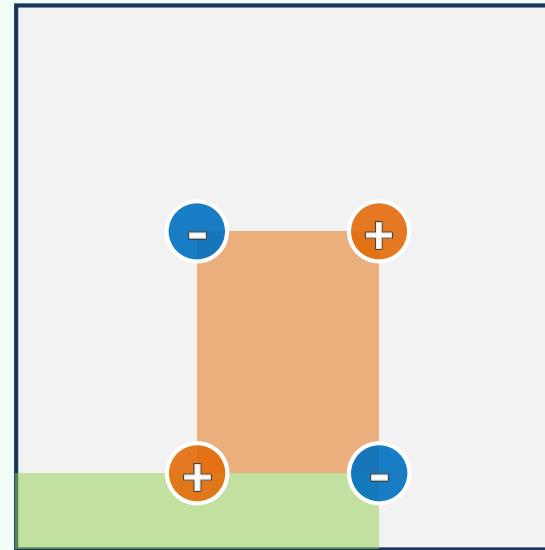
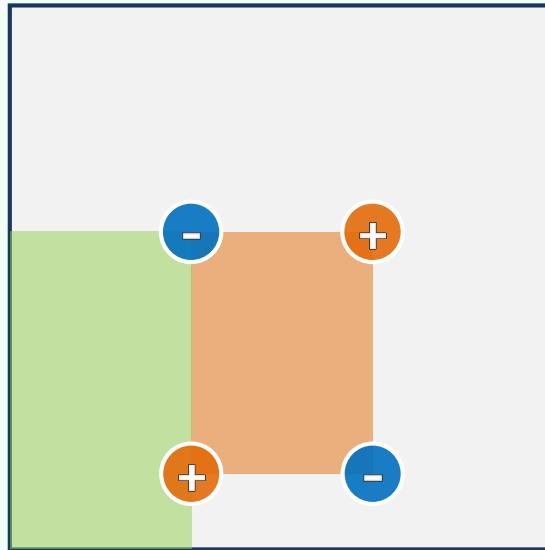
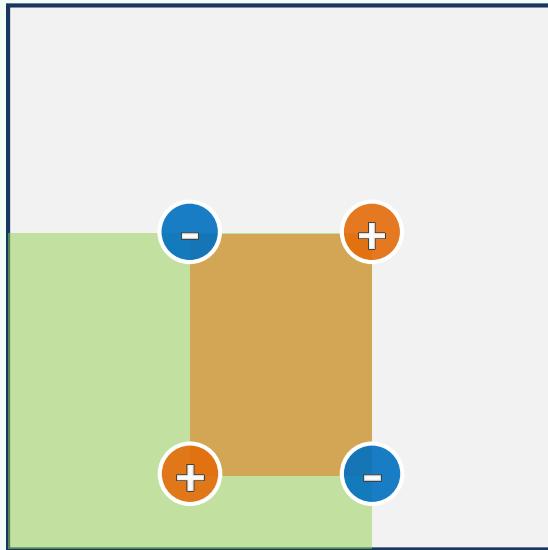
the Inclusion-Exclusion Principle



Preprocessing

❖ \forall point (x, y) , let $n(x, y) = |((0, x] \times (0, y]) \cap P|$

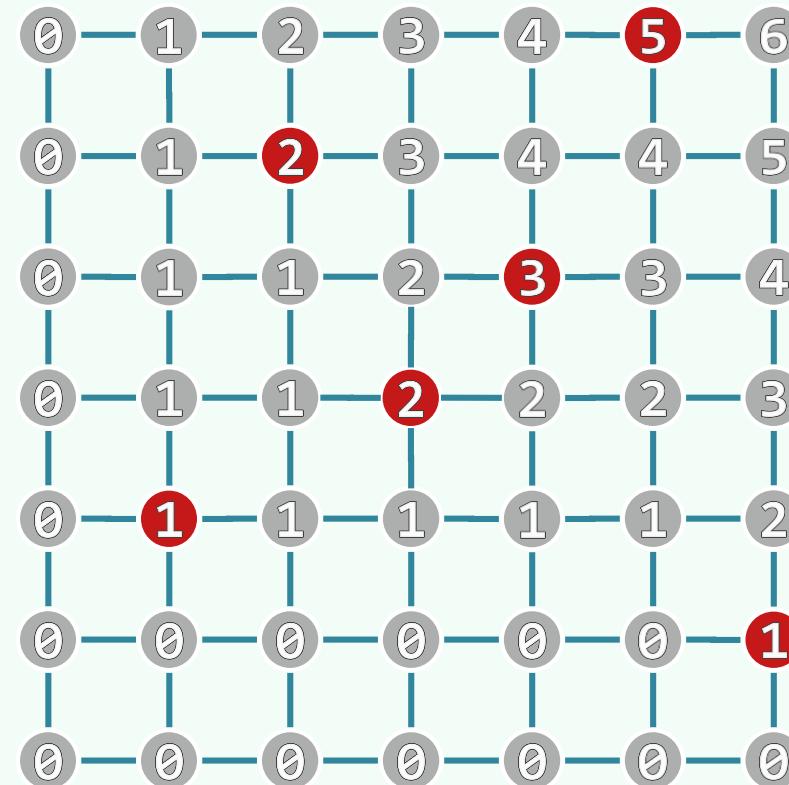
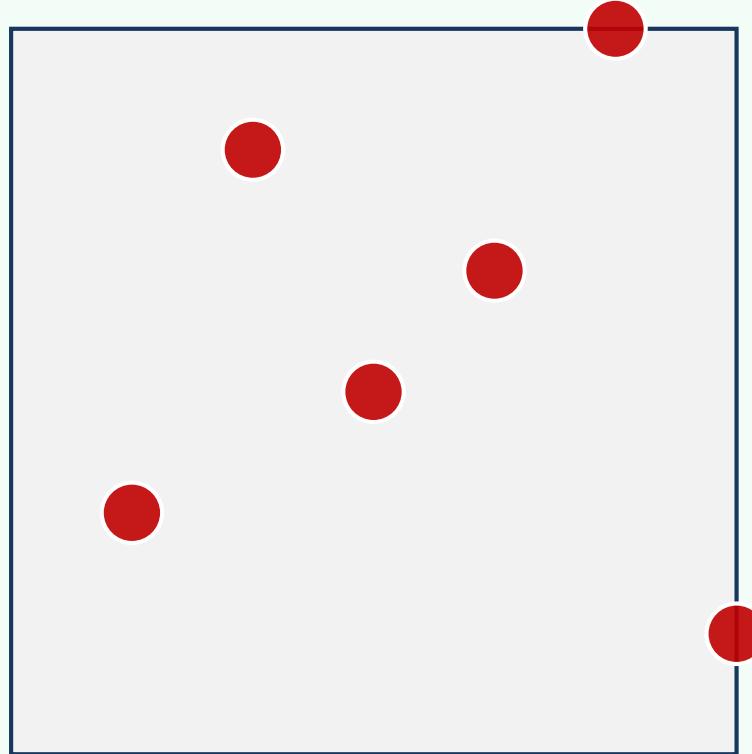
❖ This requires $\mathcal{O}(n^2)$ time/space



Domination

❖ A point (u, v) is called to be **DOMINATED** by point (x, y) if

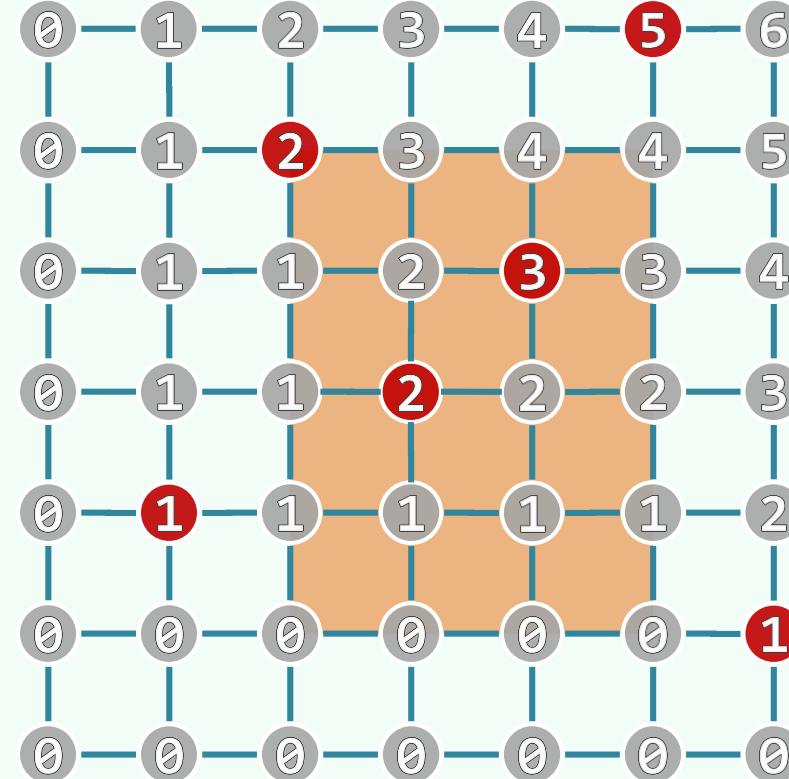
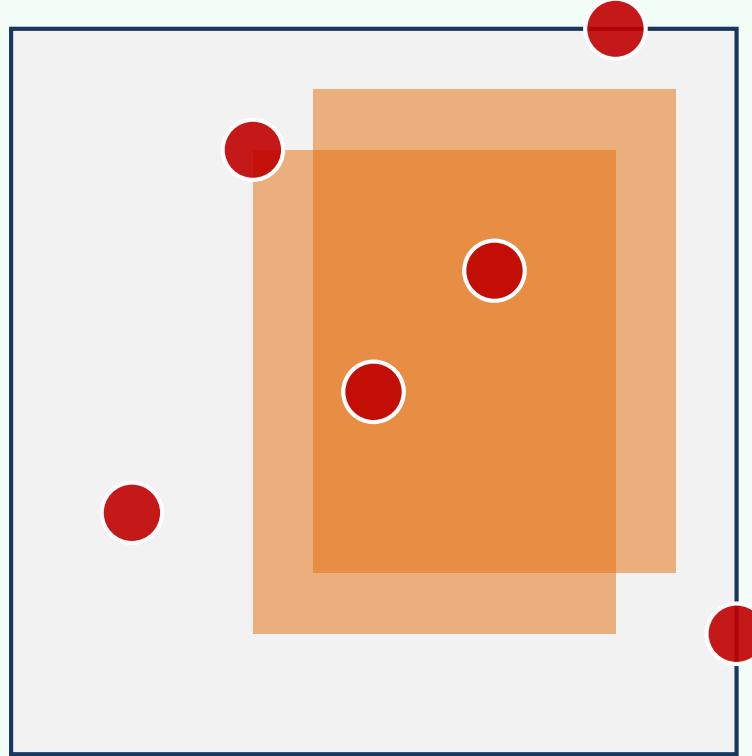
$$u \leq x \text{ and } v \leq y$$



Inclusion-Exclusion Principle

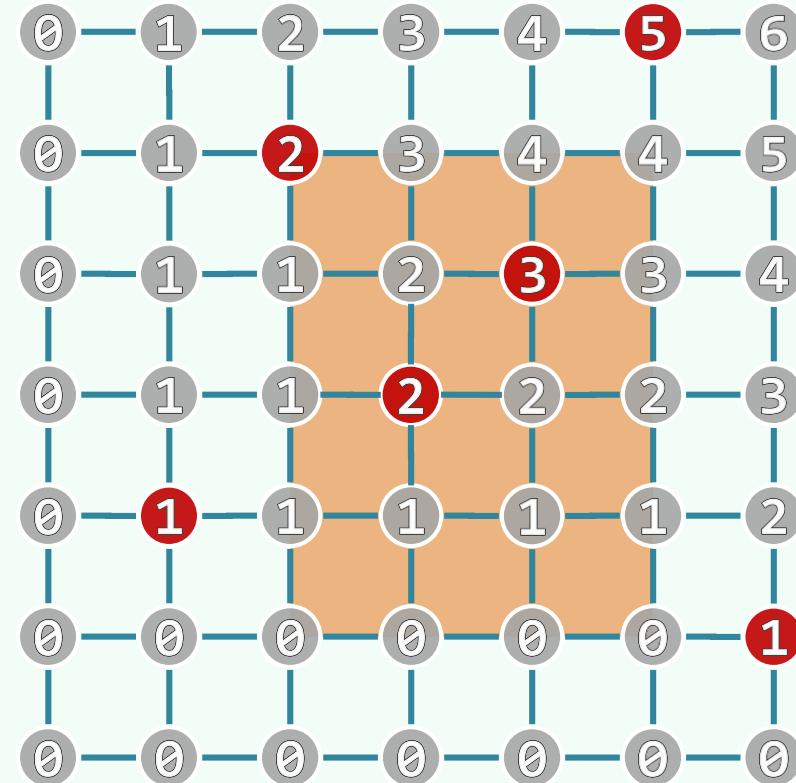
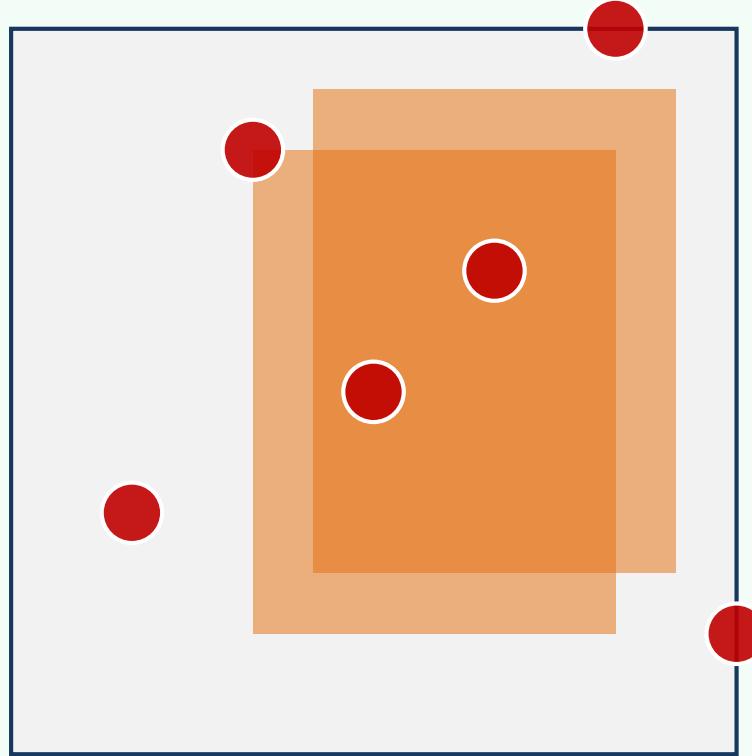
❖ Then for any rectangular range $\mathcal{R} = (x_1, x_2] \times (y_1, y_2]$, we have

$$|\mathcal{R} \cap \mathcal{P}| = n(x_1, y_1) + n(x_2, y_2) - n(x_1, y_2) - n(x_2, y_1)$$



Performance

- ❖ Each query needs only $\mathcal{O}(\log n)$ time
- ❖ Uses $\Theta(n^2)$ storage and even more for higher dimensions
- ❖ To find a better solution, let's go back to the 1D case ...



BST Application

kd-Tree: 1D

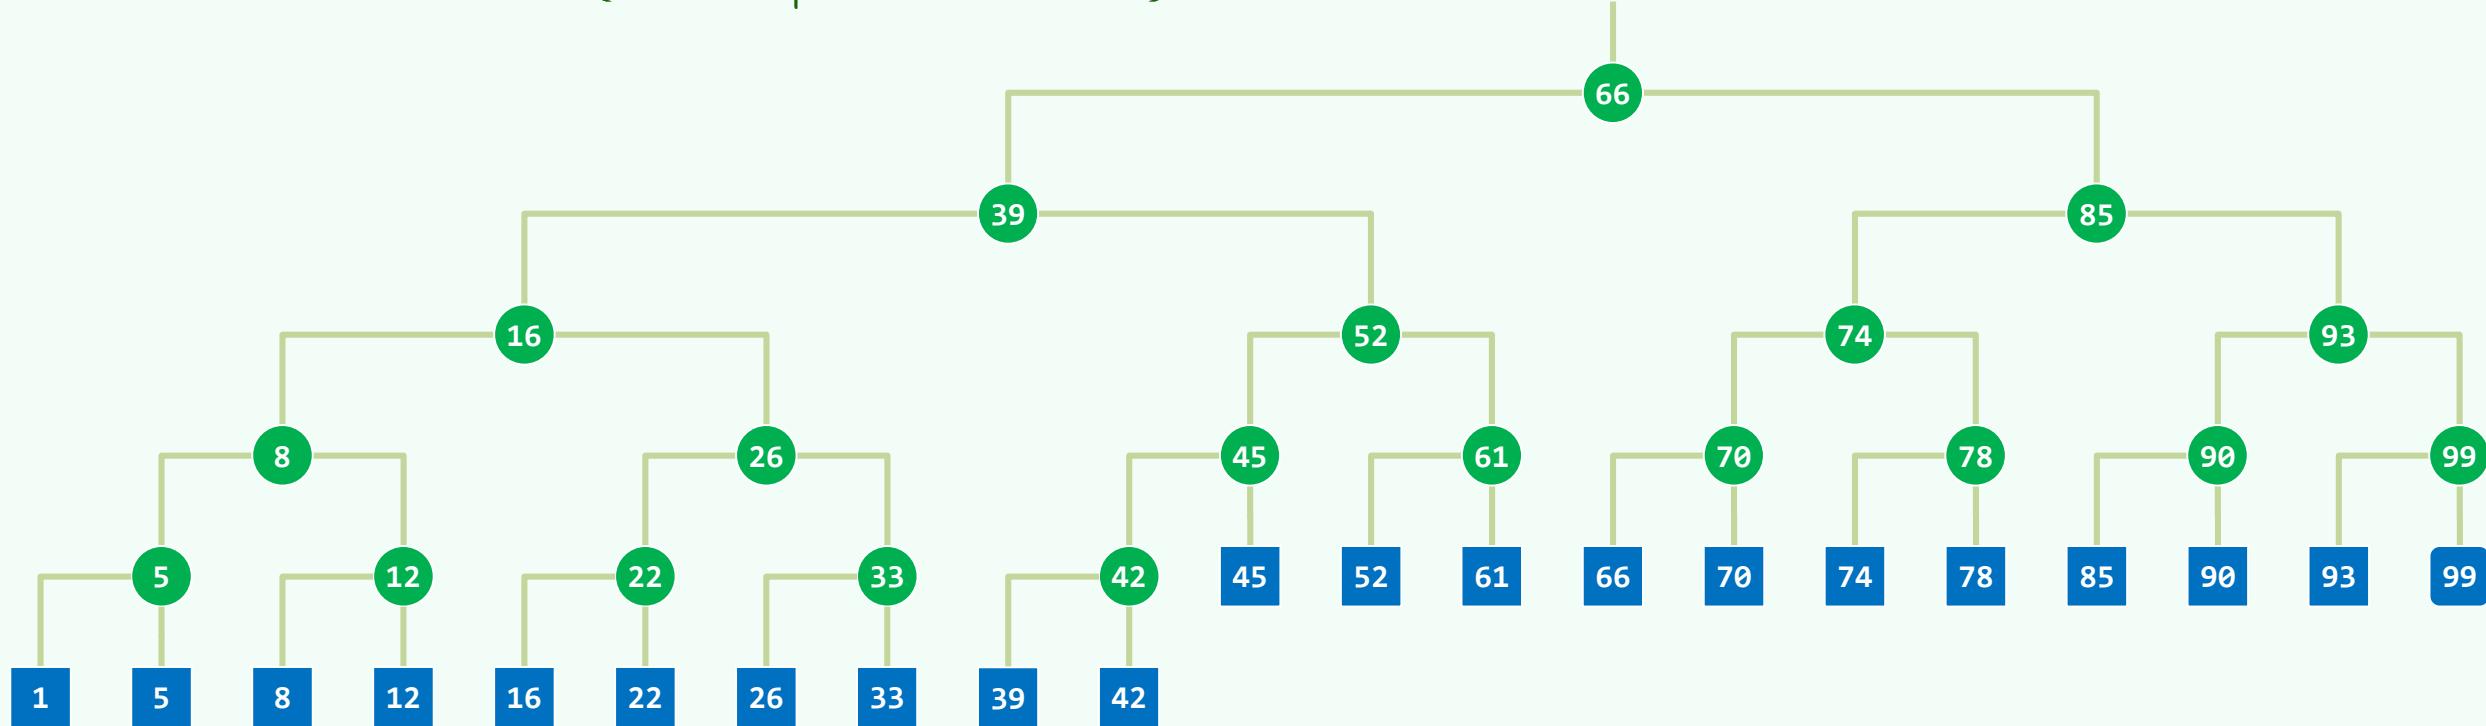
e>-B1

邓俊辉

deng@tsinghua.edu.cn

Structure: A Complete (Balanced) BST

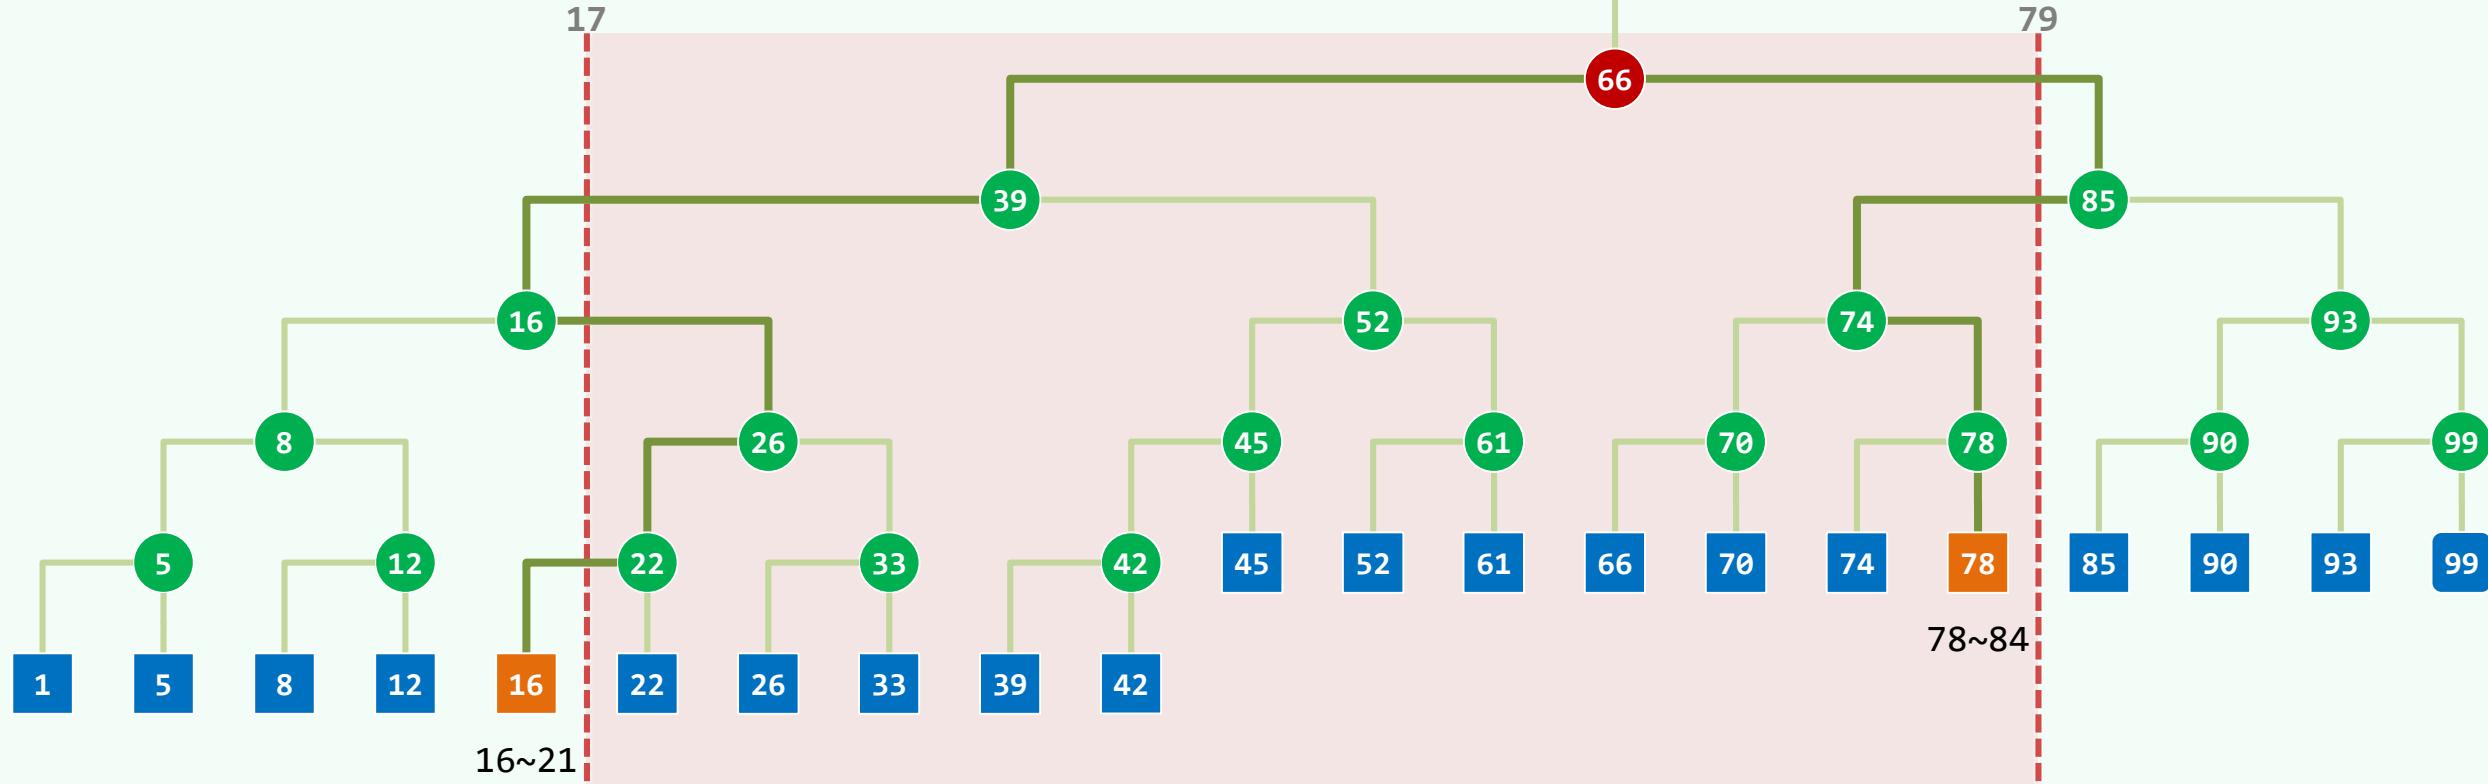
- ❖ $\forall v, v.key = \min\{ u.key \mid u \in v.rTree \} = v.succ.key$



- ❖ $\forall u \in v.lTree, u.key < v.key$
- ❖ $\forall u \in v.rTree, u.key \geq v.key$
- ❖ **search(x)** : returns the **MAXIMUM key not greater than x**

Lowest Common Ancestor

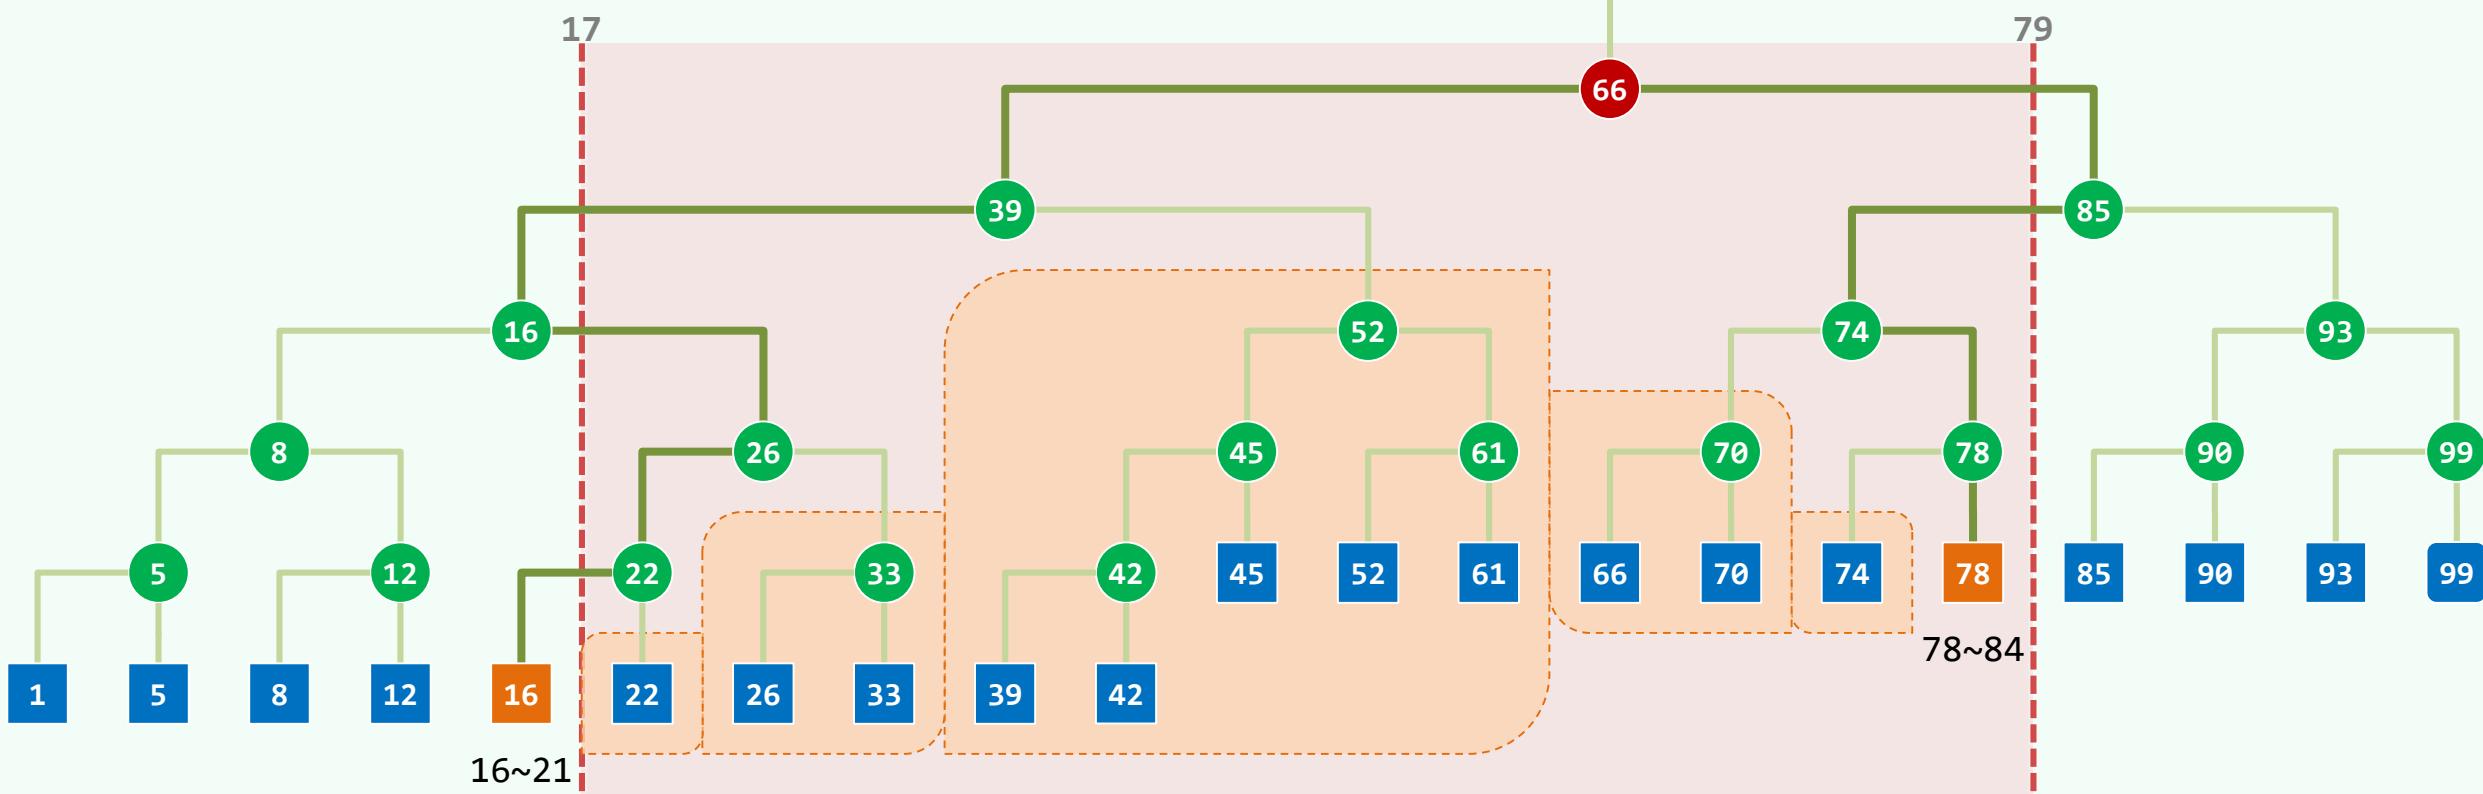
- ❖ Consider, as an example, the query for [17, 79] ...



- ❖ Do $\text{search}(17) = 16$ (might rejected) and $\text{search}(79) = 78$ (must accepted)
- ❖ Consider $\text{LCA}(16, 78) = 66 \dots$

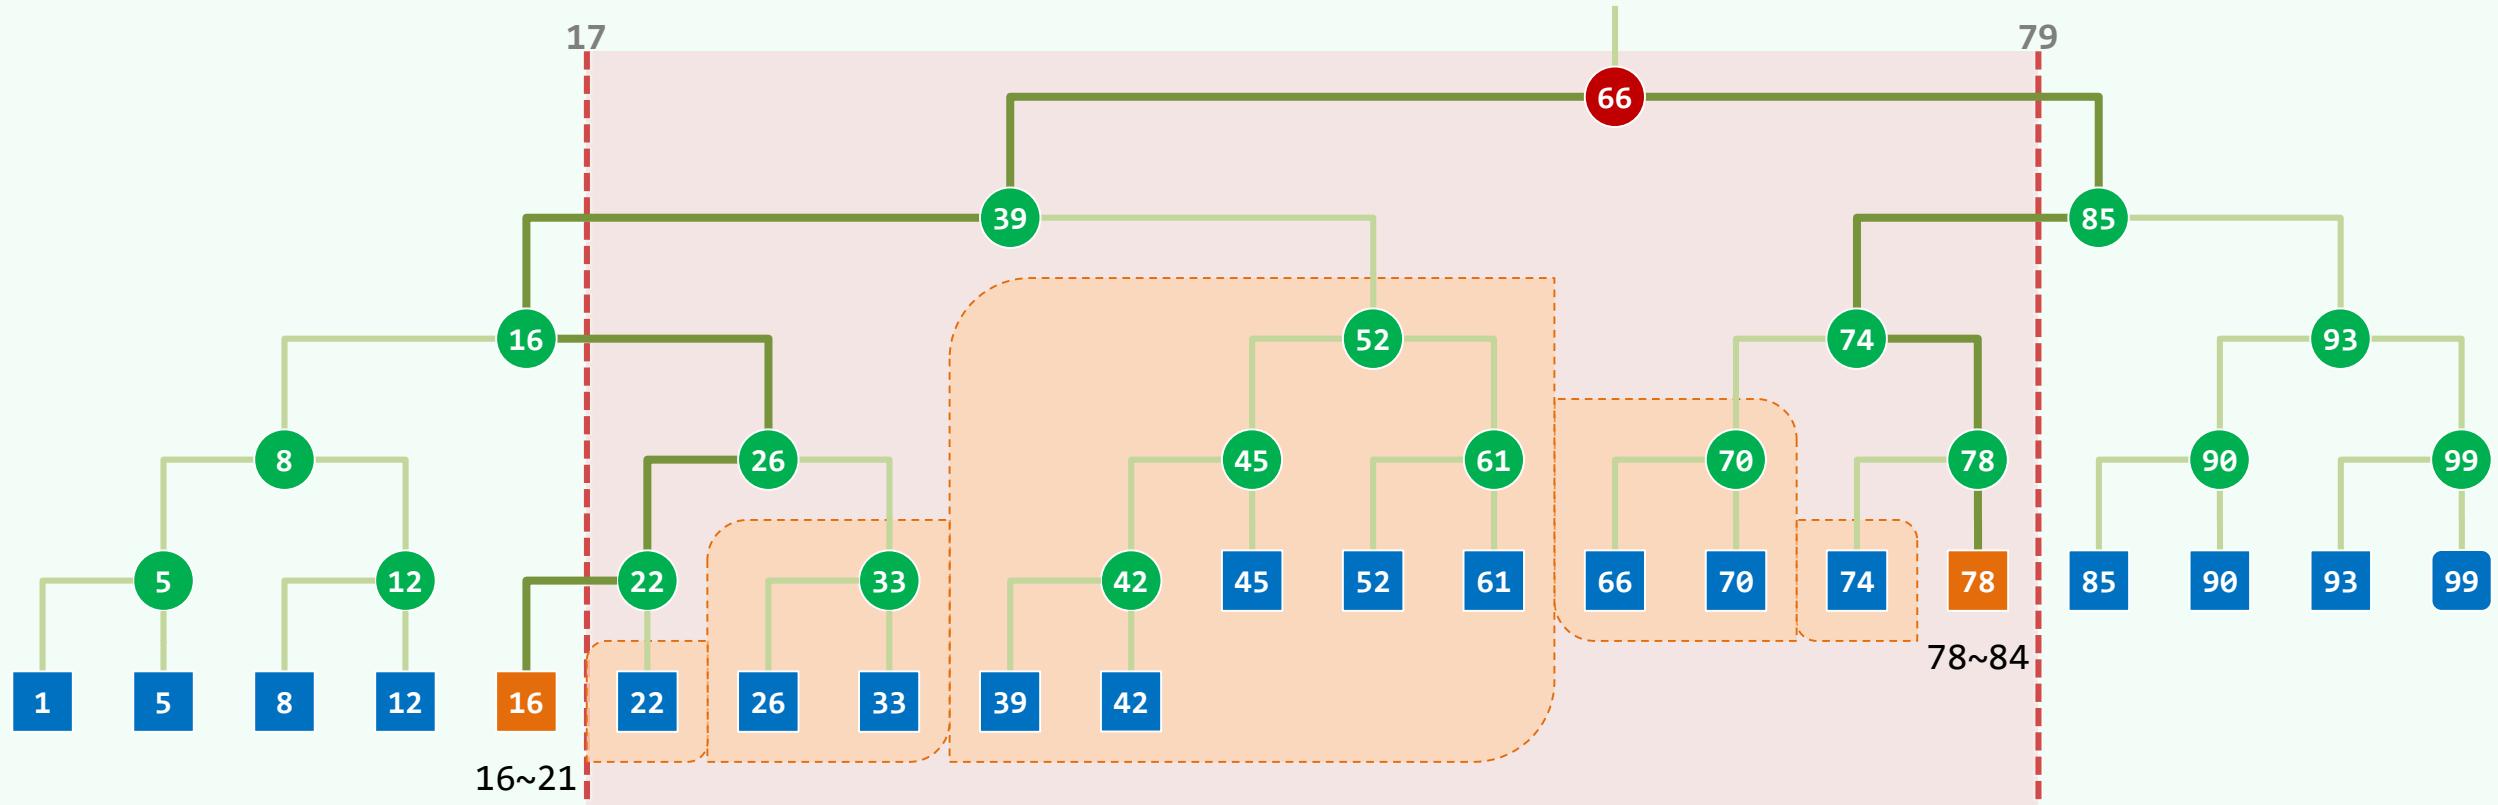
Union of $\theta(\log n)$ Disjoint Subtrees

- ❖ Starting from the LCA, traverse path(16) and path(78) once more resp.



- All R/L-turns along path(16)/path(78) are ignored and
- the R/L subtree at each L/R-turn is reported

Complexity



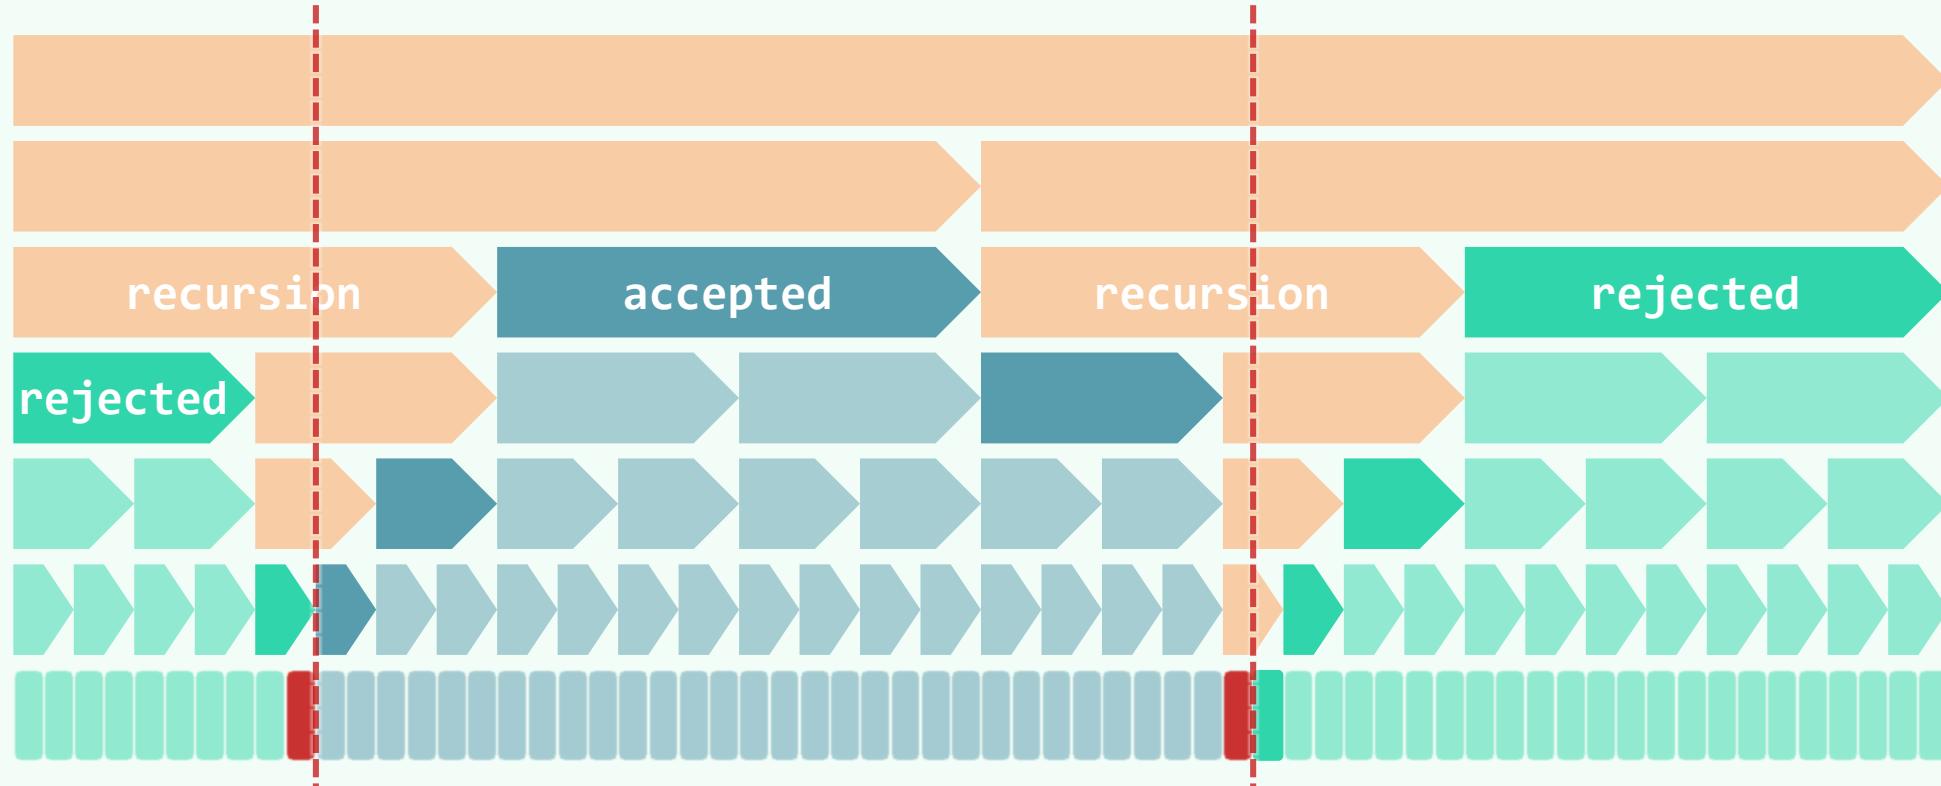
Query: $\mathcal{O}(\log n)$

Preprocessing: $\mathcal{O}(n \log n)$

Storage: $\mathcal{O}(n)$

Hot Knives Through A Chocolate Cake of Height $\theta(\log n)$

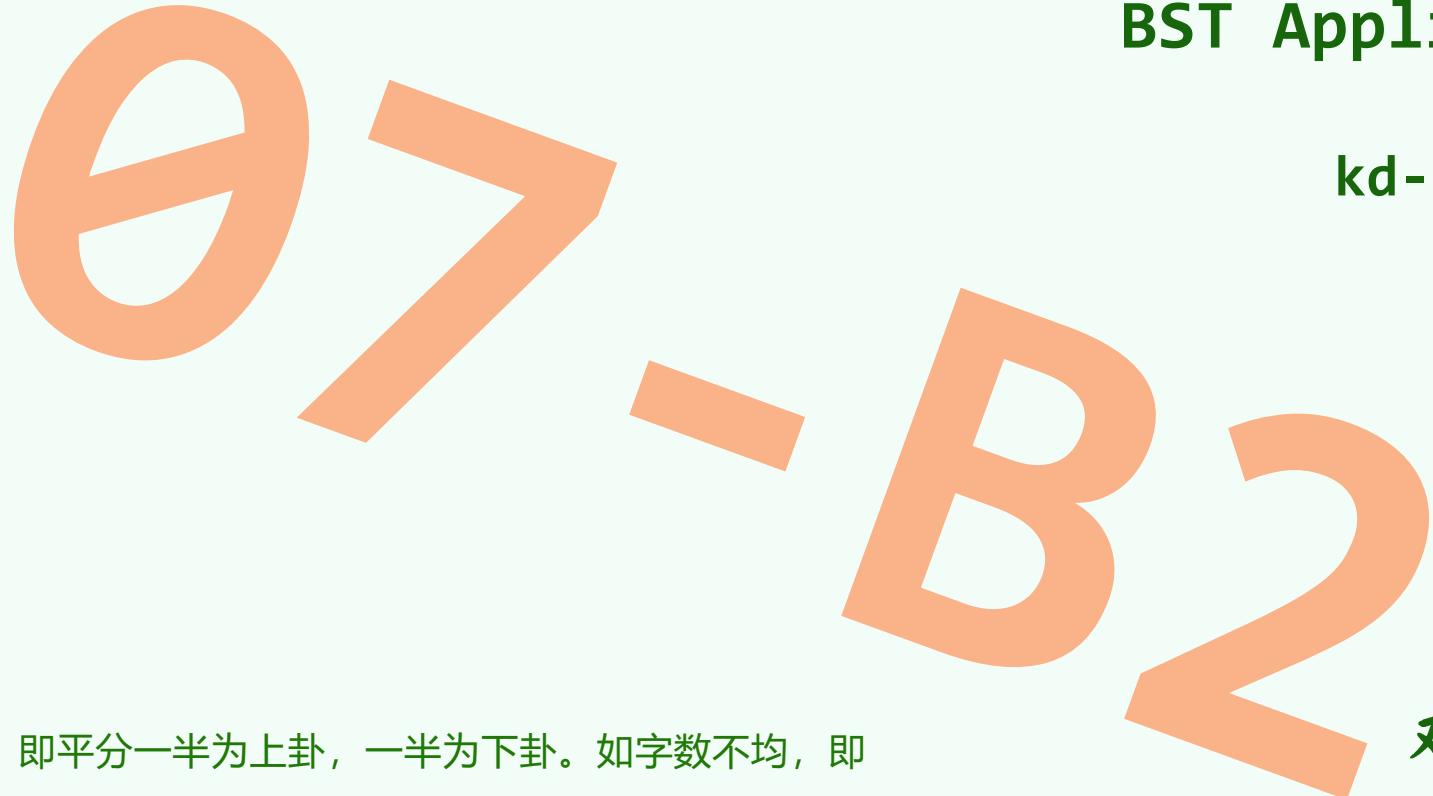
- ❖ Region(u) is enclosed by Region(v) iff u is a descendant of v in the 1d-tree
- ❖ Region(u) and Region(v) are disjoint iff neither is the ancestor of the other



- ❖ All nodes are partitioned into 3 types: accepted + rejected + recursion

BST Application

kd-Tree: 2D



凡见字数，如停匀，即平分一半为上卦，一半为下卦。如字数不均，即少一字为上卦，取天轻清之义，以多一字为下卦，取地重浊之义

邓俊辉

deng@tsinghua.edu.cn

Divide-And-Conquer

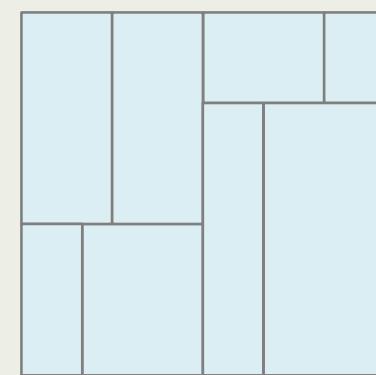
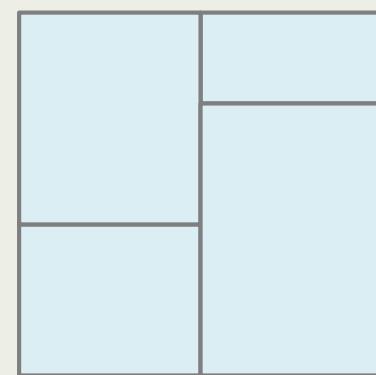
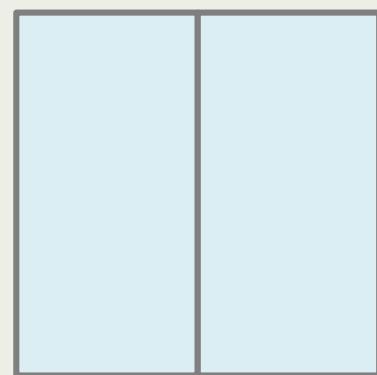
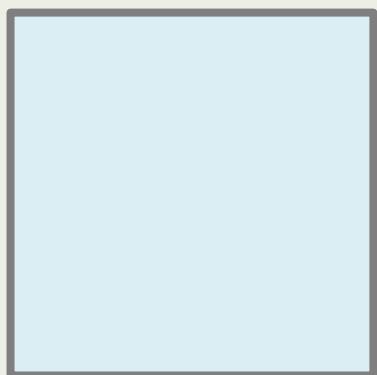
❖ To extend the BBST method to planar GRS, we

- **divide** the plane recursively and
- **arrange** the regions into a kd-tree

❖ Start with a single region (the entire plane)

Partition the region vertically/horizontally on each even/odd level

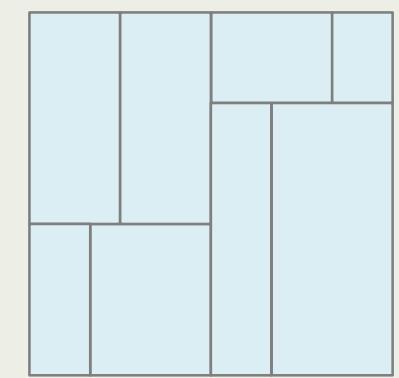
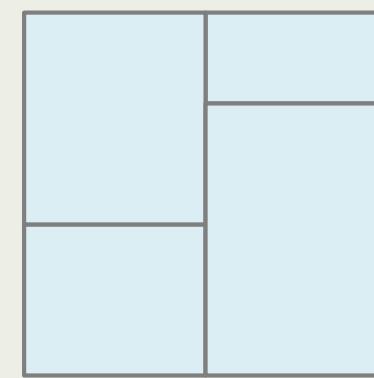
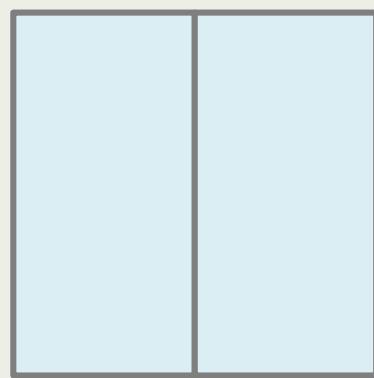
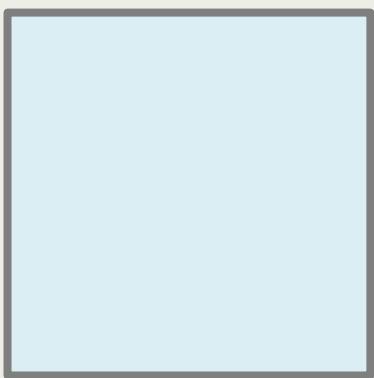
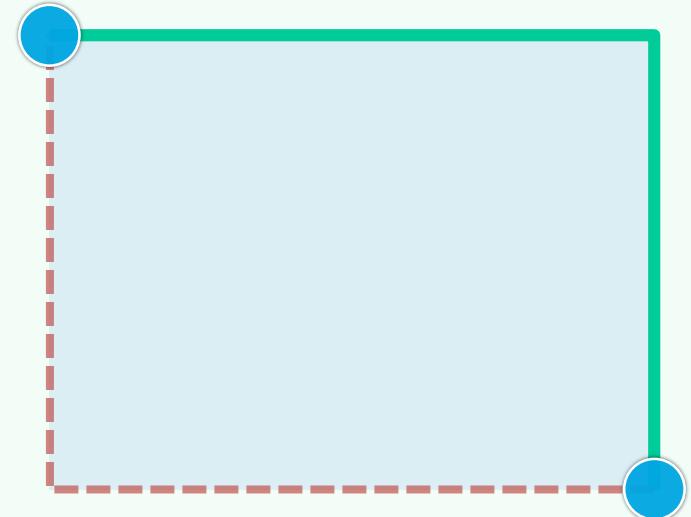
Partition the sub-regions recursively



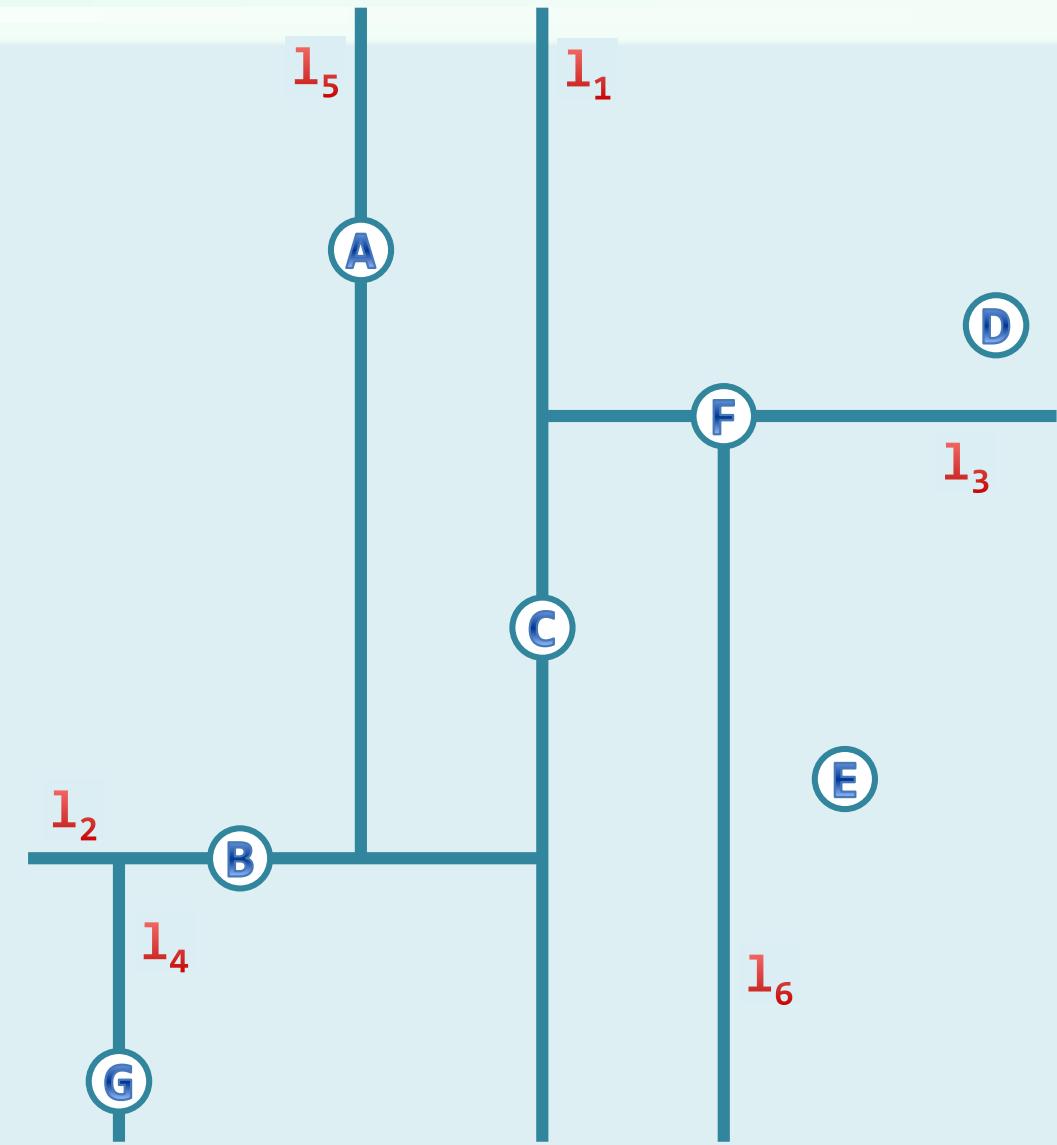
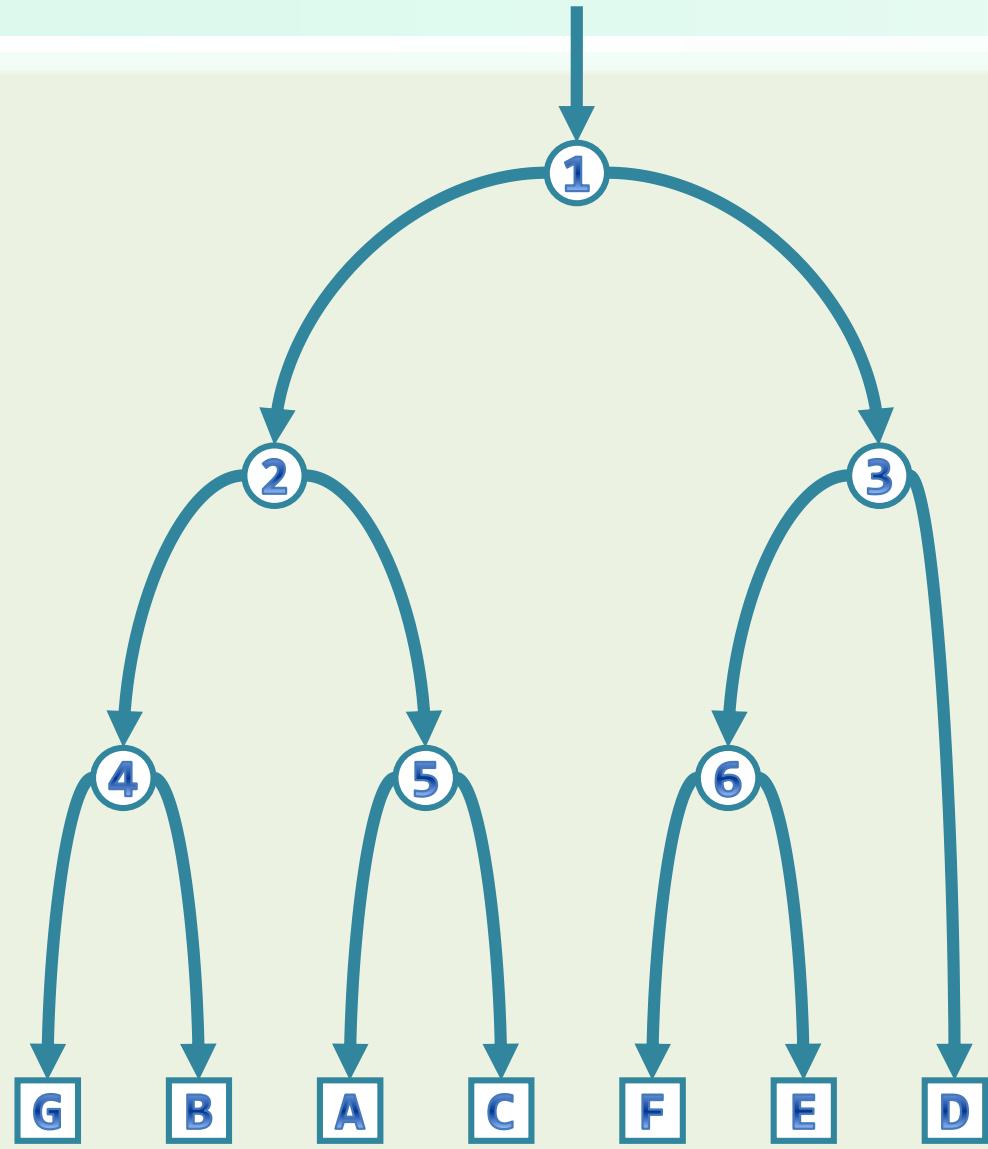
More Details

❖ To make it work,

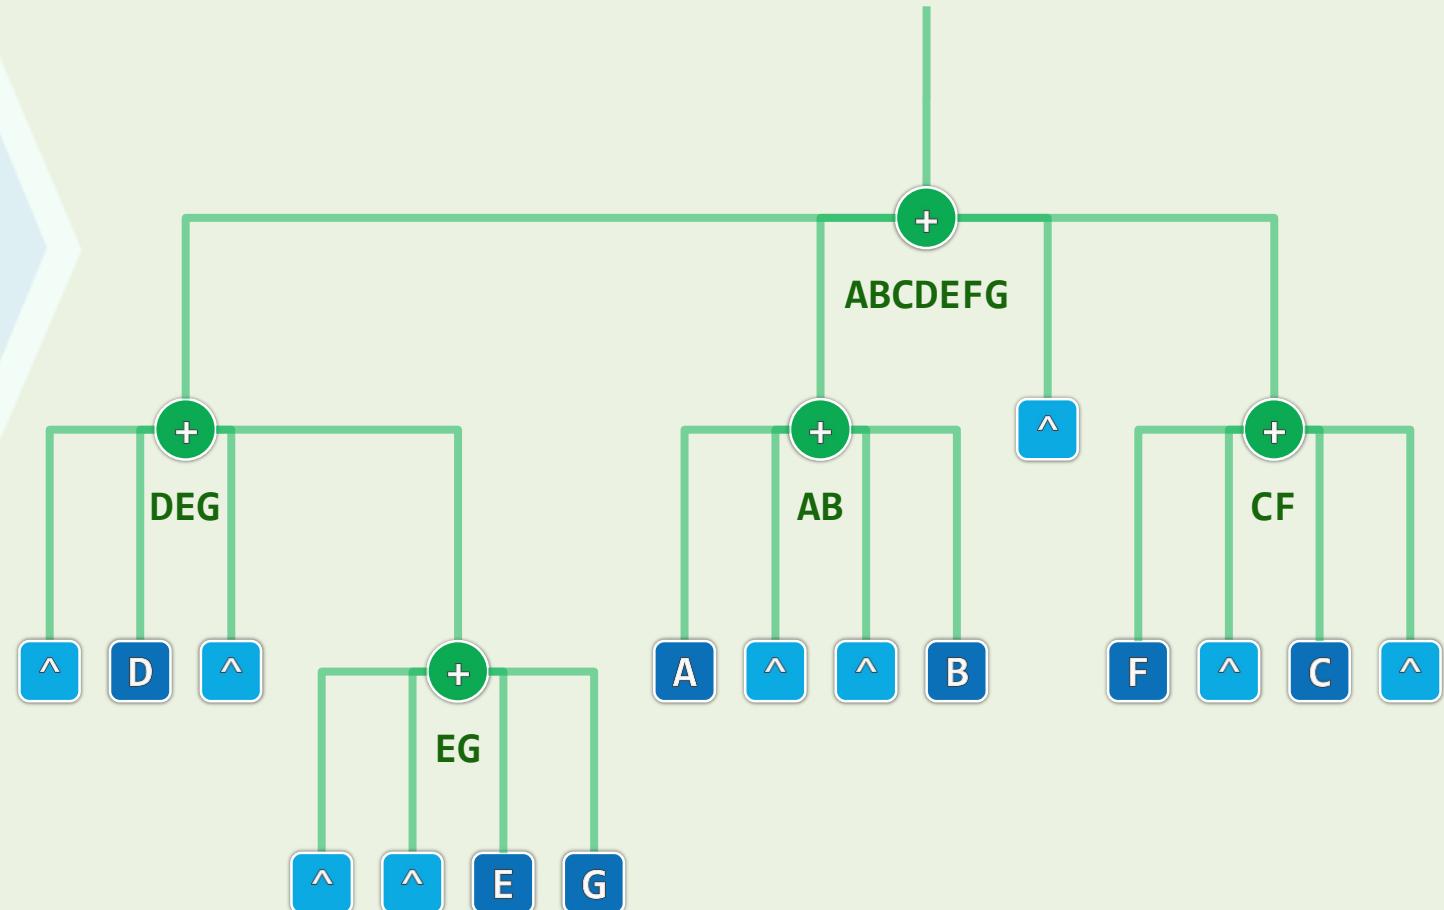
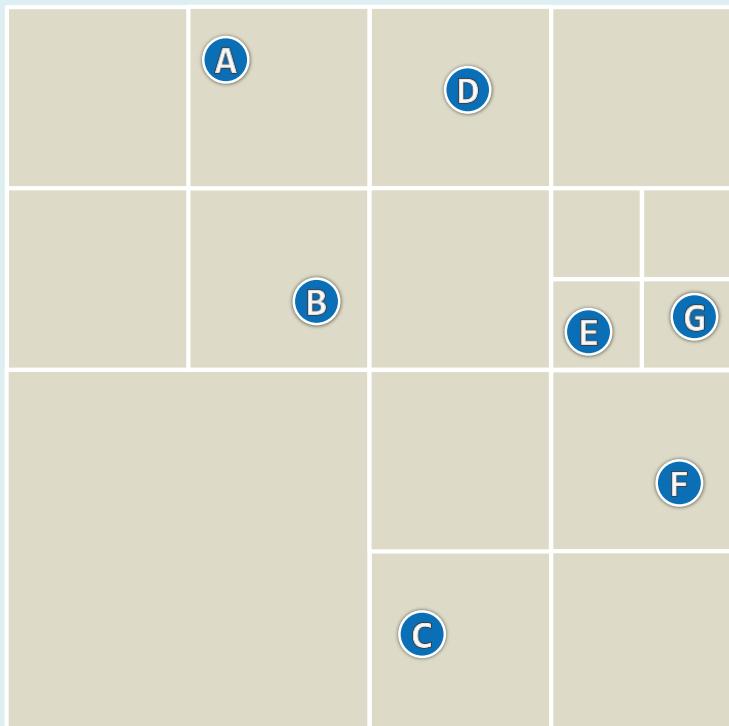
- each partition should be done as **evenly** as possible (at median)
- each region is defined to be **open/closed** on the **left-lower/right-upper** sides



Example



Quadtree



Θ>-B3

BST Application

kd-Tree: Construction

邓俊辉

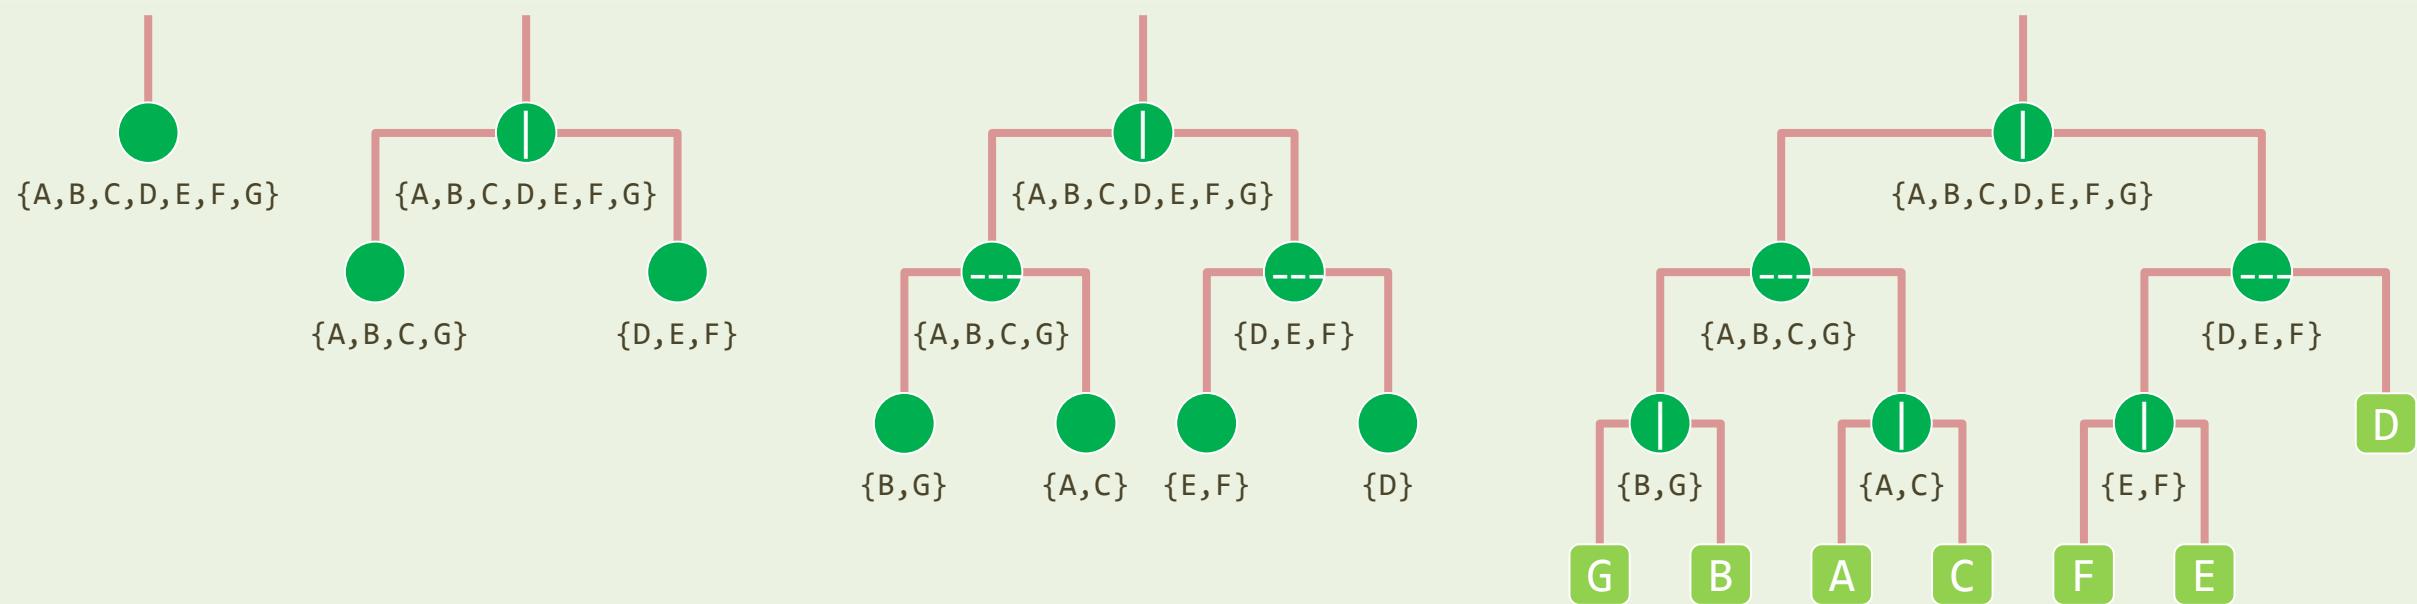
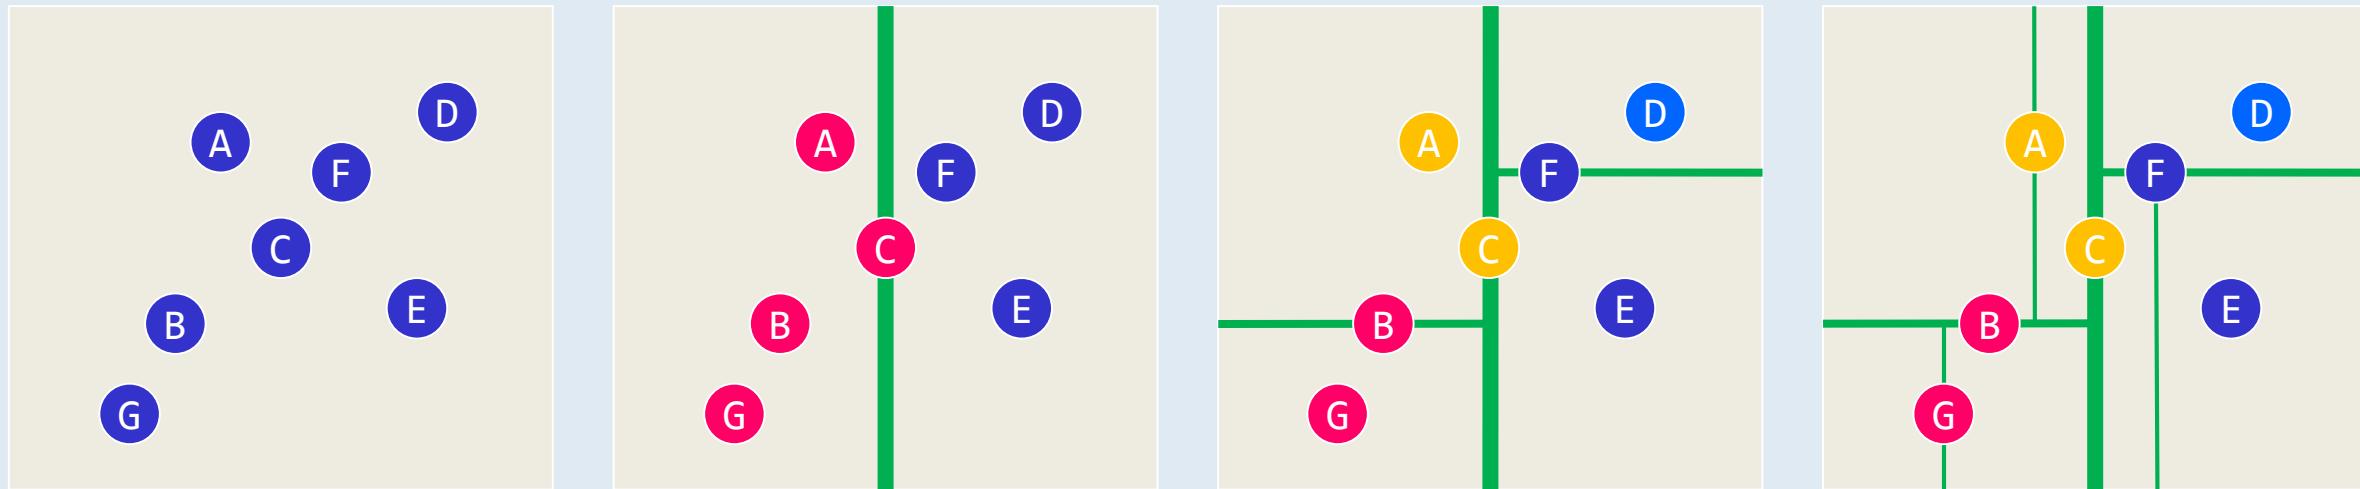
deng@tsinghua.edu.cn

God found himself by creating.

buildKdTree(P,d)

```
// construct a 2d-(sub)tree for point (sub)set P at depth d  
  
if ( P == {p} ) return CreateLeaf( p ) //base  
  
root = CreateKdNode()  
  
root->splitDirection = Even(d) ? VERTICAL : HORIZONTAL  
  
root->splitLine = FindMedian( root->splitDirection, P ) // $\mathcal{O}(n)$ !  
  
( P1, P2 ) = Divide( P, root->splitDirection, root->splitLine ) //DAC  
  
root->lc = buildKdTree( P1, d + 1 ) //recurse  
  
root->rc = buildKdTree( P2, d + 1 ) //recurse  
  
return( root )
```

Example



BST Application

kd-Tree: Canonical Subsets

07-B4

邓俊辉

deng@tsinghua.edu.cn

韦小宝跟著她走到桌边，只见桌上大白布上钉满了几千枚绣花针，几千块碎片已拼成一幅完整无缺的大地图，难得的是几千片碎皮拼在一起，既没多出一片，也没少了一片。

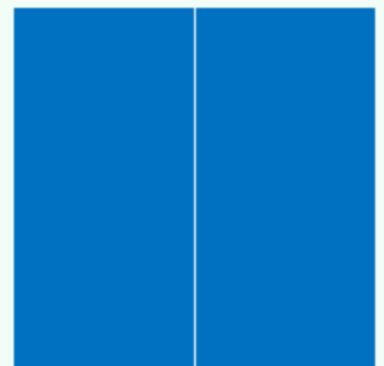
Canonical Subset

- ⦿ Each node corresponds to
 - a rectangular **sub-region** of the plane, as well as
 - the **subset** of points contained in the sub-region

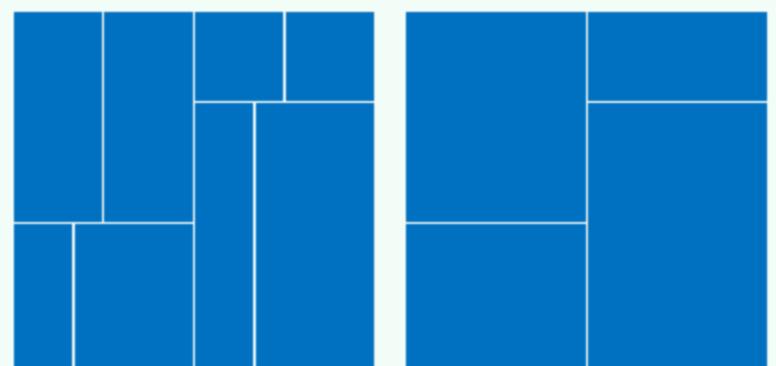


❖ Each of these subsets is called a **canonical subset**

- ⦿ For each **internal node X** with children L and R,
 $\text{region}(X) = \text{region}(L) \cup \text{region}(R)$

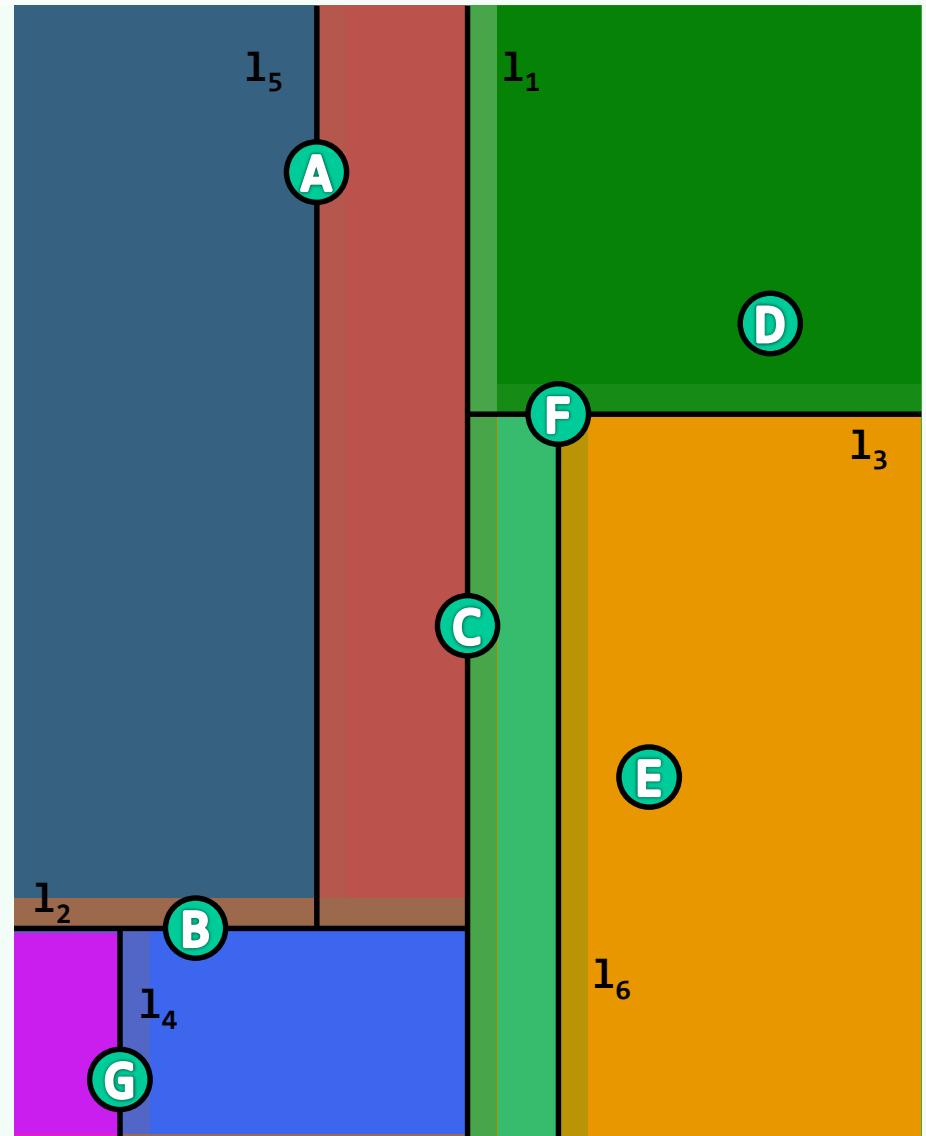
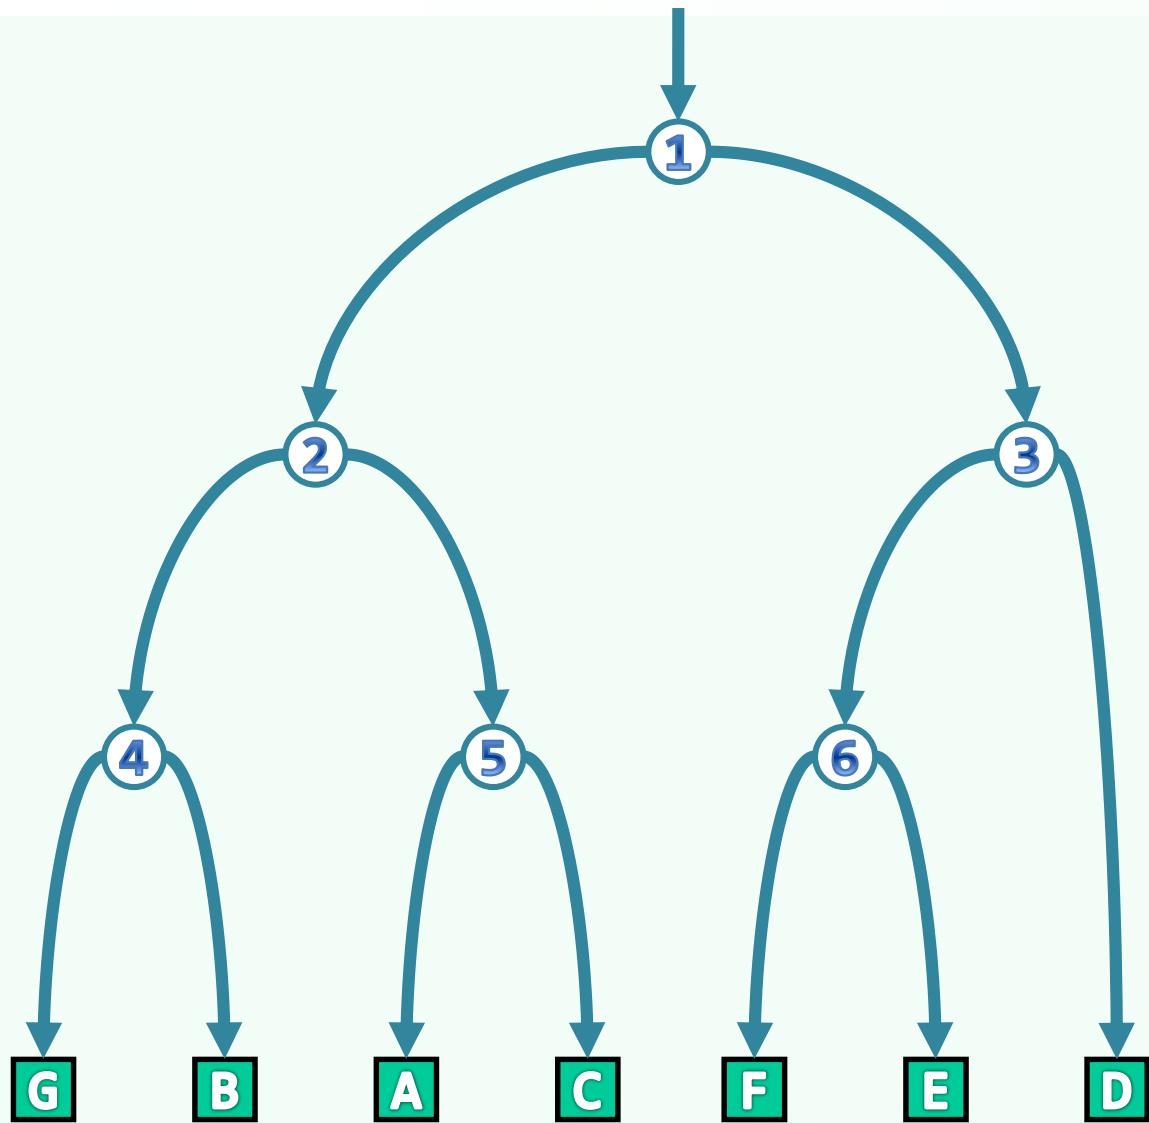


- ⦿ Sub-regions of nodes at a same depth
 - never **intersect** with each other, and
 - their **union** covers the entire plane



❖ We will see soon that each 2D GRS can be
answered by the **union** of a number of CS's

Example



BST Application

kd-Tree: Query

07-B5 邓俊辉

我们竟为这无用的找寻浪费了这么多天!

我想找寻的心上人绝对不会在这里出现。

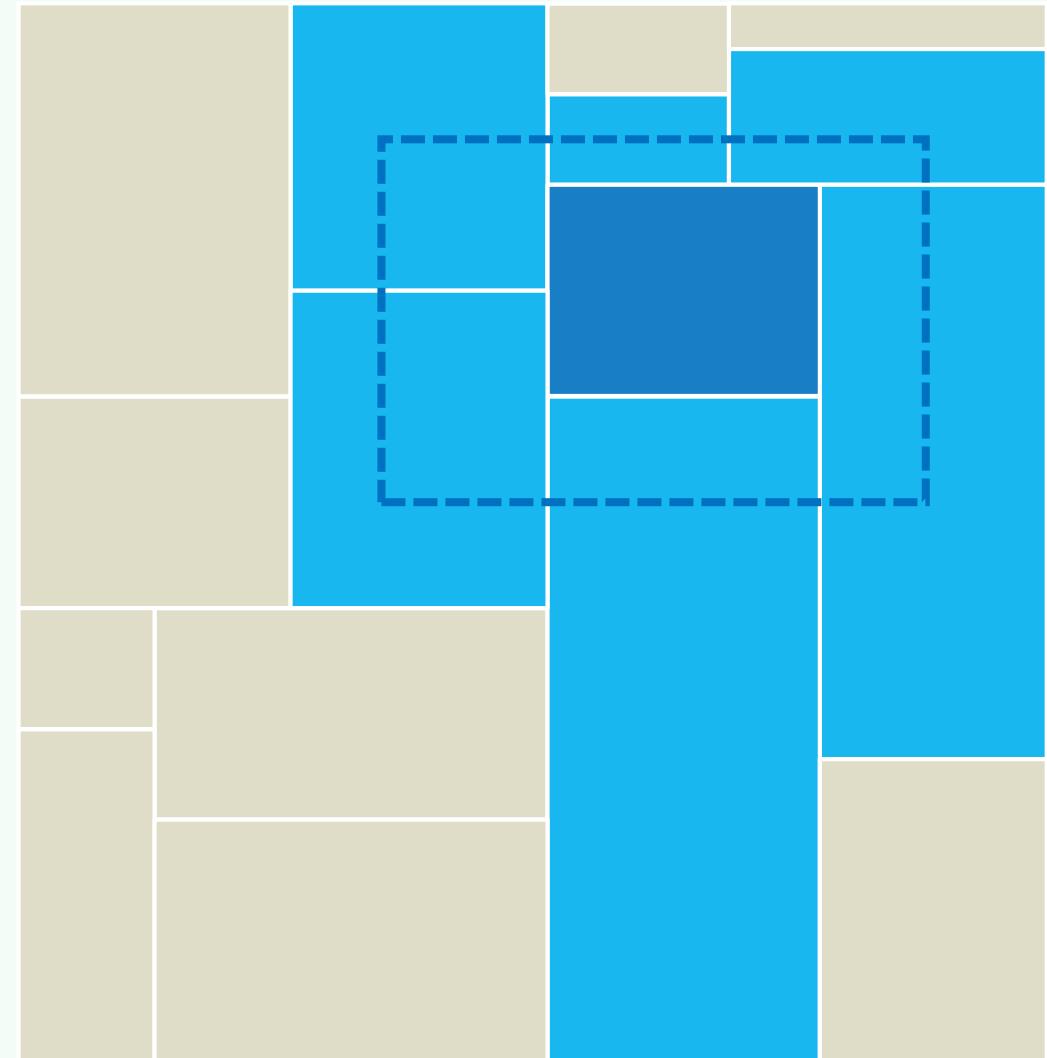
deng@tsinghua.edu.cn

kdSearch(v, R): 热刀来切~~干~~ (logn) 层巧克力

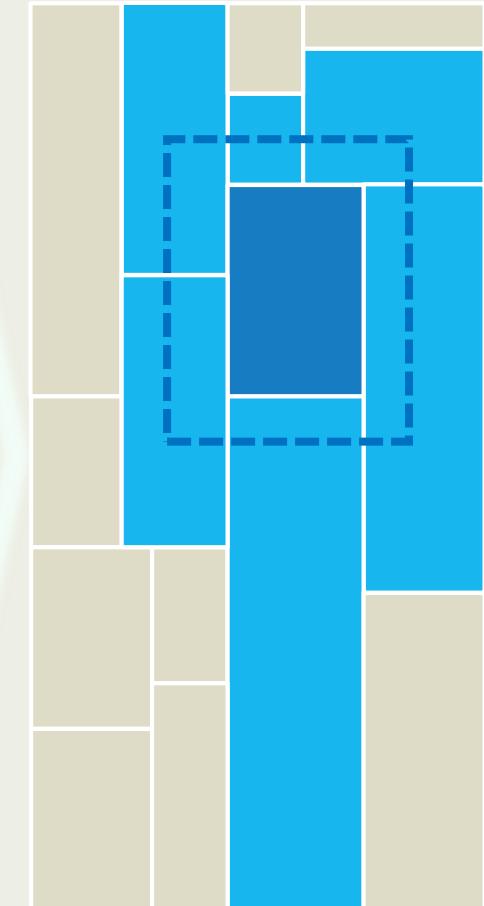
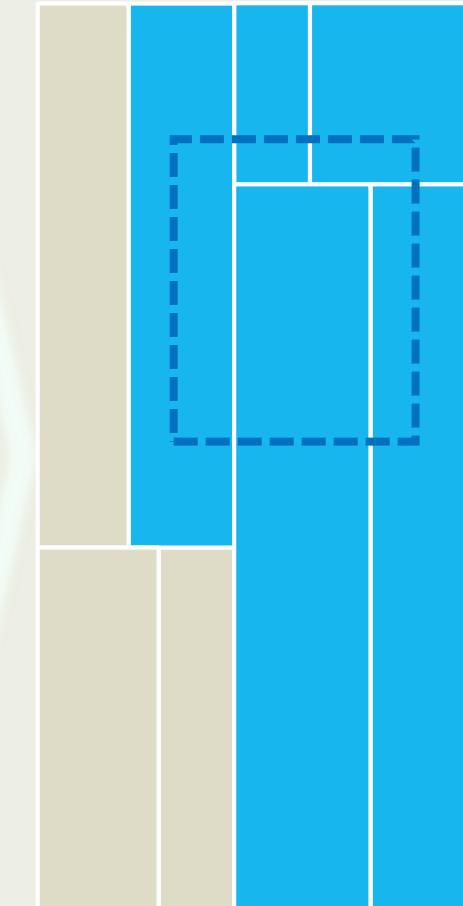
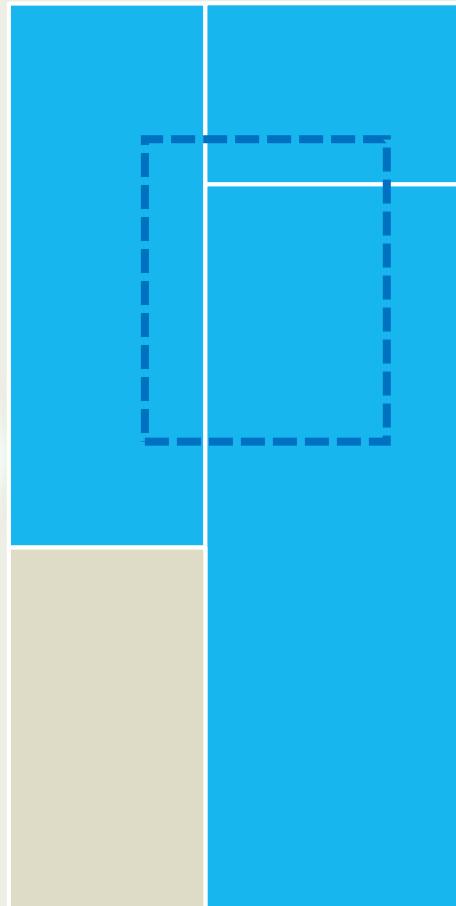
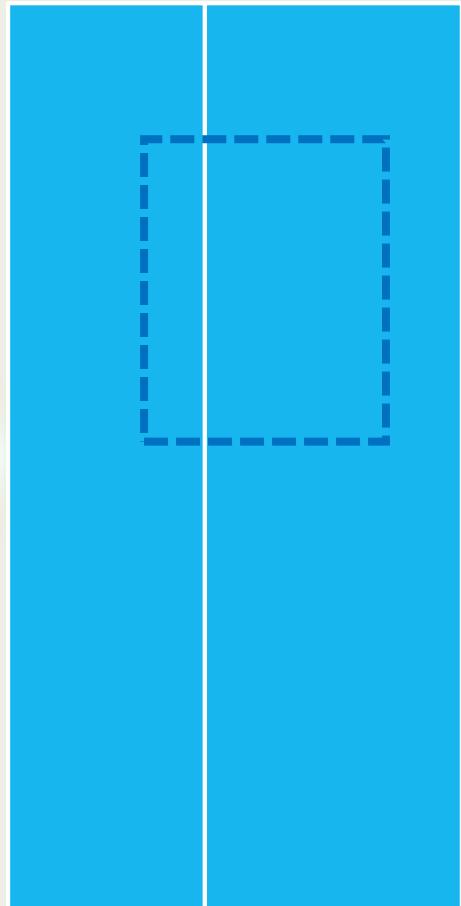
```
❖ if ( isLeaf( v ) )
    if ( inside( v, R ) ) report(v)
    return

❖ if ( region( v->lc ) ⊆ R )
    reportSubtree( v->lc )
else if ( region( v->lc ) ∩ R ≠ ∅ )
    kdSearch( v->lc, R )

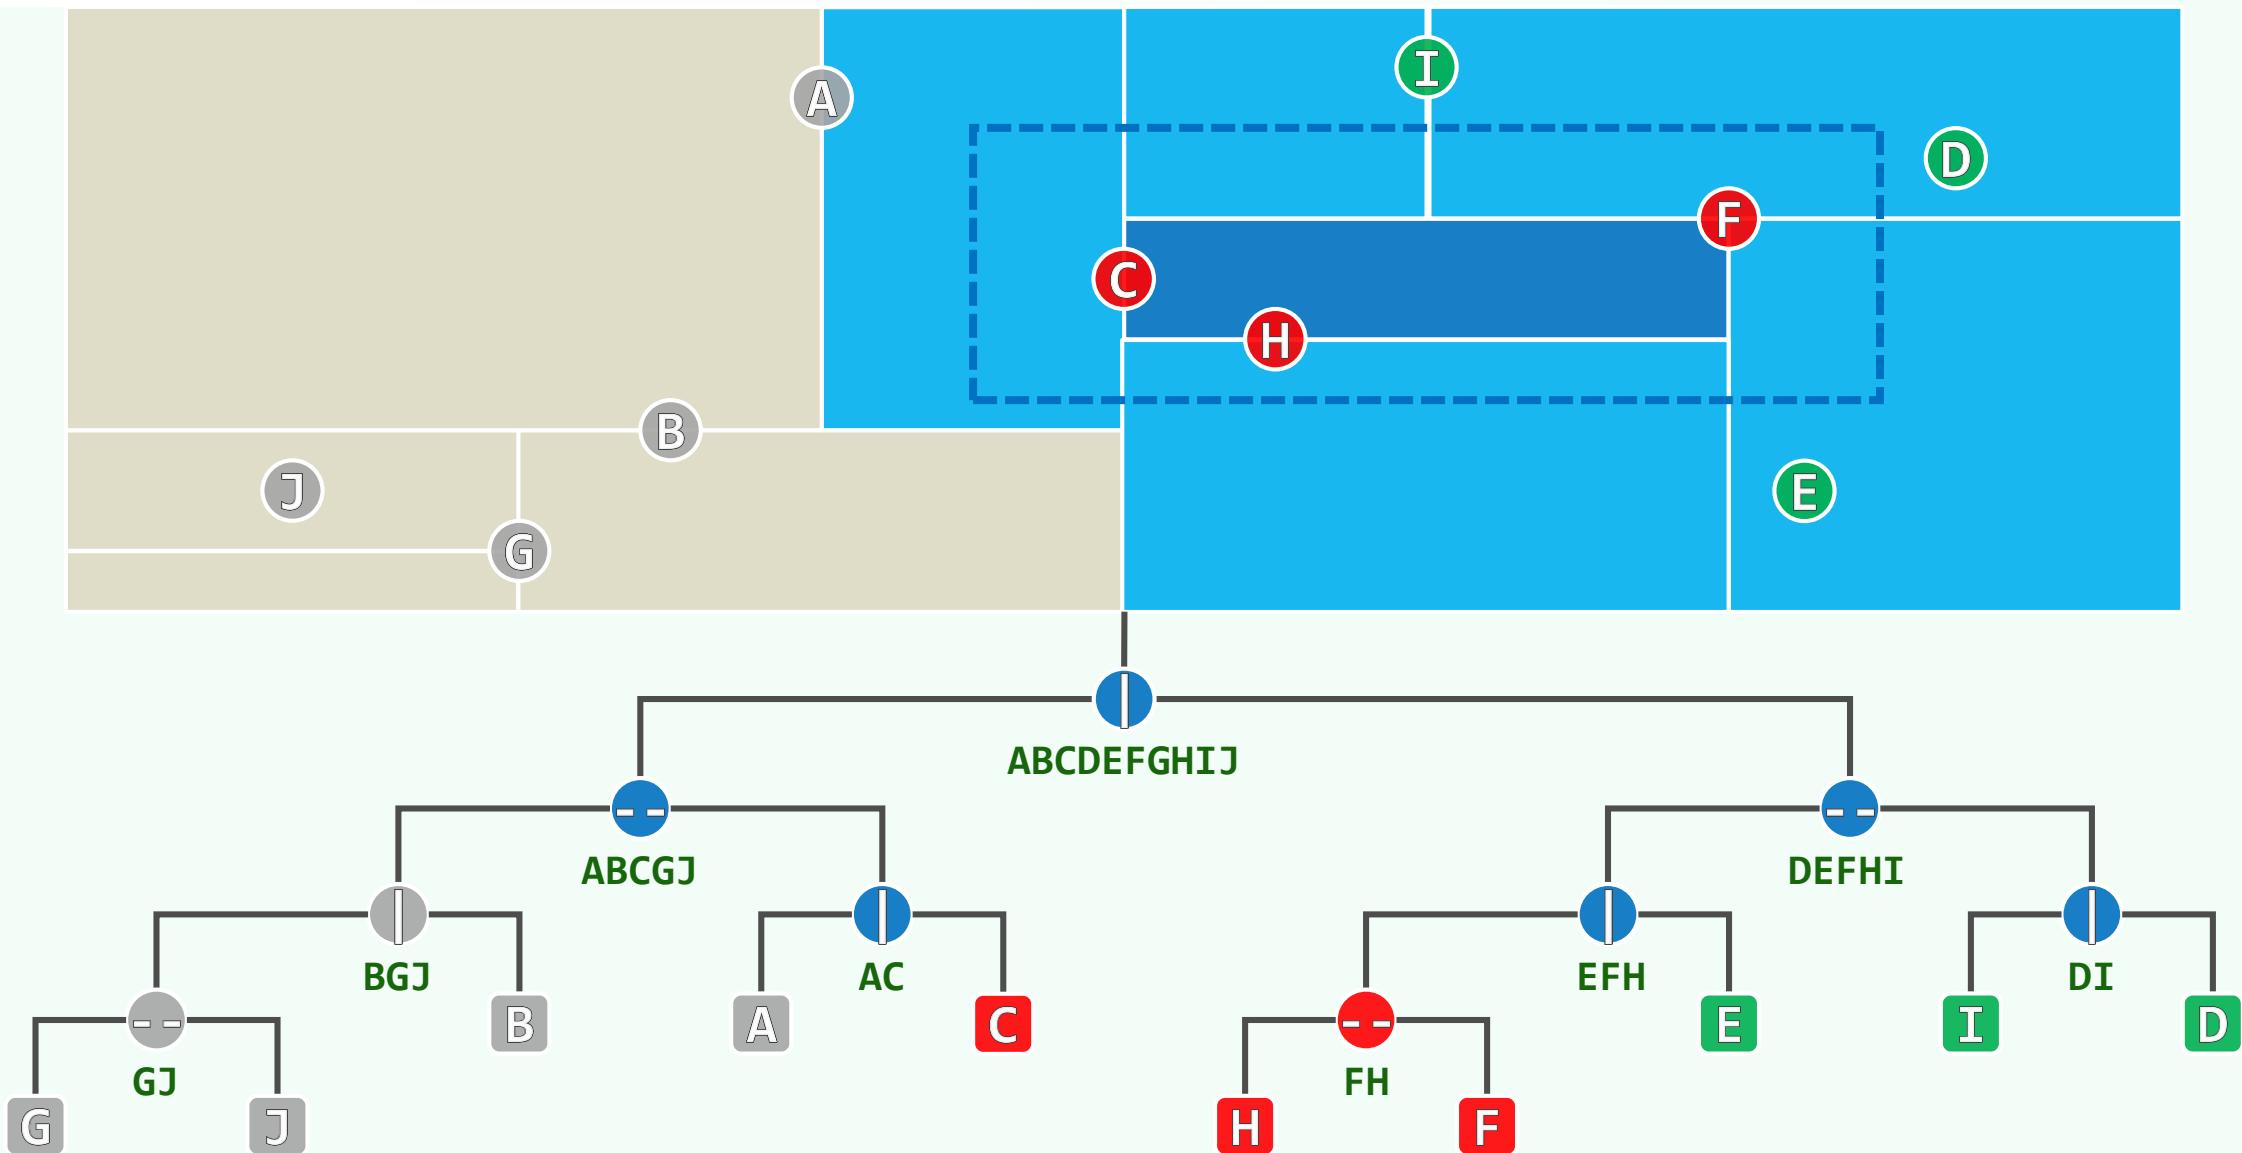
❖ if ( region( v->rc ) ⊆ R )
    reportSubtree( v->rc )
else if ( region( v->rc ) ∩ R ≠ ∅ )
    kdSearch( v->rc, R )
```



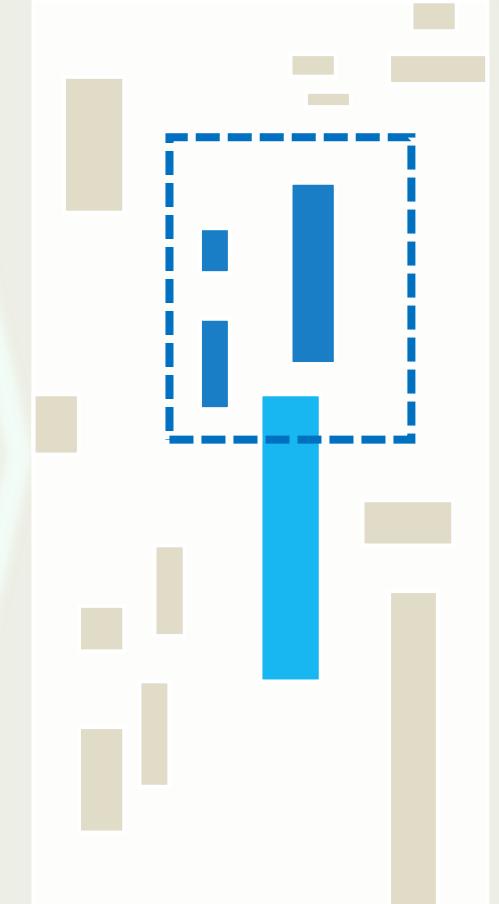
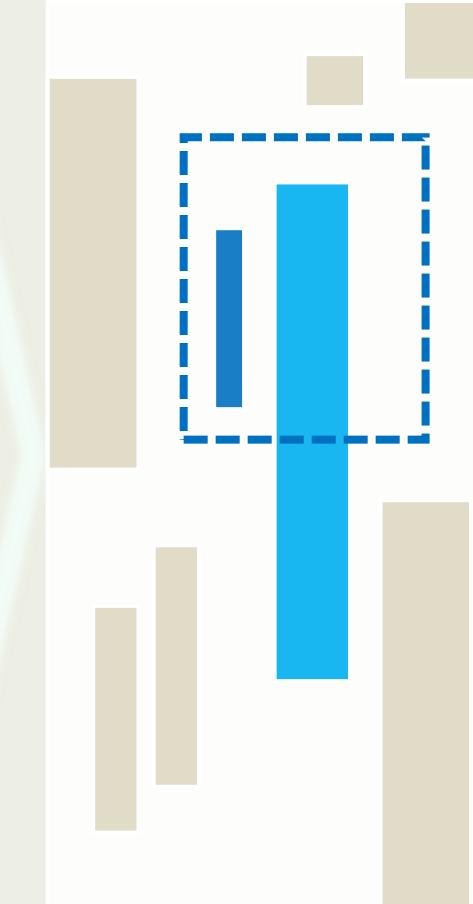
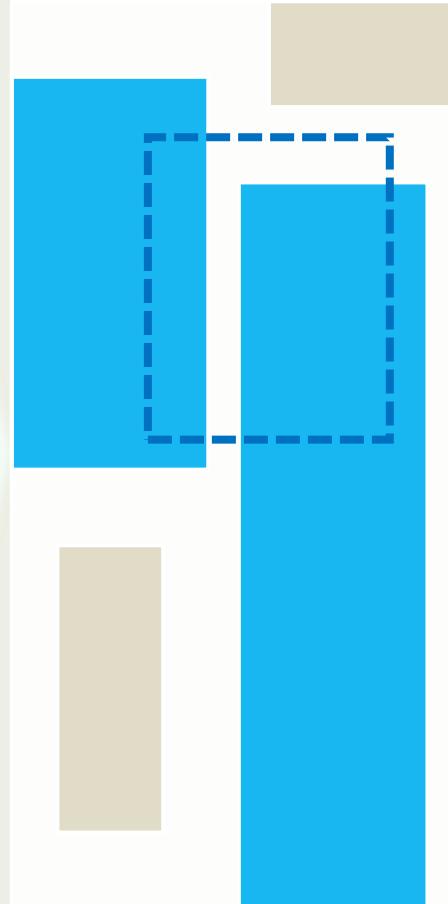
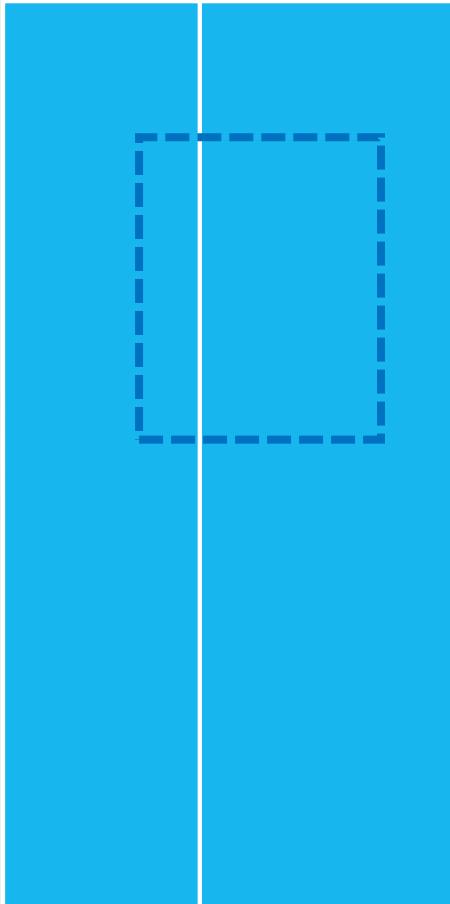
Example



Example



Bounding Box



BST Application

kd-Tree: Complexity

θ>B6

肉眼看不清细节，但他们都知道那是木星所在的位置，这颗太阳系最大的行星已经坠落到二维平面上了。

有人嘲笑这种体系说：为了能发现这个比例中项并组成政府共同体，按照我的办法，只消求出人口数字的平方根就行了。

邓俊辉

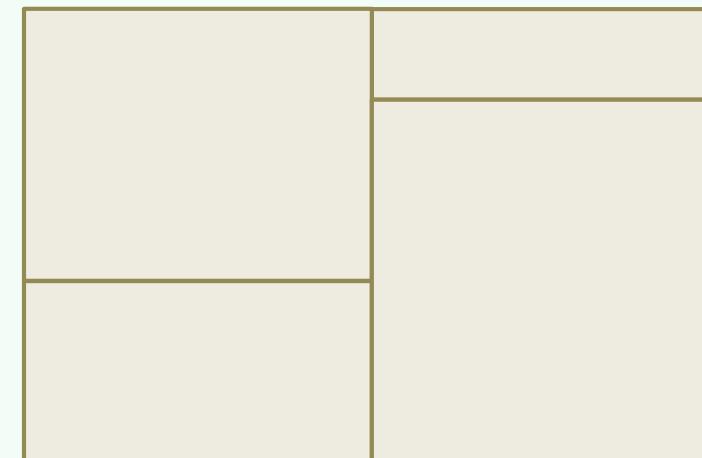
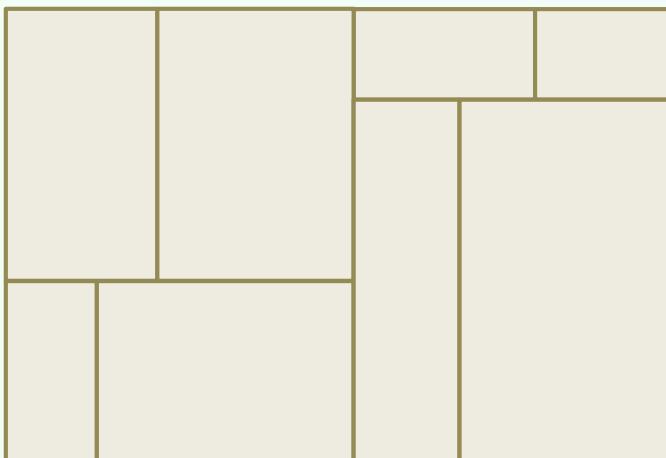
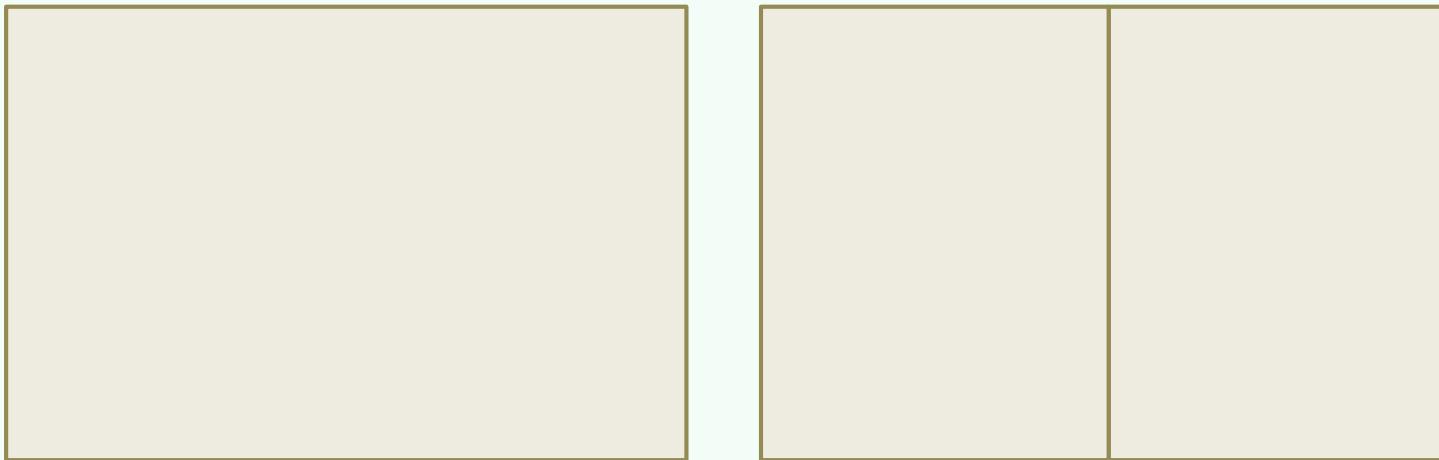
deng@tsinghua.edu.cn

Preprocessing

❖ $T(n)$

$$= 2*T(n/2) + O(n)$$

$$= O(n \log n)$$



Storage

❖ The tree has a height

of $\Theta(\log n)$

❖

1

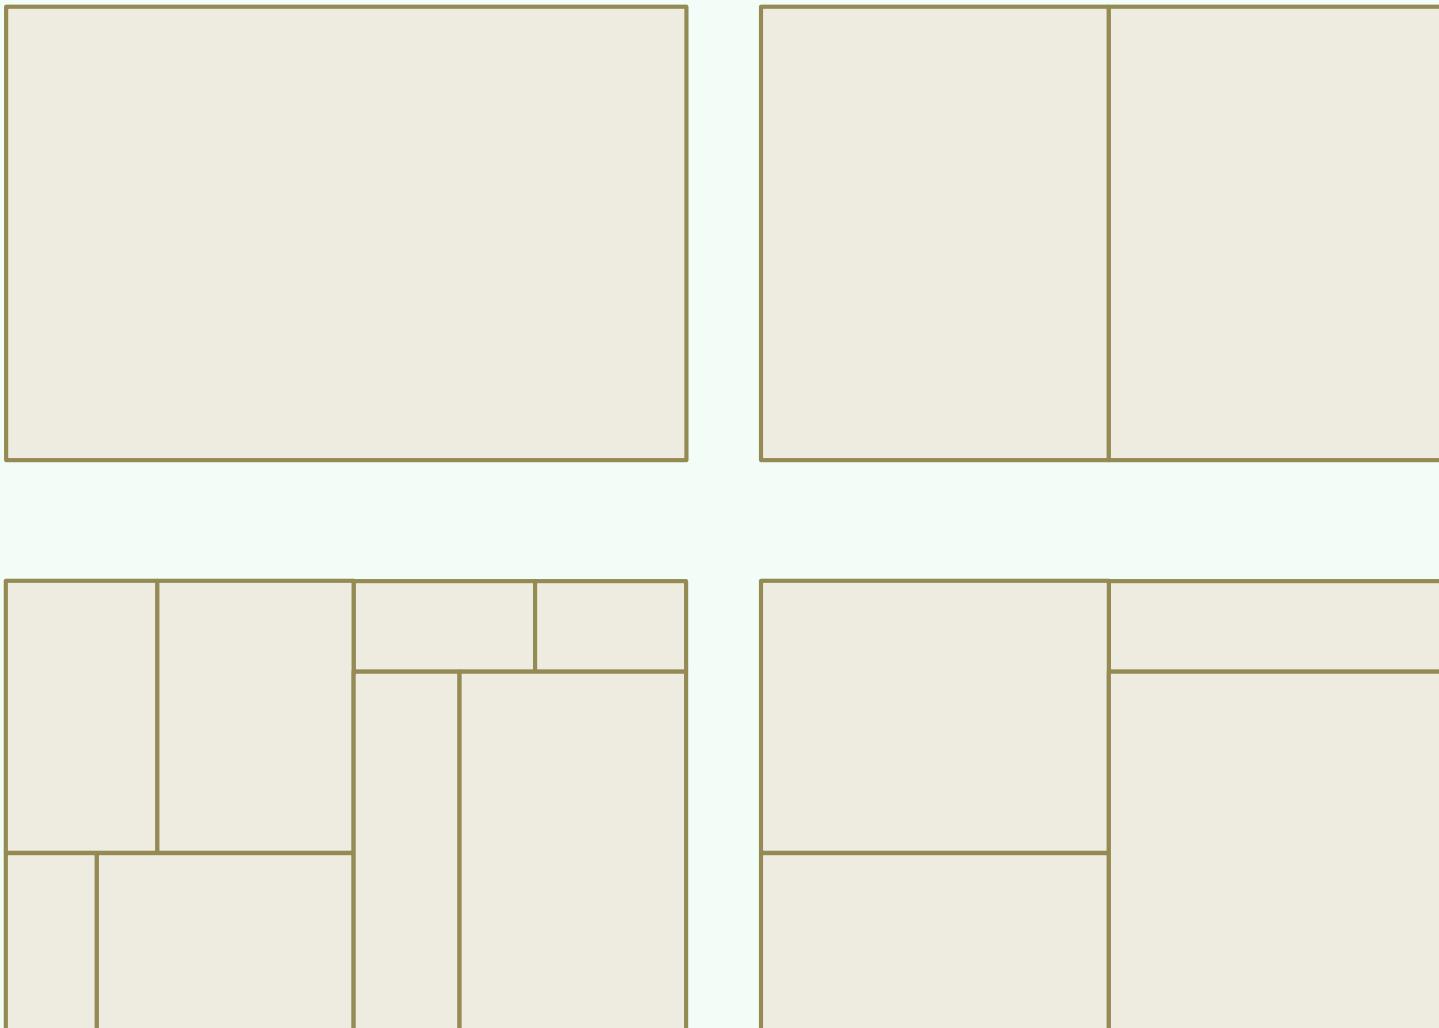
+ 2

+ 4

+ ...

+ $\Theta(2^{\log n})$

= $\Theta(n)$



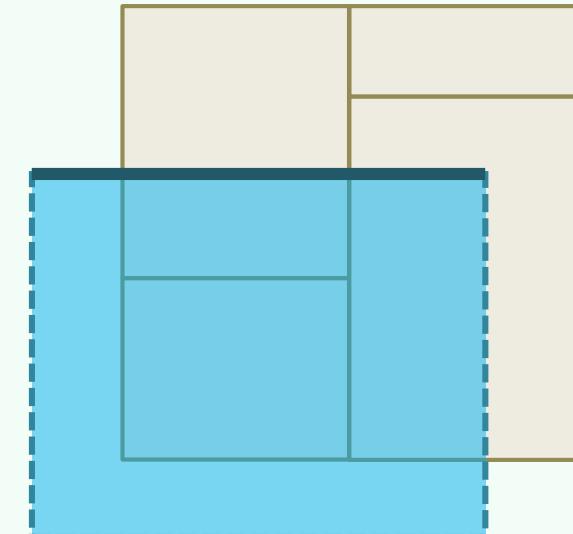
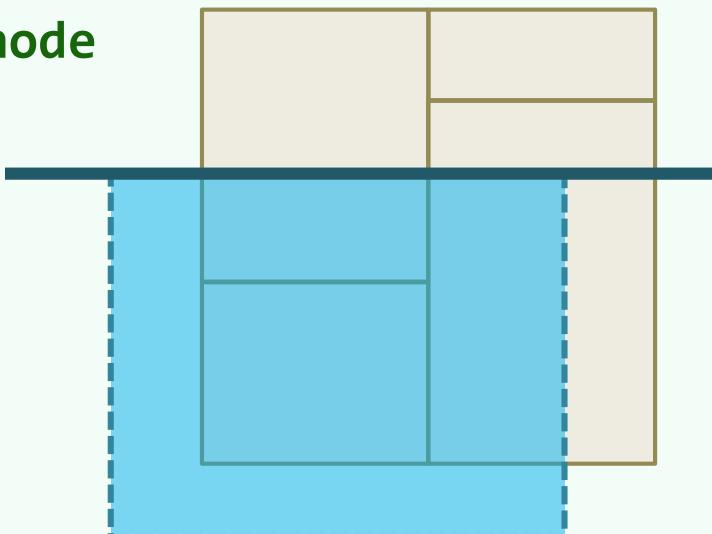
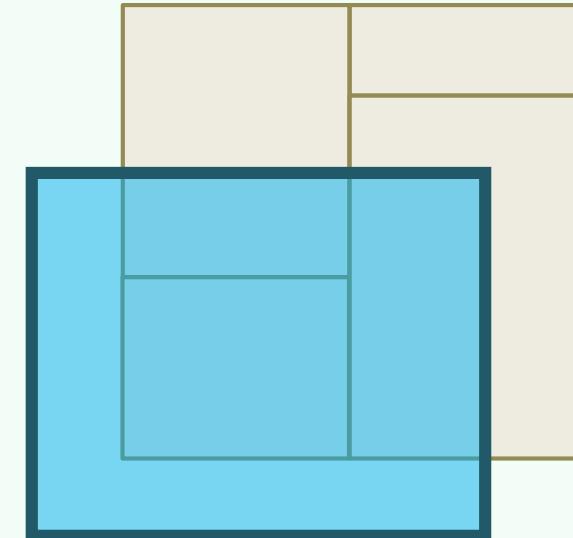
Query Time

- ❖ **Claim:** Report + Search = $\mathcal{O}(r + \sqrt{n})$
- ❖ The searching time depends on $Q(n)$, the number of
 - recursive calls, or
 - sub-regions **intersecting with R** (at all levels)

❖ No more than **2** of
the **4** grandchildren of each node

will recurse

- $Q(1) = \mathcal{O}(1)$
 - $Q(n) = 2 \cdot Q(n/4) + \mathcal{O}(1)$
- ❖ Solve to $Q(n) = \mathcal{O}(\sqrt{n})$



Beyond 2D

❖ Can 2d-tree be extended to kd-tree and help **HIGHER** dimensional GRS?

If yes, how efficiently can it help?

❖ A kd-tree in k-dimensional space is constructed by

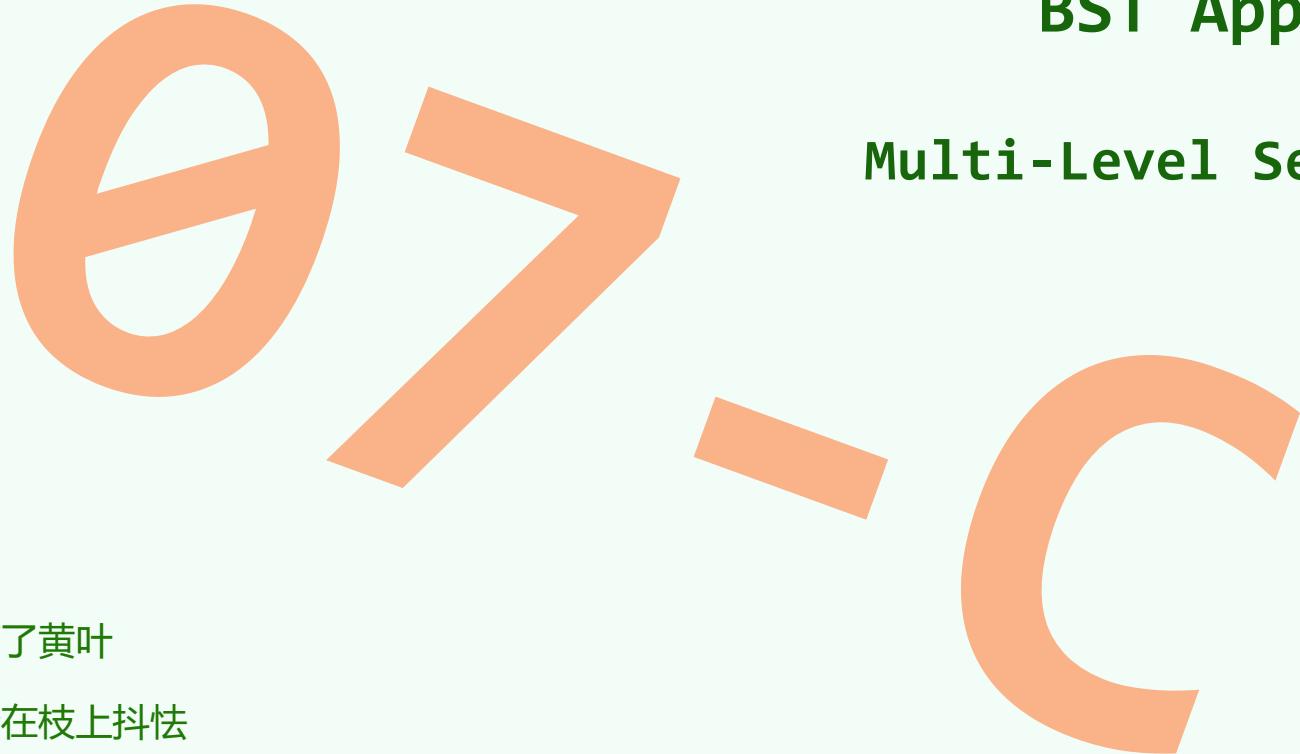
recursively divide \mathcal{E}^d along the $1^{\text{st}}, 2^{\text{nd}}, \dots, k^{\text{th}}$ dimensions

❖ An orthogonal range query on a set of n points in \mathcal{E}^d

- can be answered in $\mathcal{O}(r + n^{1-1/d})$ time,
- using a kd-tree of size $\mathcal{O}(n)$, which
- can be constructed in $\mathcal{O}(n \log n)$ time

BST Application

Multi-Level Search Tree



几株不知名的树，已脱下了黄叶

只有那两三片，多么可怜在枝上抖怯

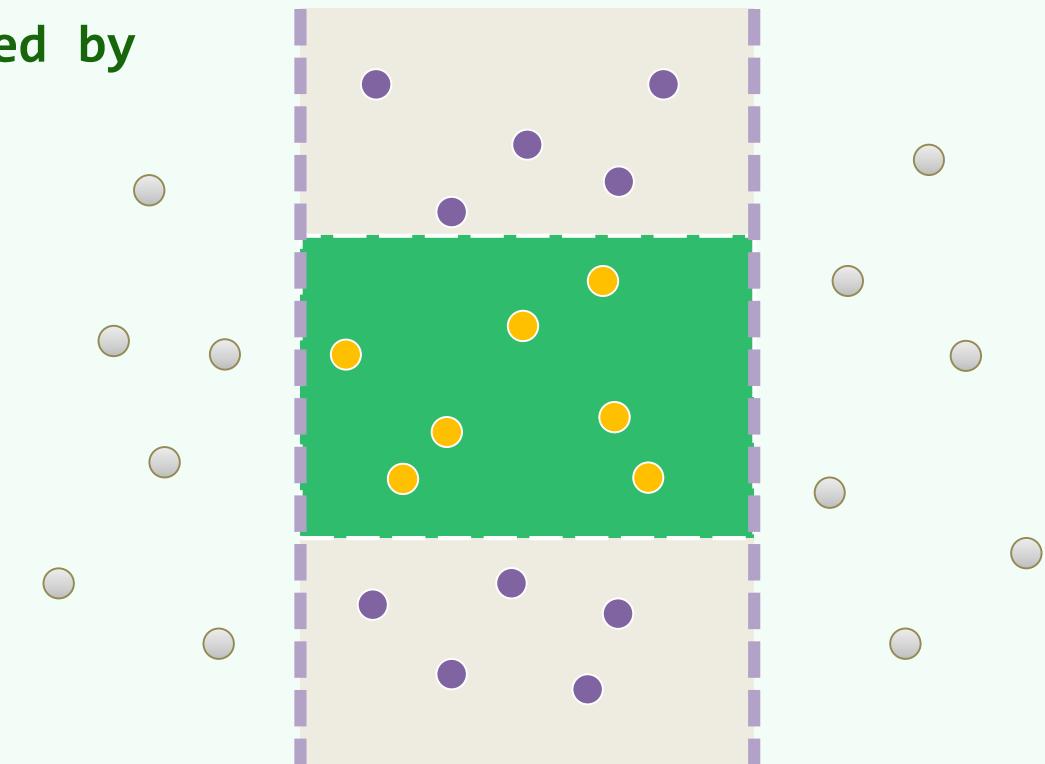
它们感到秋来到，要与世间离别

邓俊辉

deng@tsinghua.edu.cn

2D Range Query = x-Query + y-Query

- ❖ Is there any structure which answers range query **FASTER** than kd-trees?
- ❖ An m-D orthogonal range query can be answered by the **INTERSECTION** of m 1D queries
- ❖ For example, a 2D range query can be divided into two 1D range queries:
 - find all points in $[x_1, x_2]$; and then
 - find in these candidates those lying in $[y_1, y_2]$



Worst Cases

❖ Using kd-trees needs $\mathcal{O}(1 + \sqrt{n})$ time. But here ...

❖ The x-query returns

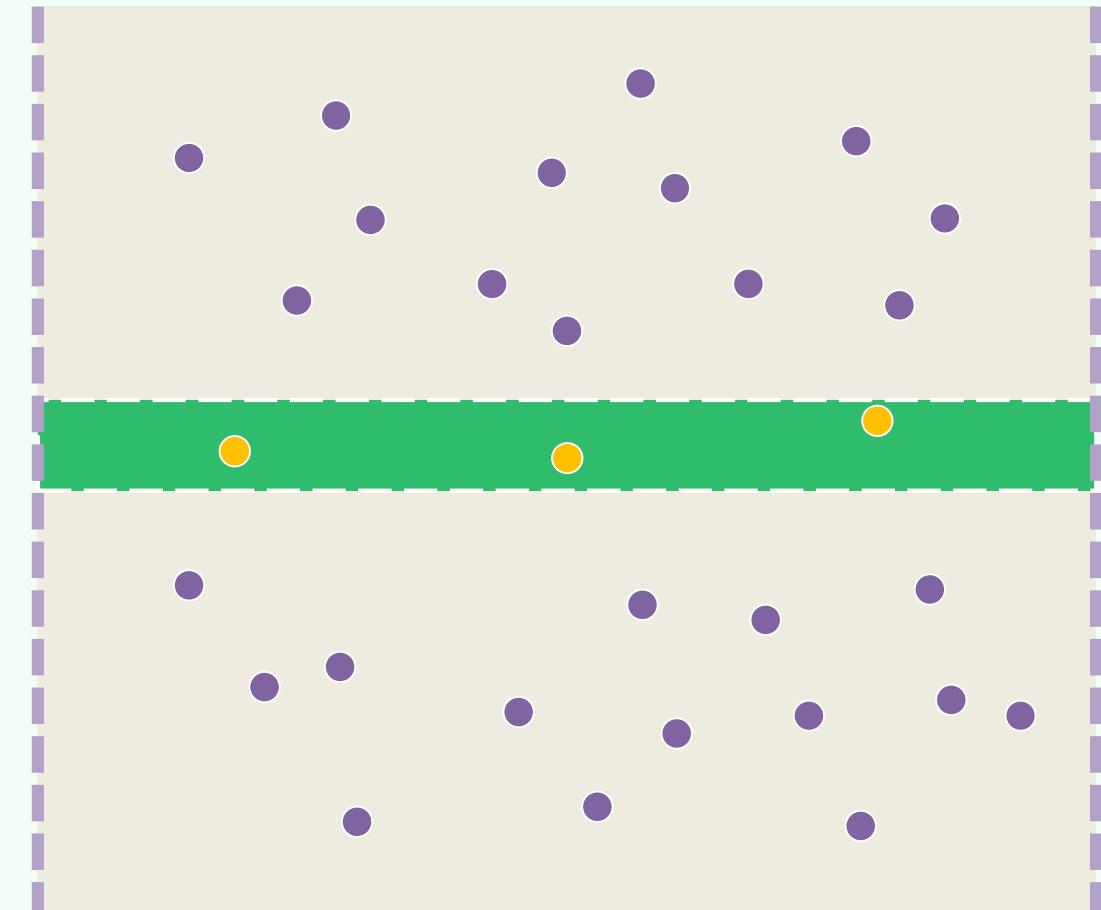
(almost) all points whereas

the y-query rejects

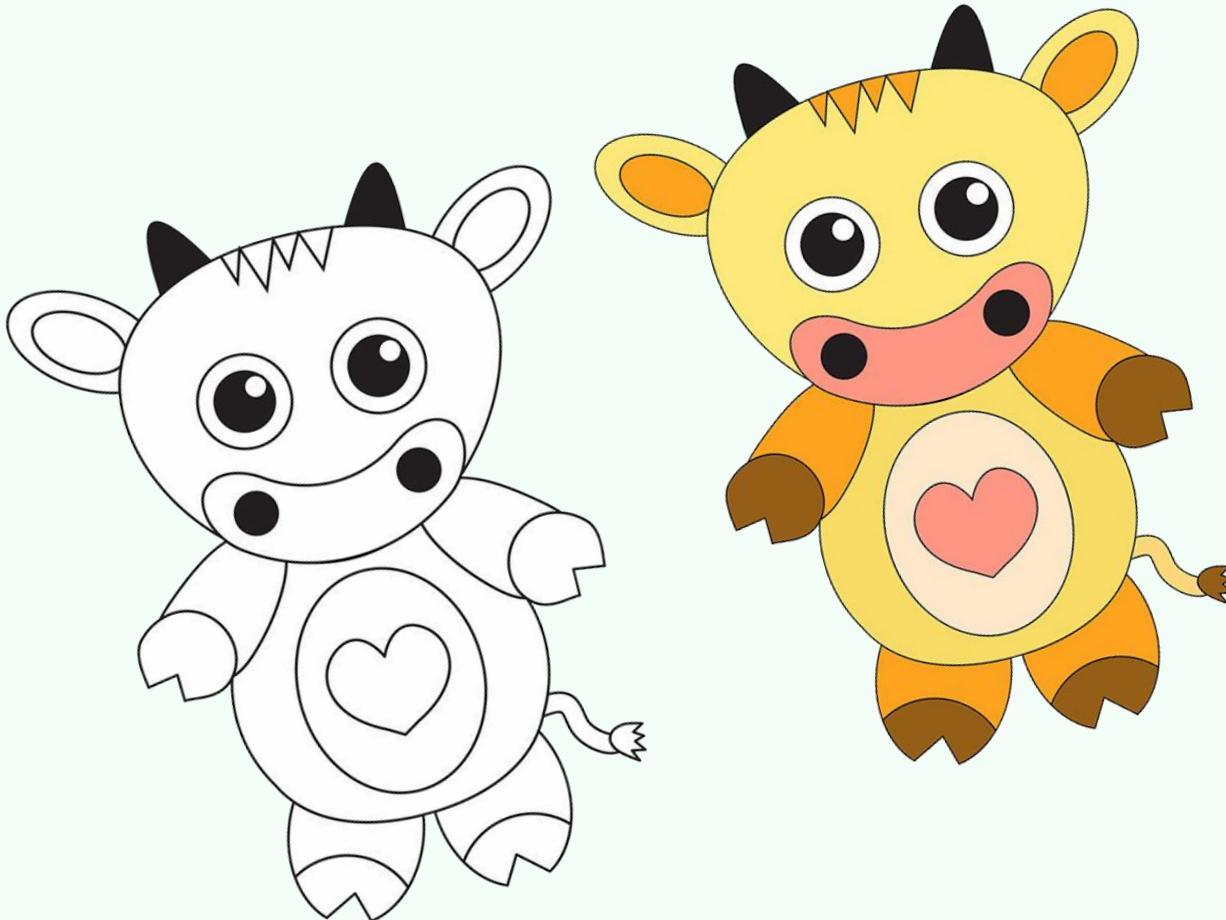
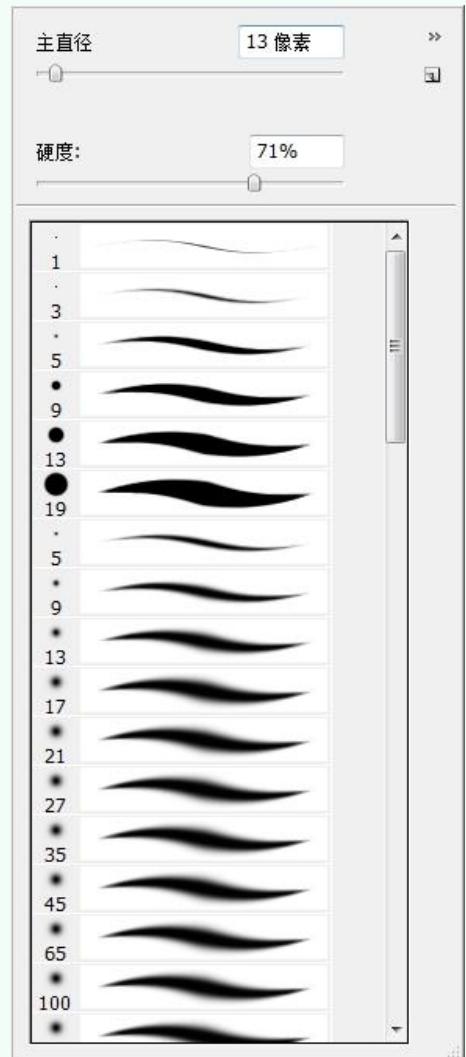
(almost) all

❖ We spent $\Omega(n)$ time

before getting $r = 0$ points



Painters' Strategy



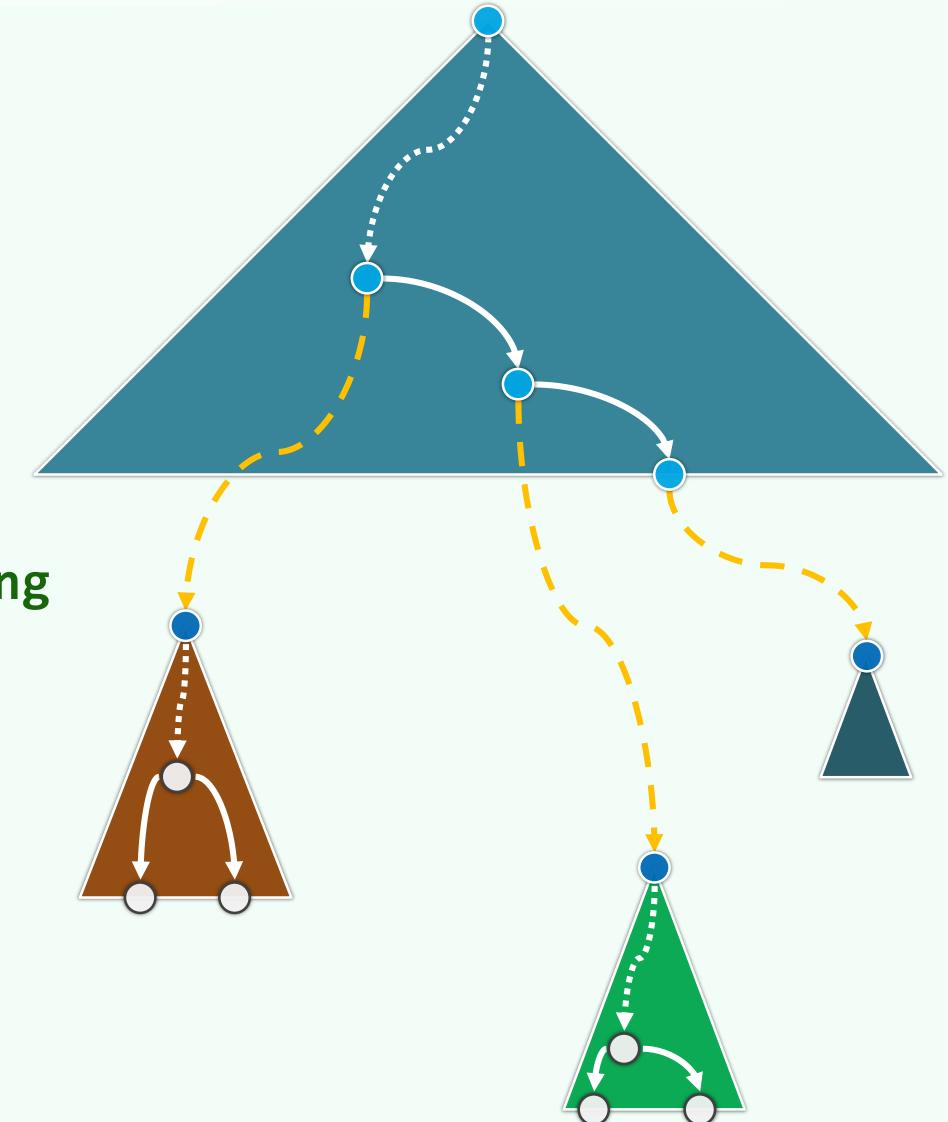
$$2D \text{ Range Query} = x\text{-Query} * y\text{-Query}$$

❖ Tree of trees

- build a 1D BBST (called **x-tree**)
for the first range query (**x-query**);
- for each node v in the x-range tree,
build a y-coordinate BBST (**y-tree**), containing
the canonical subset associate with v

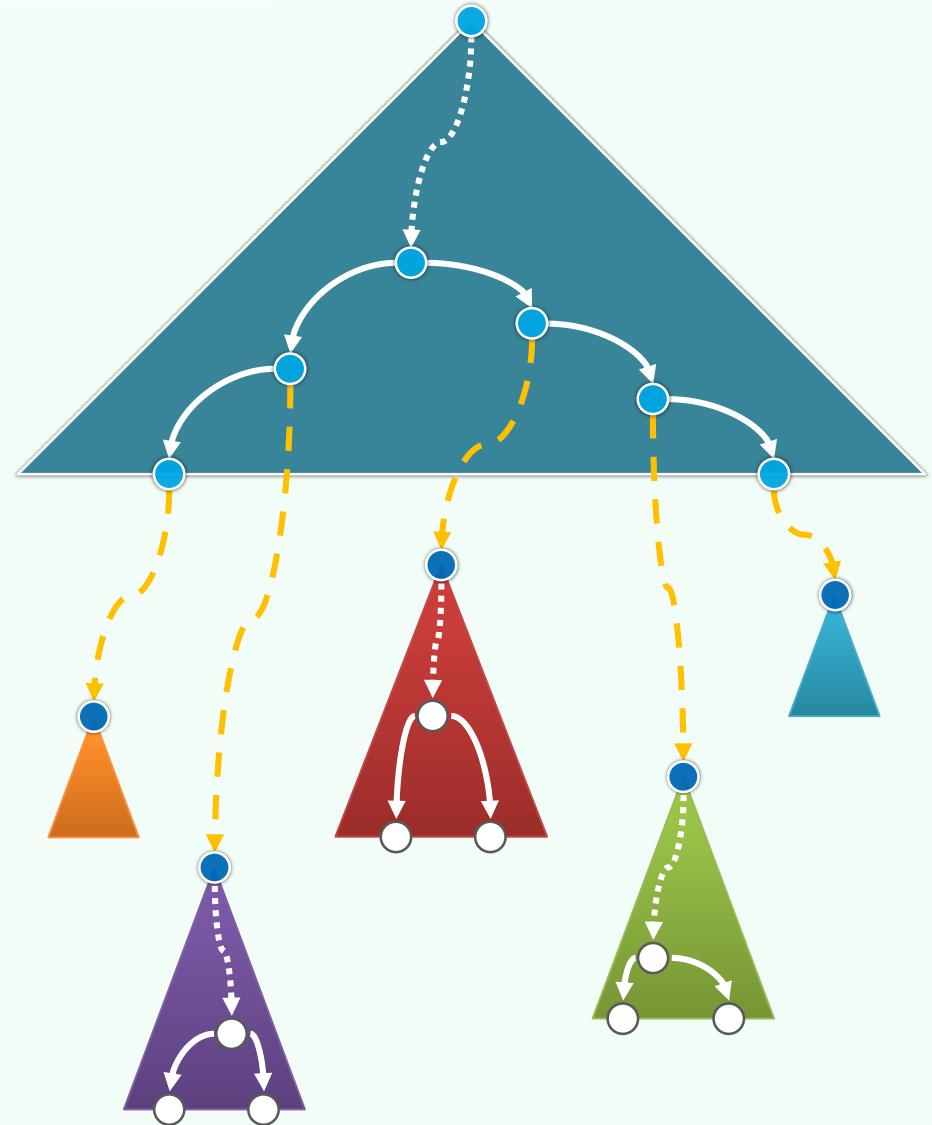
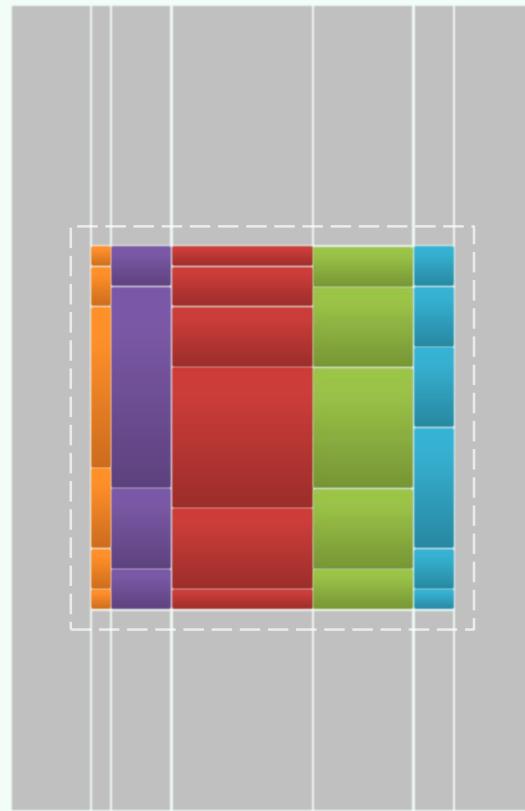
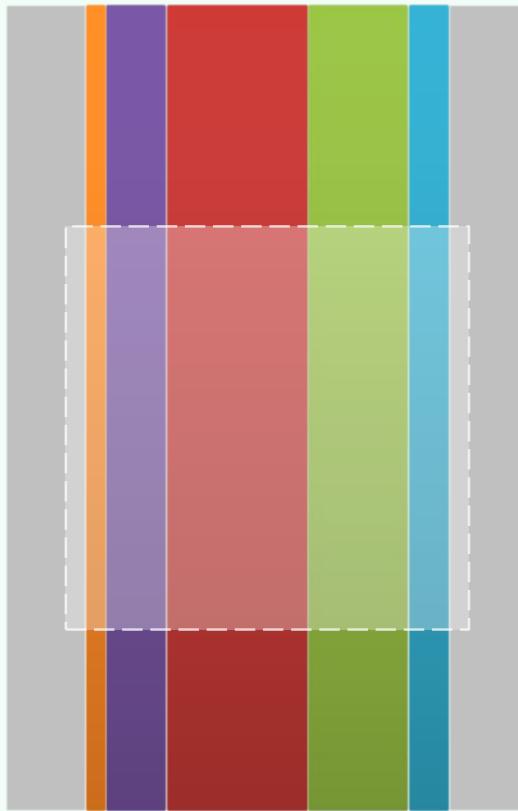
❖ It's an **x-tree** of (a number of) **y-trees**,
called a Multi-Level Search Tree

❖ How to answer range queries with such an MLST?



2D Range Query = x-Query * y-Queries

❖ **Query Time** = $\mathcal{O}(r + \log^2 n)$ ~ $\mathcal{O}(r + \log n)$



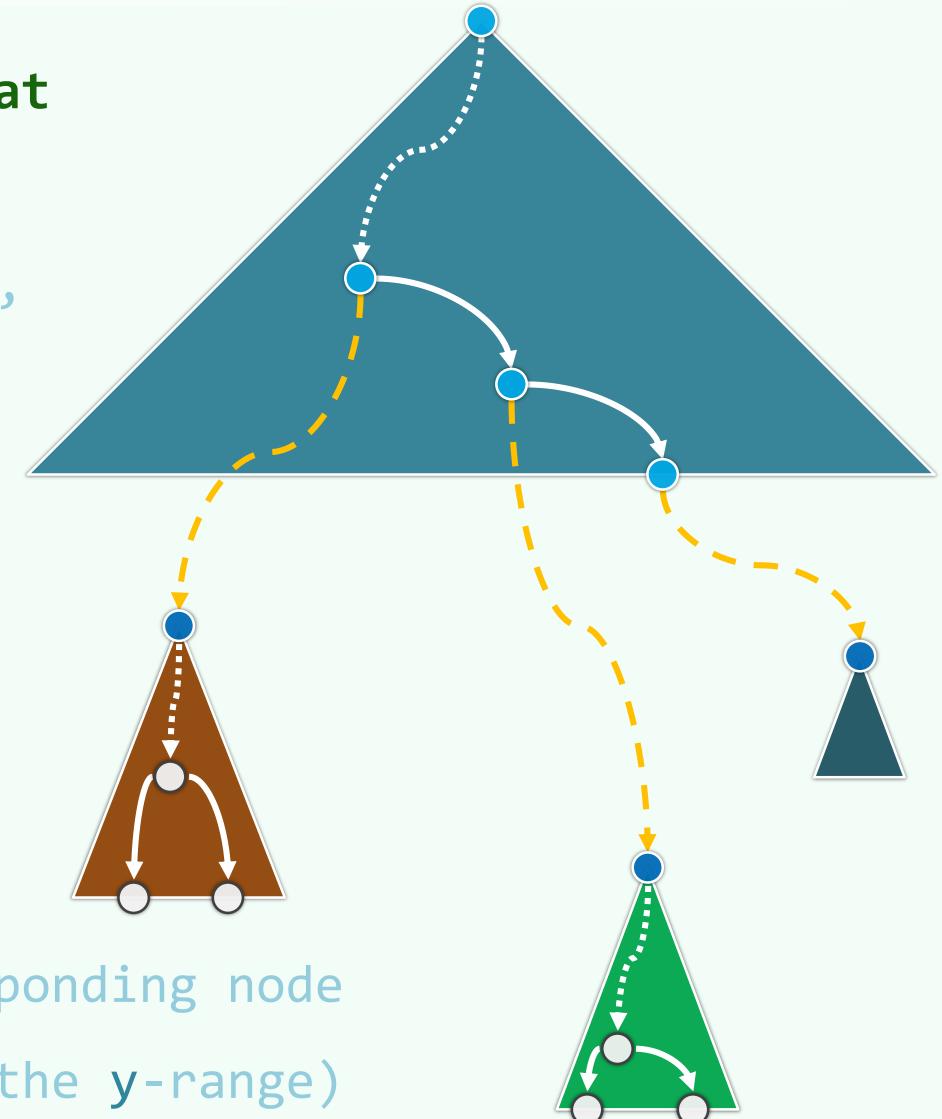
Query Algorithm

1. Determine the canonical subsets of points that satisfy the first query

```
// there will be  $\mathcal{O}(\log n)$  such canonical sets,  
// each of which is just represented as  
// a node in the x-tree
```

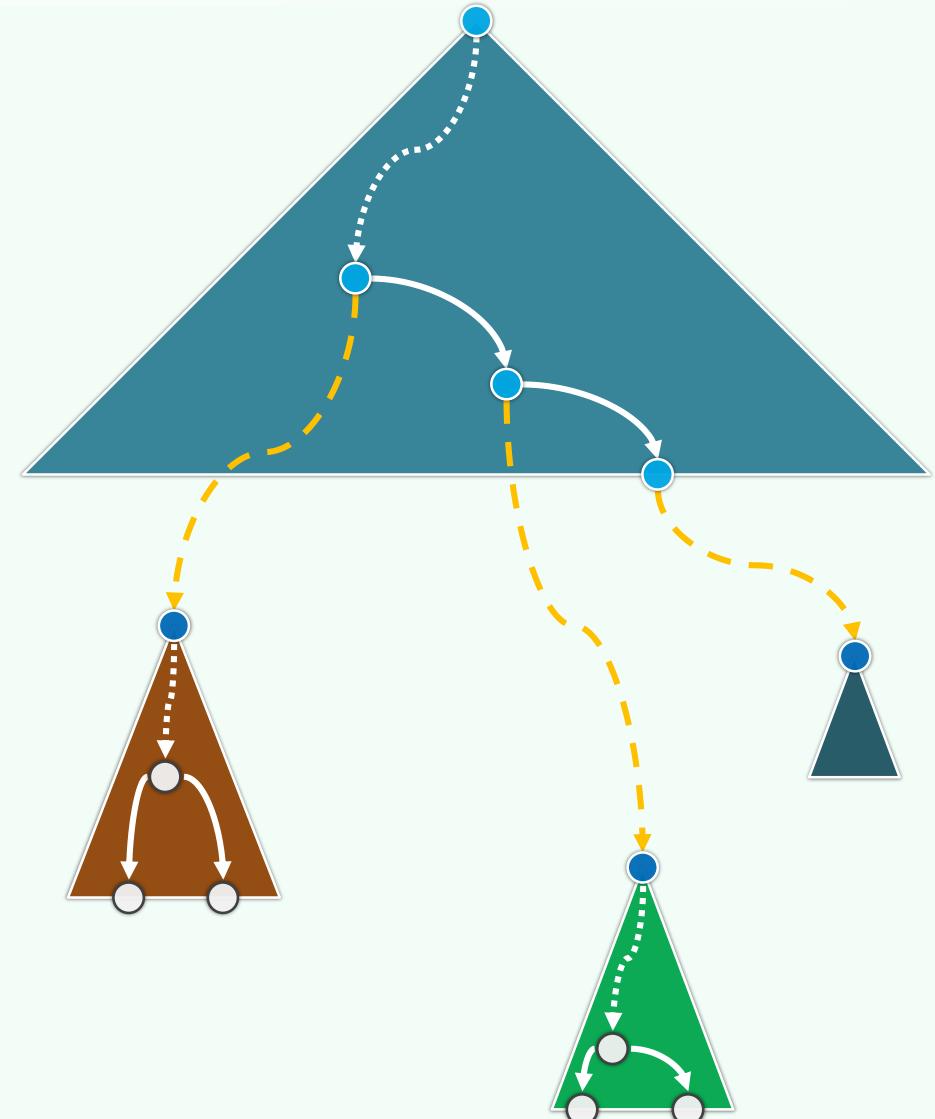
2. Find out from each canonical subset which points lie within the y-range

```
// To do this,  
// for each canonical subset,  
// we access the y-tree for the corresponding node  
// this will be again a 1D range search (on the y-range)
```



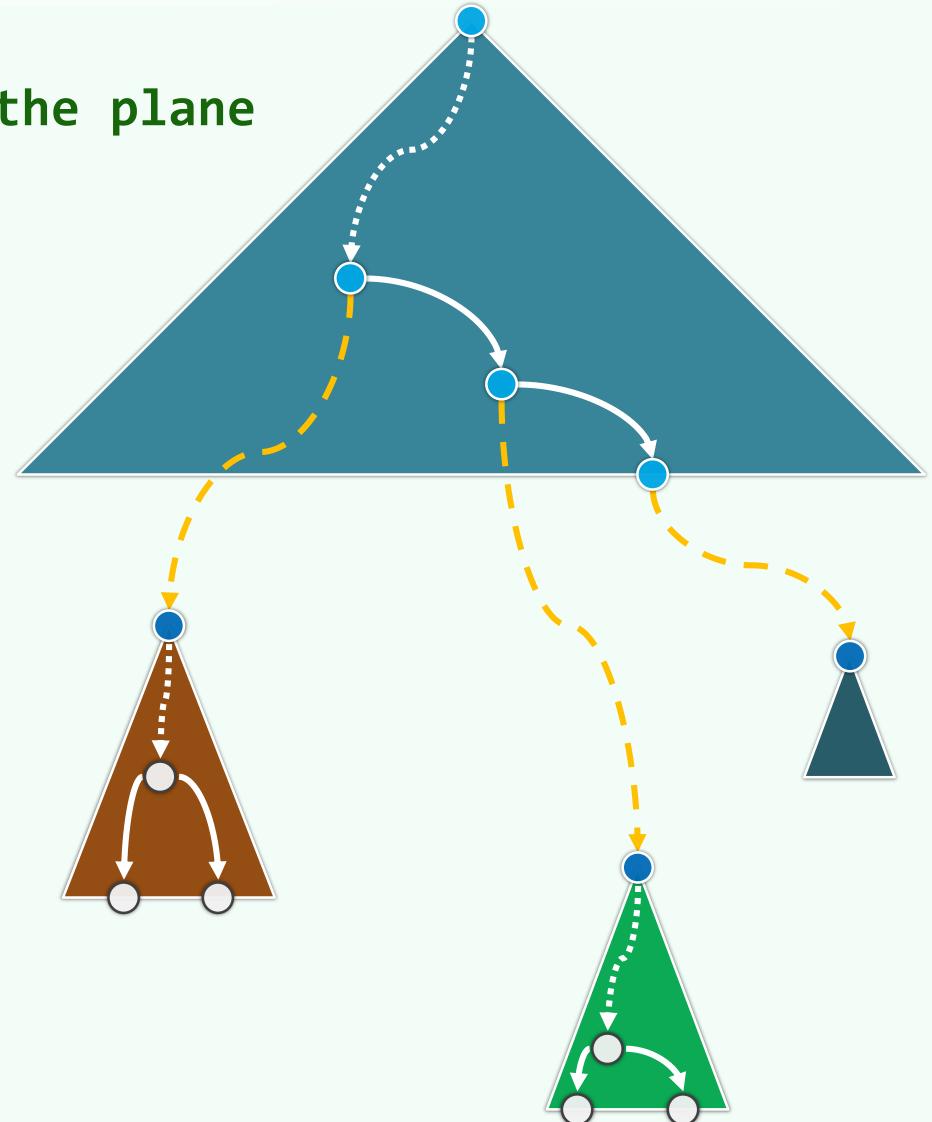
Complexity: Preprocessing Time + Storage

- ❖ A 2-level search tree
 - for n points in the plane
 - can be built
 - in $\mathcal{O}(n \log n)$ time
- ❖ Each input point is stored in $\mathcal{O}(\log n)$ y-trees
- ❖ A 2-level search tree
 - for n points in the plane
 - needs $\mathcal{O}(n \log n)$ space



Complexity: Query Time

- ❖ **Claim:** A 2-level search tree for n points in the plane answers each planar range query in $\mathcal{O}(r + \log^2 n)$ time
- ❖ The **x-range query** needs $\mathcal{O}(\log n)$ time to locate the $\mathcal{O}(\log n)$ nodes representing the canonical subsets
- ❖ Then for each of these nodes, a **y-range search** is invoked, which needs $\mathcal{O}(\log n)$ time

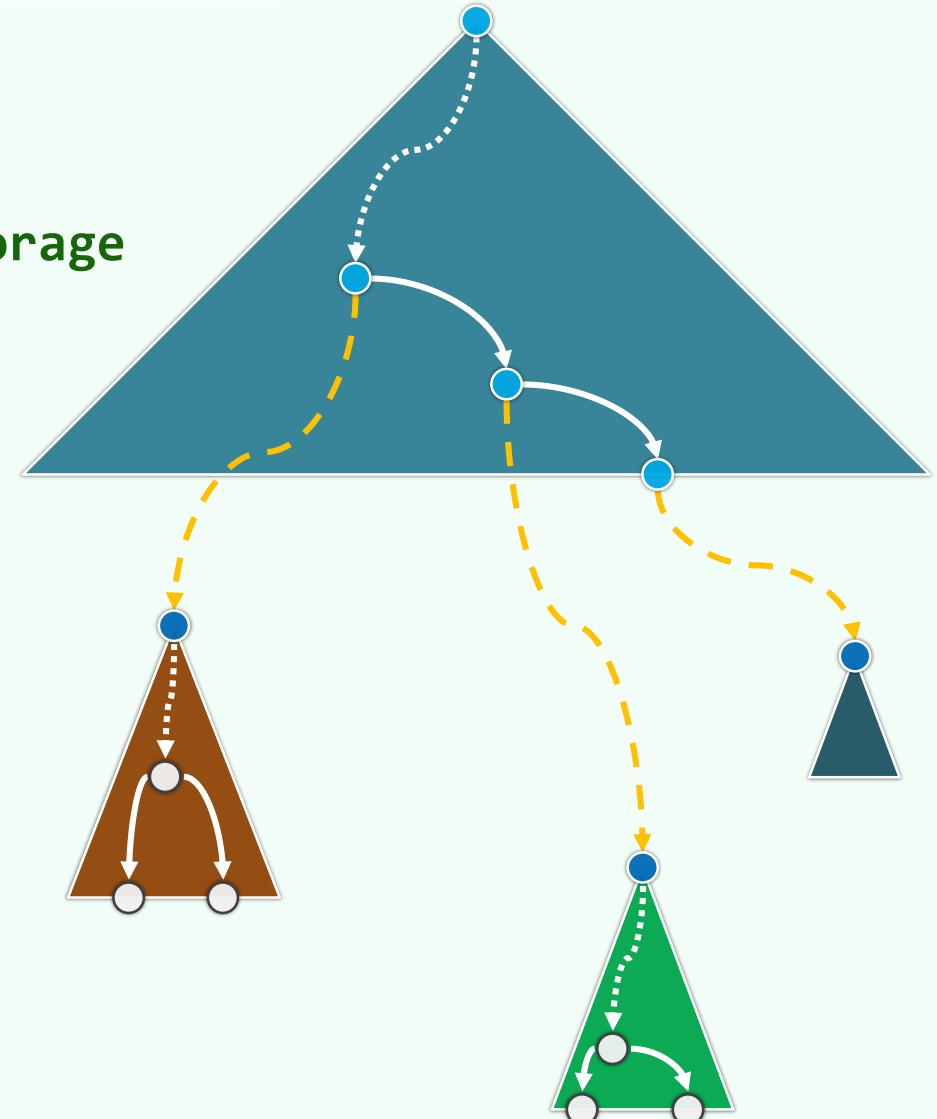


Beyond 2D

- ❖ Let S be a set of n points in \mathcal{E}^d , $d \geq 2$

- A d -level tree for S uses $\mathcal{O}(n \cdot \log^{d-1} n)$ storage
- Such a tree can be constructed in $\mathcal{O}(n \cdot \log^{d-1} n)$ time
- Each orthogonal range query of S can be answered in $\mathcal{O}(r + \log^d n)$ time

- ❖ For planar case, can the query time be improved to, say, $\mathcal{O}(\log n)$?



BST Application

Range Tree

e > -D

邓俊辉

deng@tsinghua.edu.cn

顺藤摸瓜

Coherence

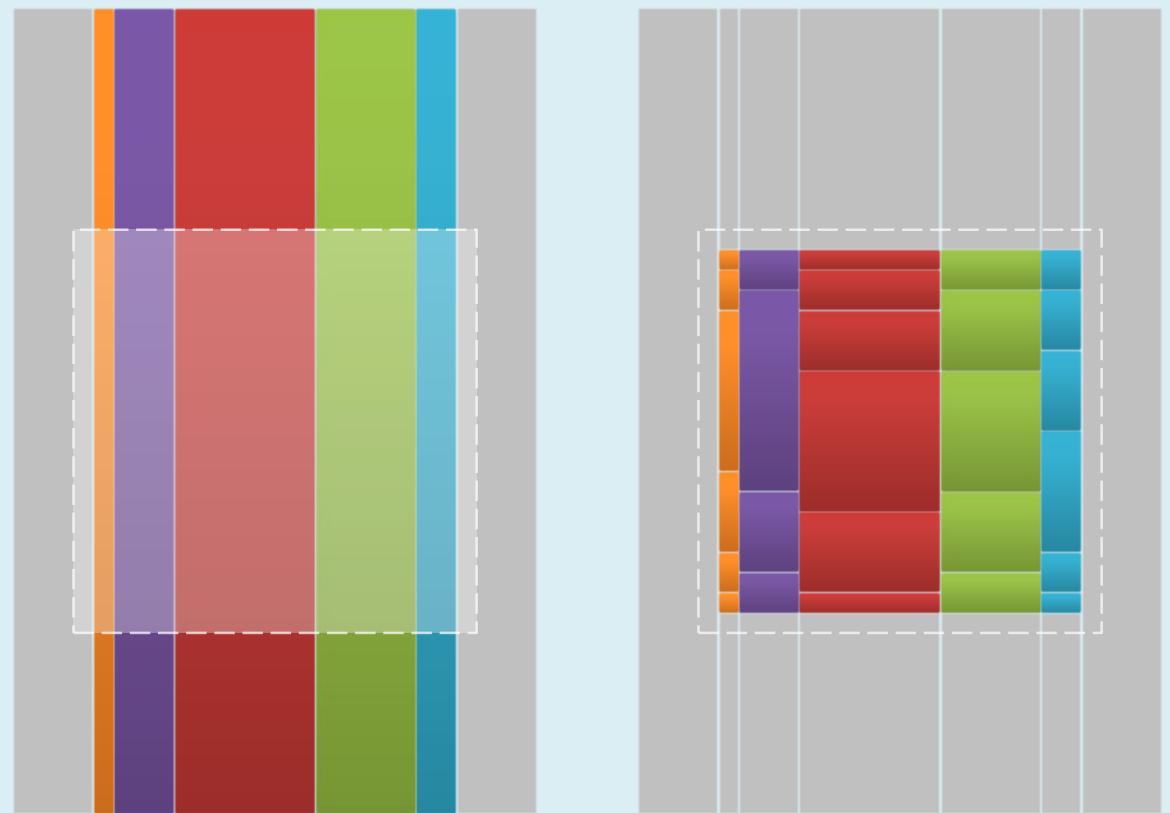
❖ For each query, we

- need to repeatedly search

DIFFERENT y-lists,

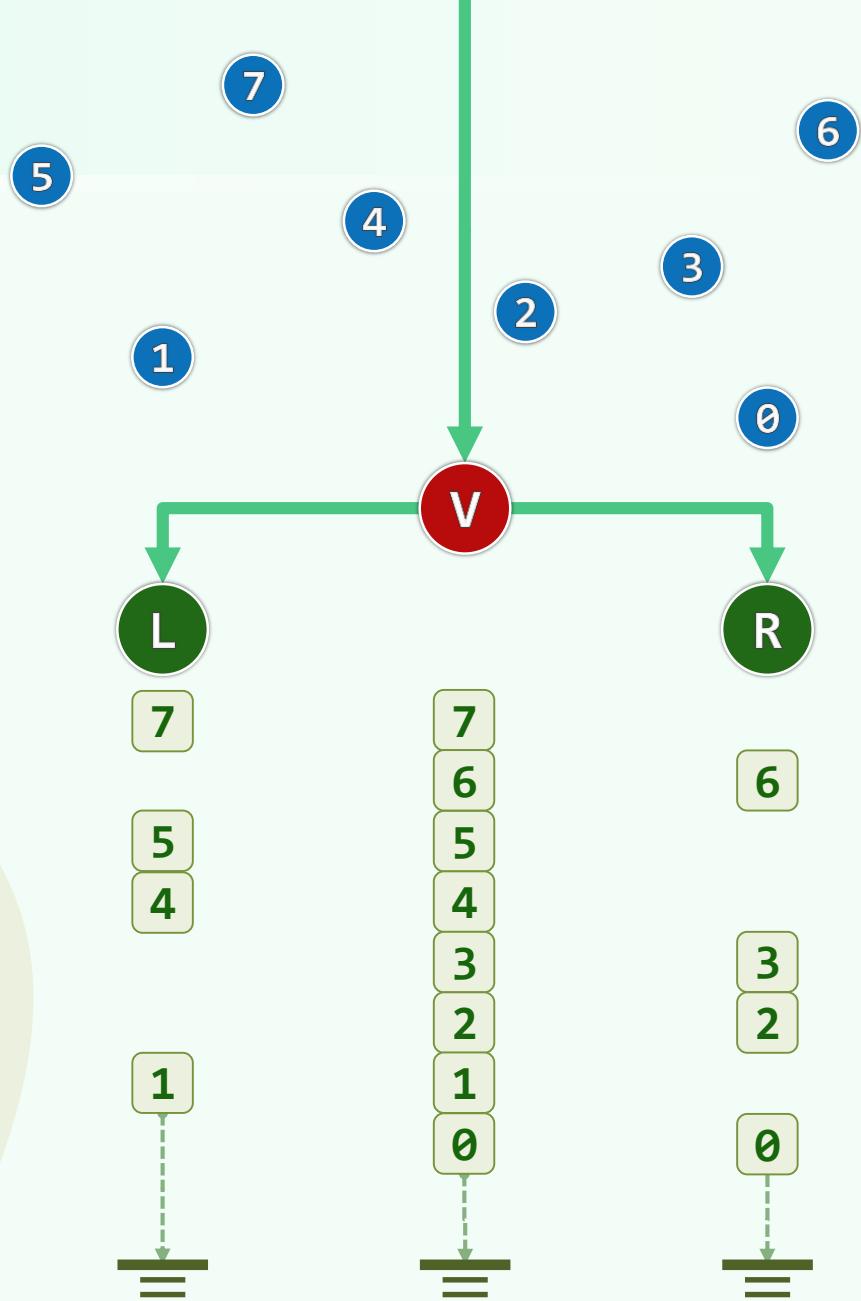
- but always with

the SAME key



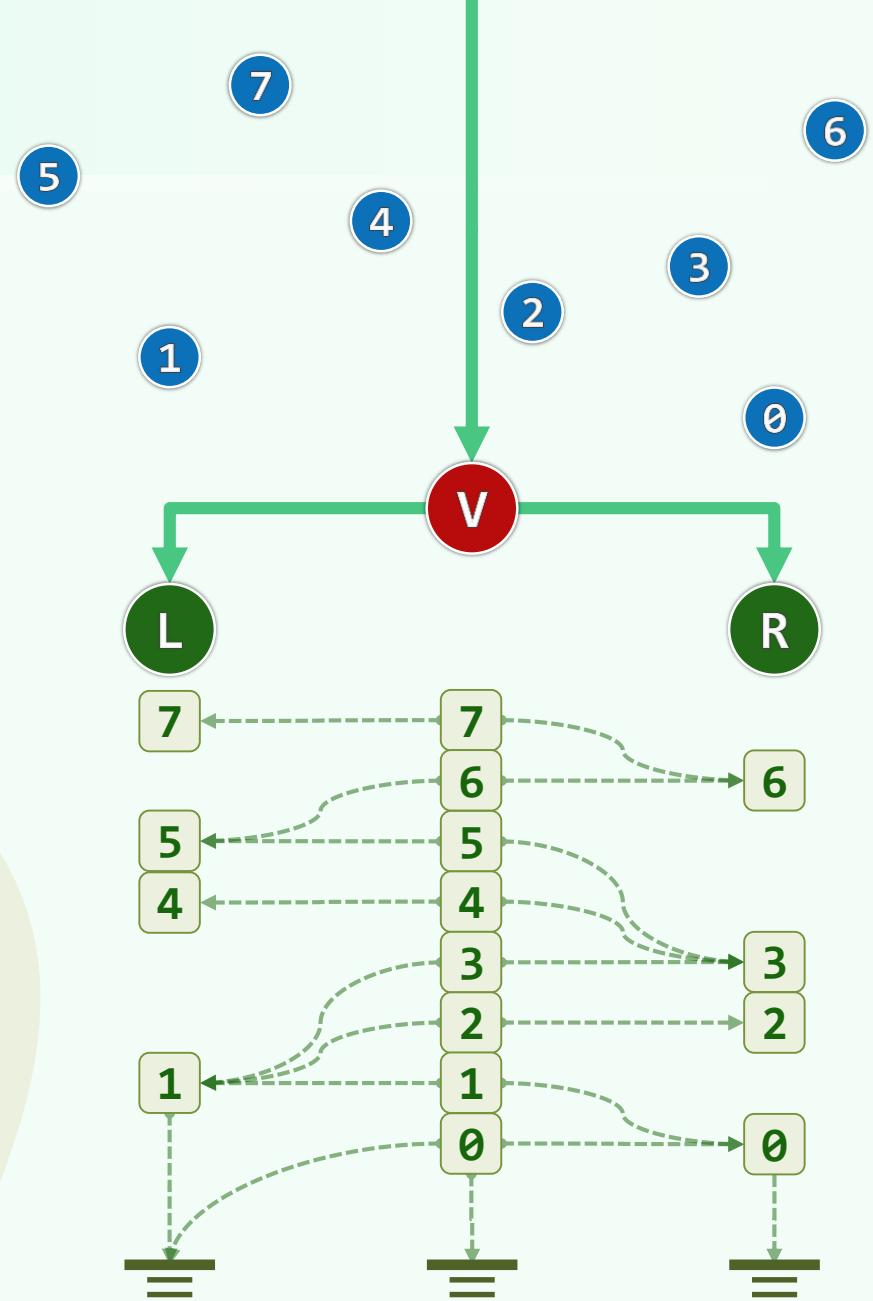
BBST<BBST<T>> --> **BBST<List<T>>**

- ❖ Each y-search is just a **1D query without further recursions**
- ❖ So it not necessary to store each canonical subset as a BBST
- ❖ Instead, a sorted y-list simply works



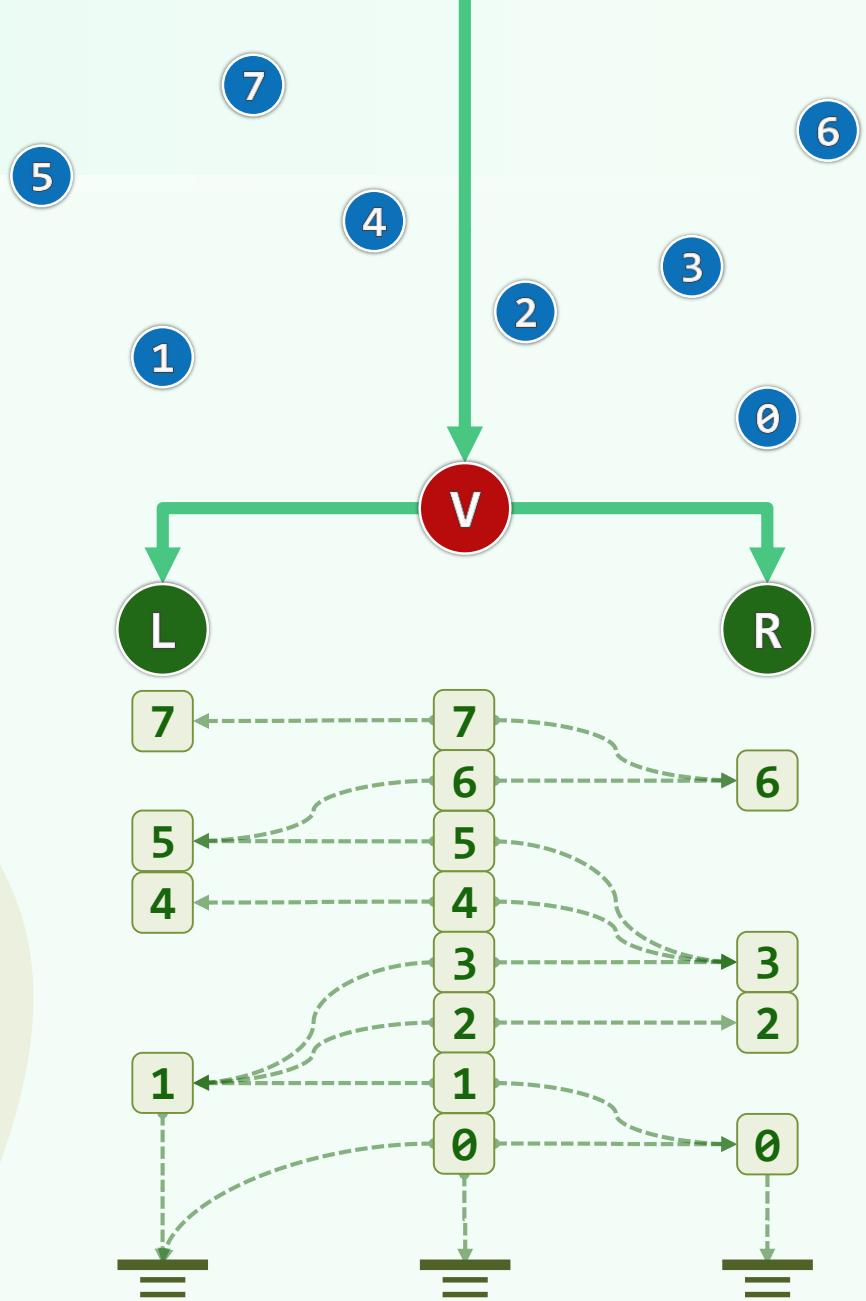
Links Between Lists

- ❖ We can **CONNECT** all the different lists into a **SINGLE** massive list
- ❖ Thus, once a parent y-list is searched, we can get, in $\mathcal{O}(1)$ time, the entry for child y-list by following the link between them



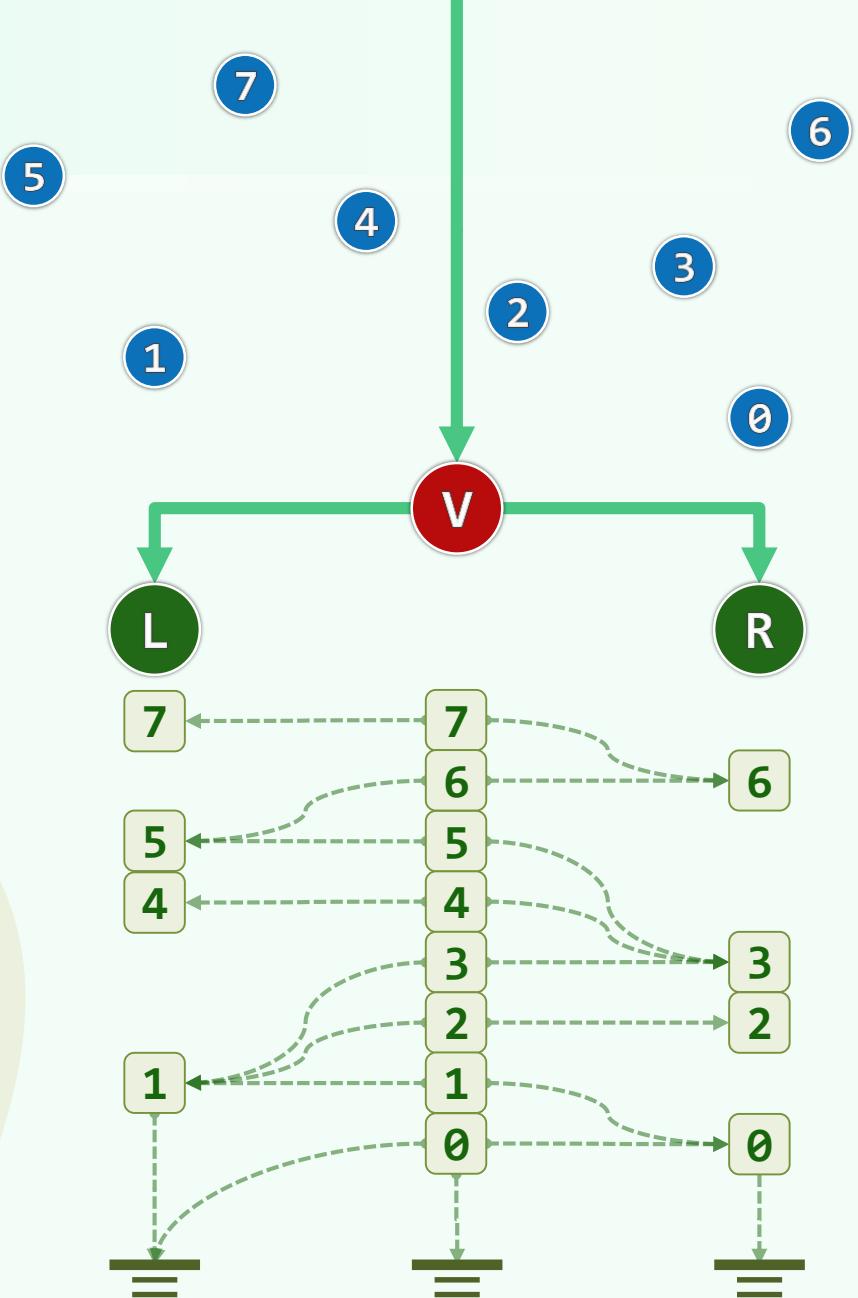
Using Coherence

- ❖ To answer a 2D range query, we will do an $\mathcal{O}(\log n)$ search on the **y-list** for the **LCA**
- ❖ Thereafter, while descending the **x-tree**, we can keep track of the position of **y-range** in each successive list in $\mathcal{O}(1)$ time
- ❖ This technique is called



Fractional Cascading

- ❖ For each item in A_v ,
we store two pointers to
the item of **NLT** value
in A_L and A_R resp.
- ❖ Hence for any y-query with q_y ,
once we know its entry in A_v , we can
determine its entry in either A_L or A_R
in $\Theta(1)$ additional time



Construction By 2-Way Merging

❖ Let v be an internal node in the x -tree

with L/R its left/right child resp.

❖ Let A_v be the y -list for v and

A_L/A_R be the y -lists for its children

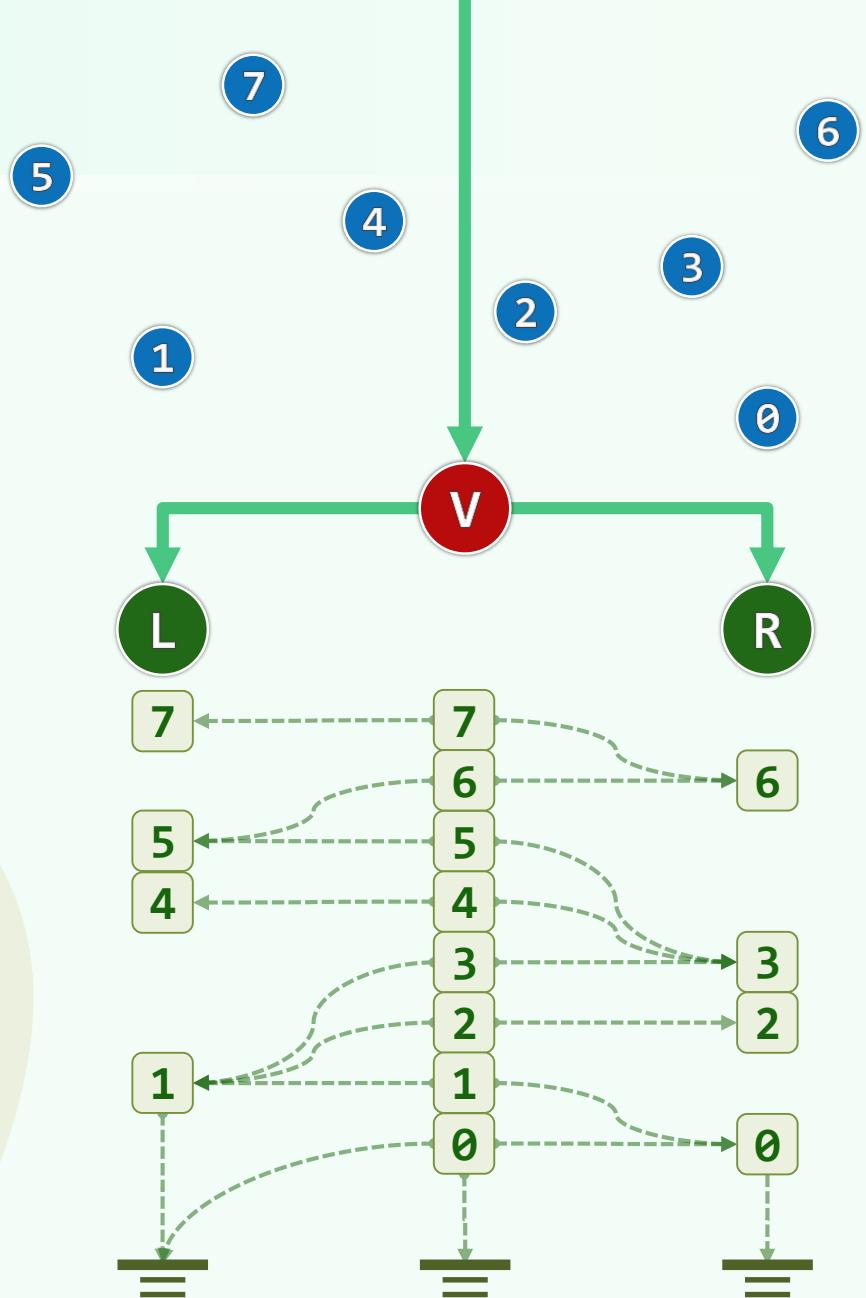
❖ Assuming no duplicate y -coordinates, we have

- A_v is the disjoint union

- of A_L and A_R , and hence

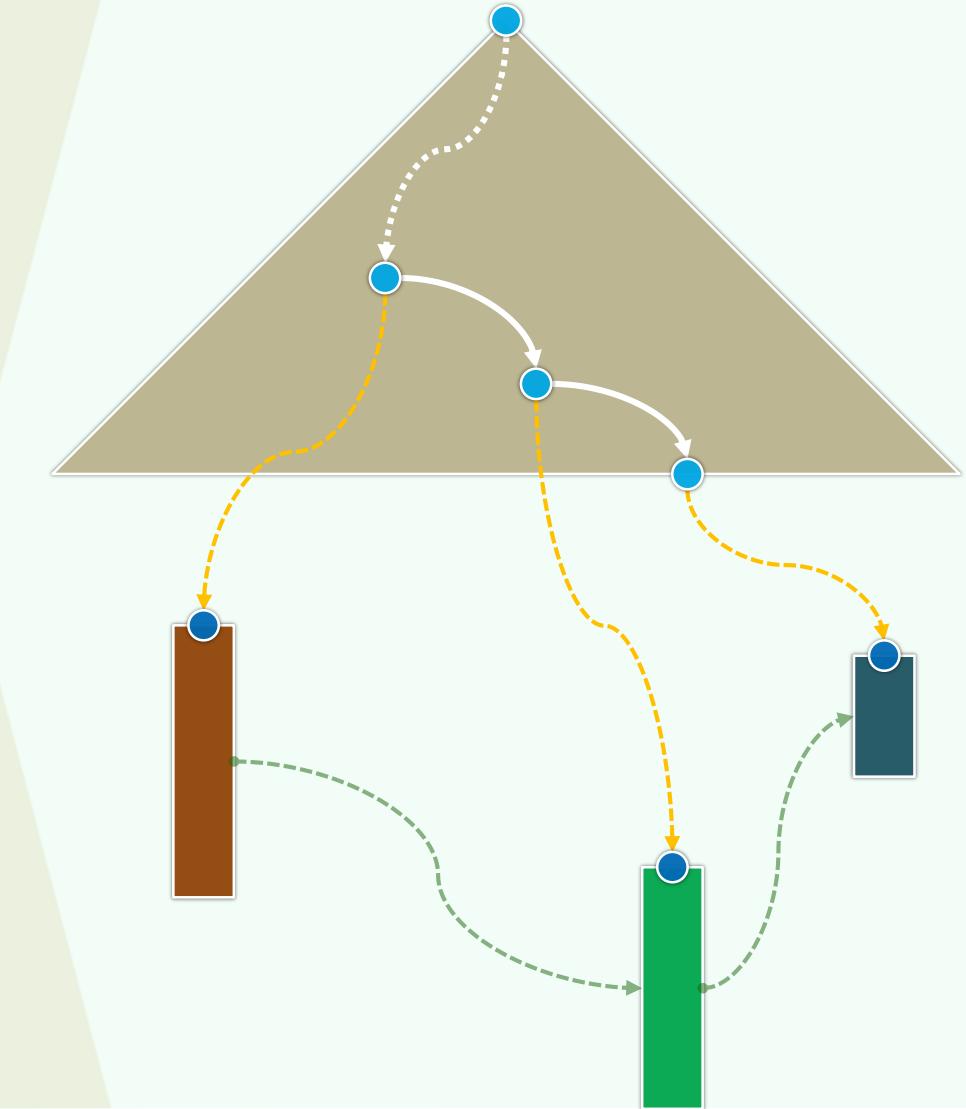
- A_v can be obtained by

- merging A_L and A_R (in linear time)



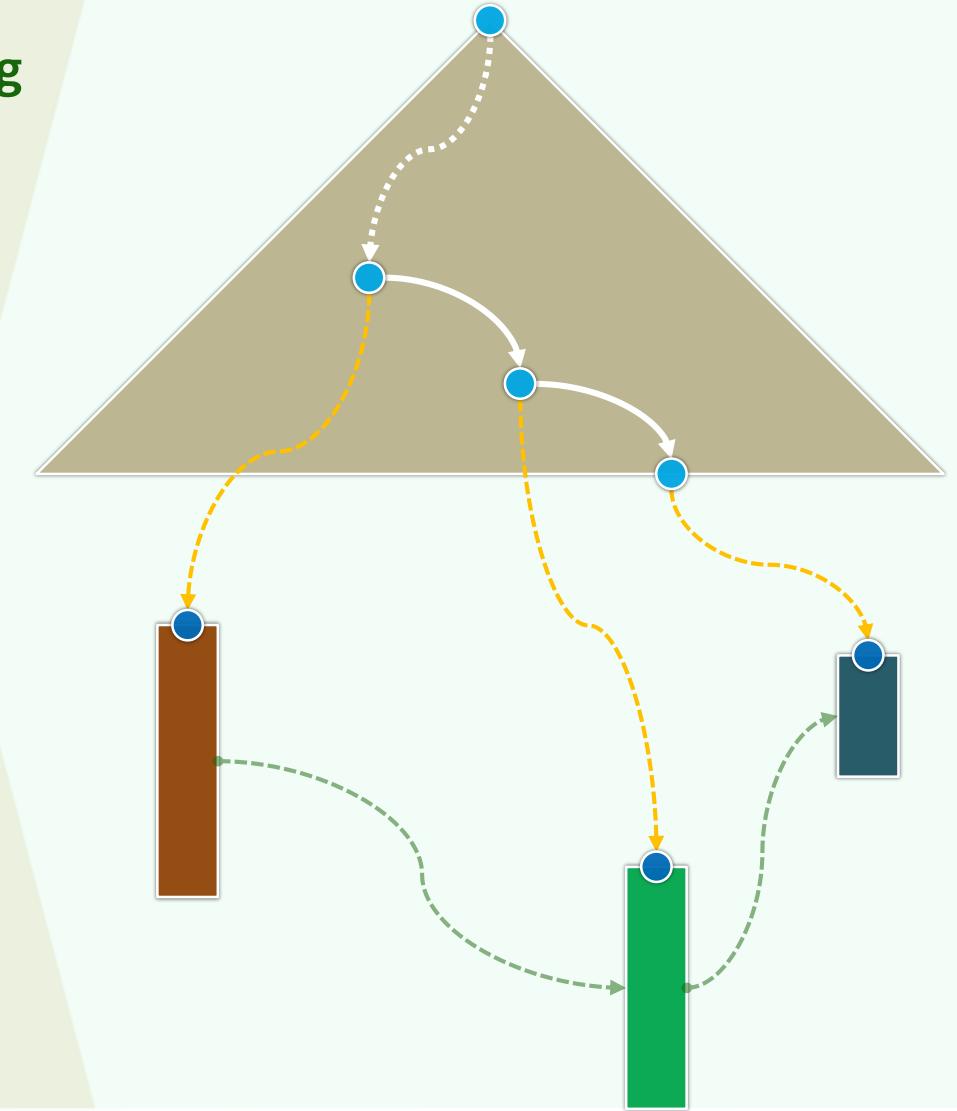
Complexity

- ❖ An MLST with fractional cascading is called a **range tree**
- ❖ A **y-search** for root is done in $\mathcal{O}(\log n)$ time
- ❖ To drop down to each next level, we can get, in $\mathcal{O}(1)$ time, the current **y-interval** from that of the **prior level**
- ❖ Hence, each 2D orthogonal range query
 - can be answered in $\mathcal{O}(r + \log n)$ time
 - from a data structure of size $\mathcal{O}(n \cdot \log n)$,
 - which can be constructed in $\mathcal{O}(n \cdot \log n)$ time



Beyond 2D

- ❖ Unfortunately, the trick of fractional cascading can **ONLY** be applied to the **LAST** level the search structure
- ❖ Given a set of n points in \mathcal{E}^d , an orthogonal range query
 - can be answered in $\mathcal{O}(r + \log^{d-1} n)$ time
 - from a data structure of size $\mathcal{O}(n \cdot \log^{d-1} n)$,
 - which can be constructed in $\mathcal{O}(n \cdot \log^{d-1} n)$ time



BST Application

Interval Tree



Your instinct, rather than precision stabbing, is more
about just random bludgeoning.

邓俊辉

deng@tsinghua.edu.cn

Stabbing Query

❖ Given a set of intervals in general position

on the **x-axis**:

$$S = \{s_i = [x_i, x'_i] \mid 1 \leq i \leq n\}$$

and a query point q_x

❖ Find all intervals that contain q_x

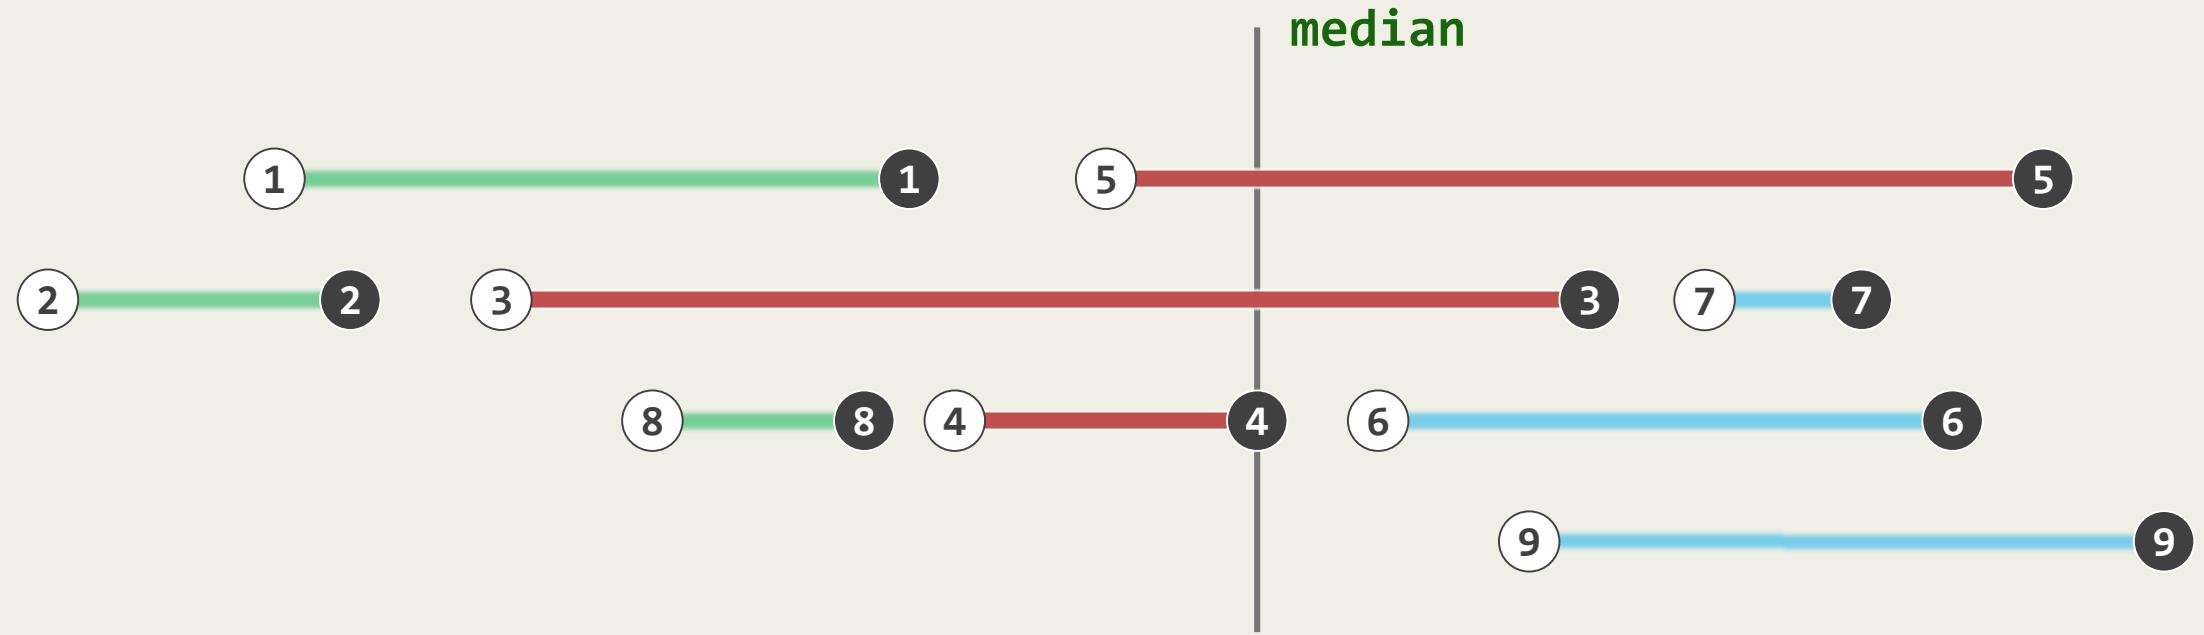
$$\{s_i = [x_i, x'_i] \mid x_i \leq q_x \leq x'_i\}$$

❖ To solve this query,

we will use the so-called interval tree ...

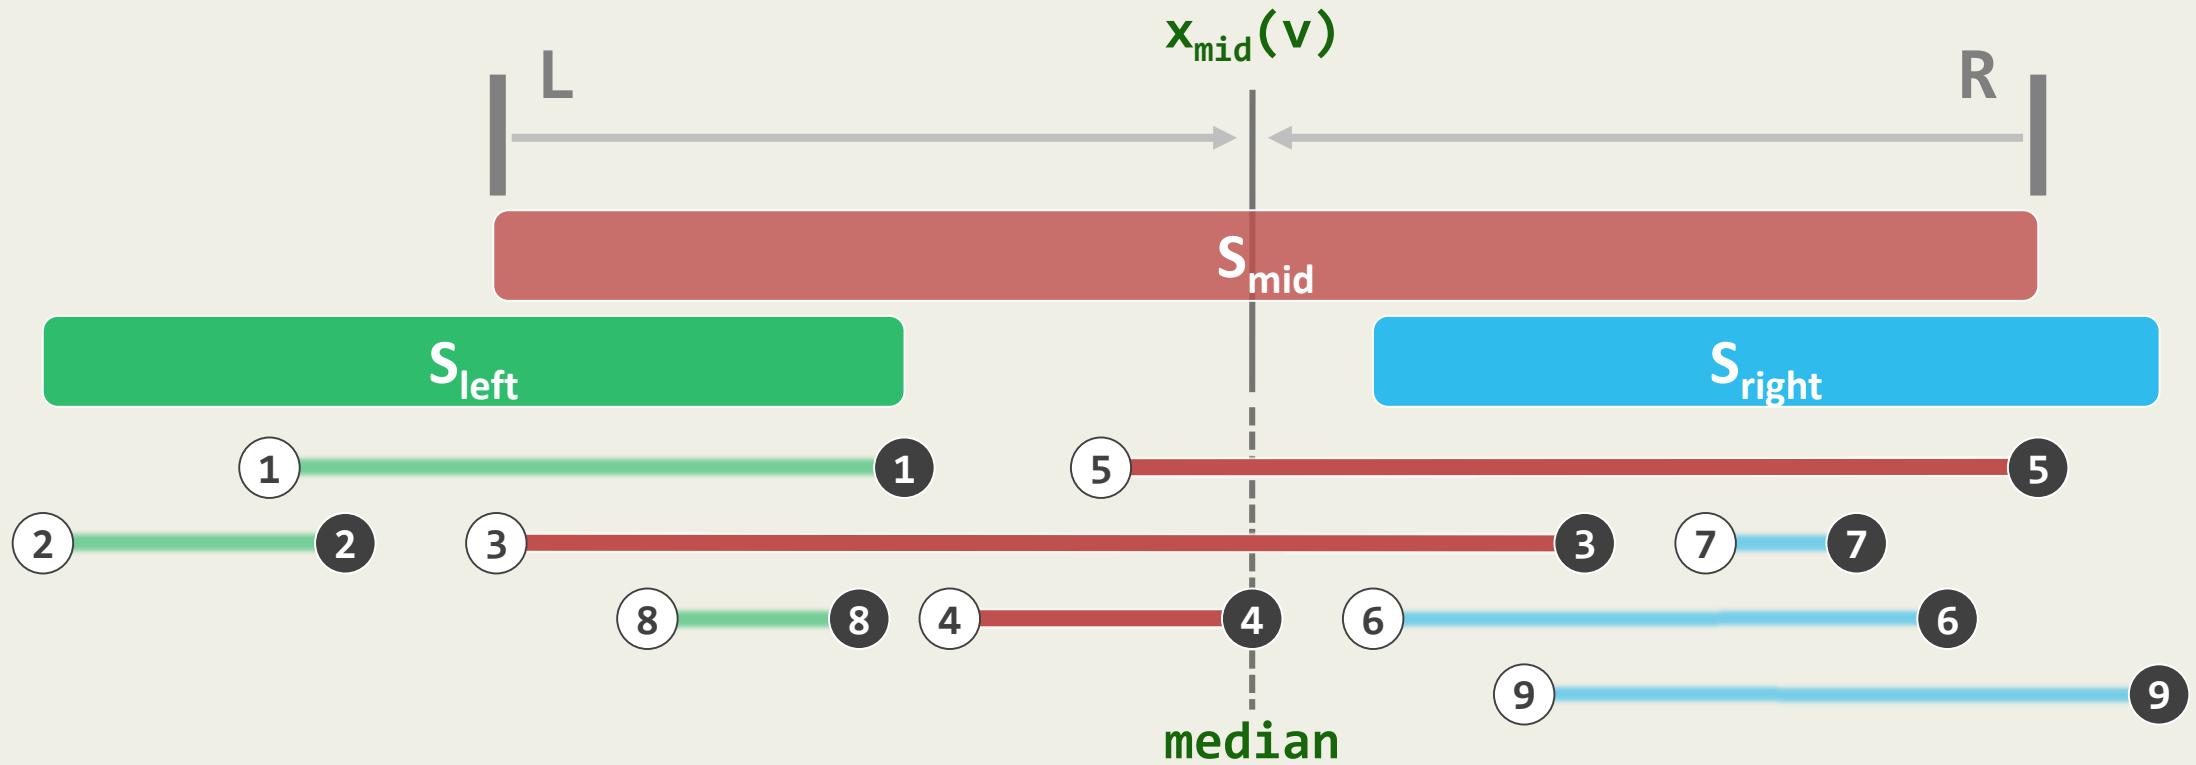


Median



- ❖ Let $P = \partial S$ be the set of all endpoints
 - (By general position assumption, $|P| = 2n$)
- ❖ Let $x_{mid} = \text{median}(P)$ be the median of P

Partitioning



❖ All intervals can be then categorized into 3 subsets:

$$S_{left} = \{ S_i \mid x'_i < x_{mid} \} \quad S_{mid} = \{ S_i \mid x_i \leq x_{mid} \leq x'_i \} \quad S_{right} = \{ S_i \mid x_{mid} < x_i \}$$

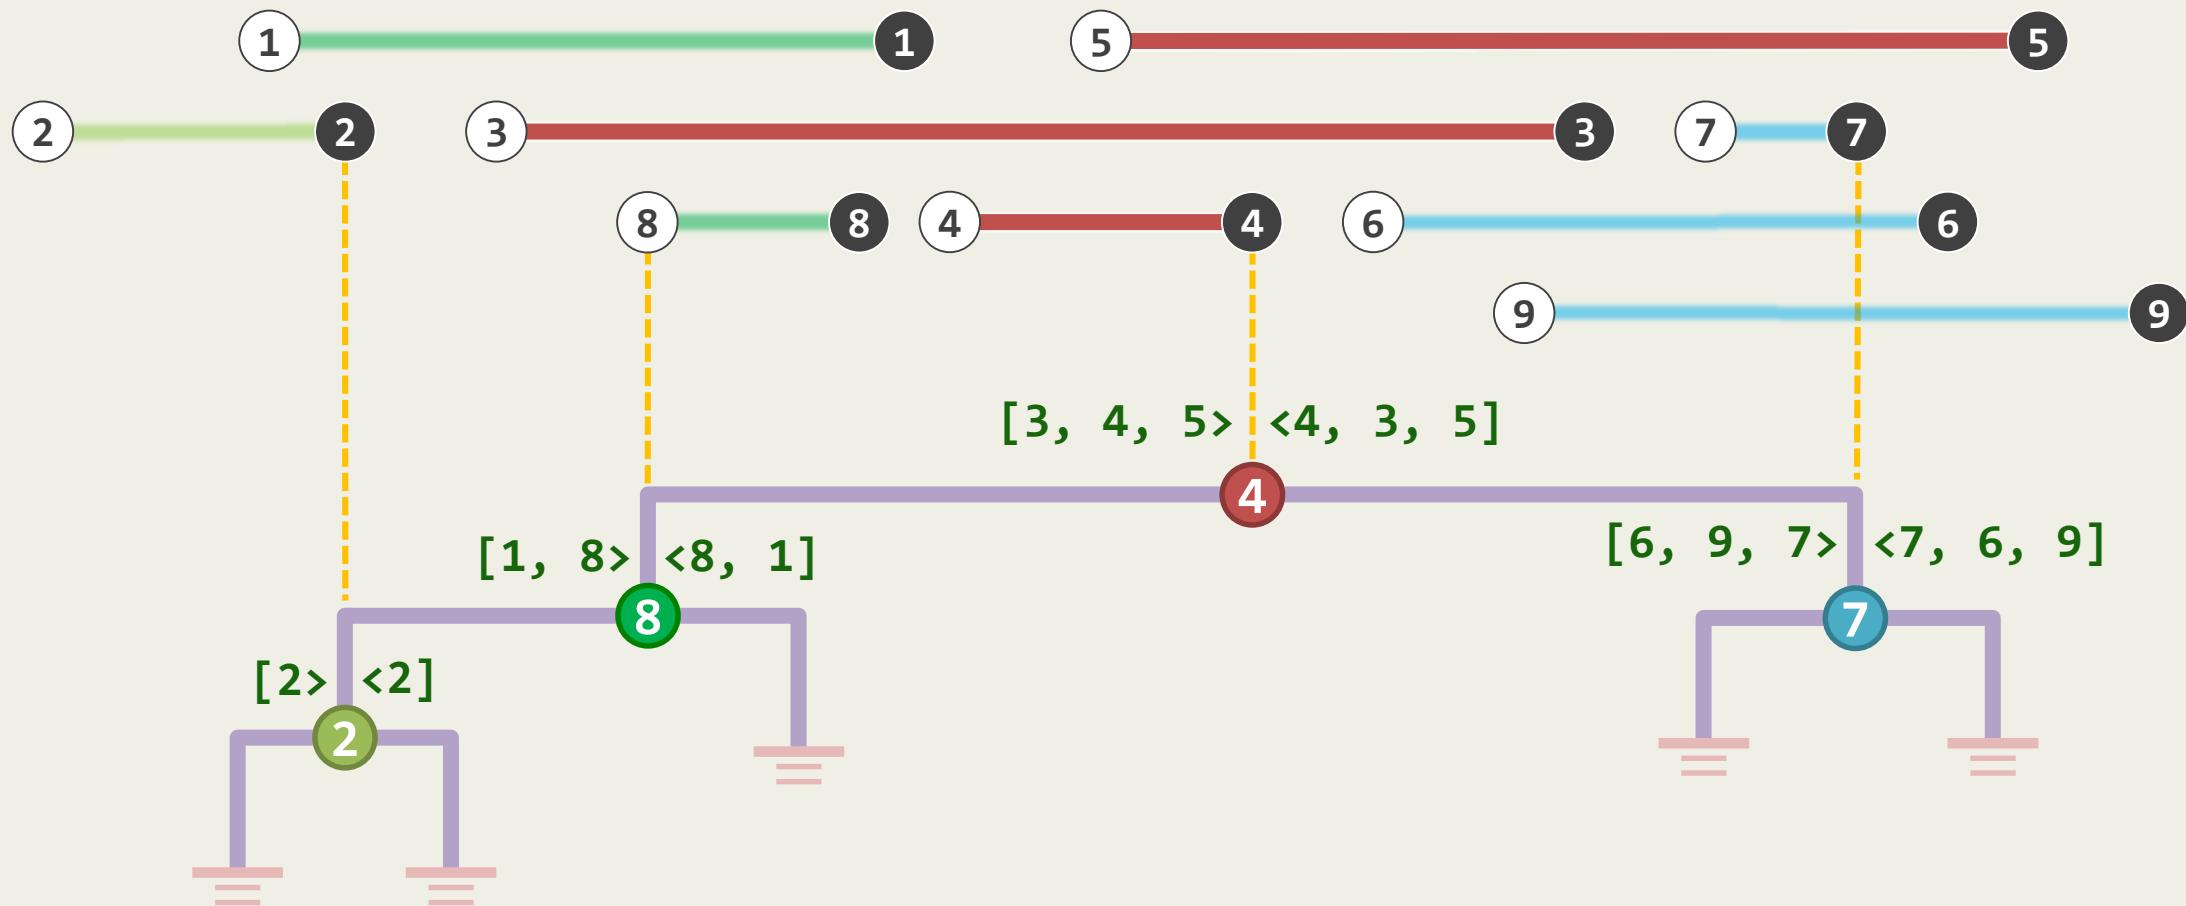
❖ $S_{left/right}$ will be recursively partitioned until they are empty (leaves)

Balance & $\mathcal{O}(\log n)$ Depth

$$\max\{ |S_{left}|, |S_{right}| \} \leq n/2$$

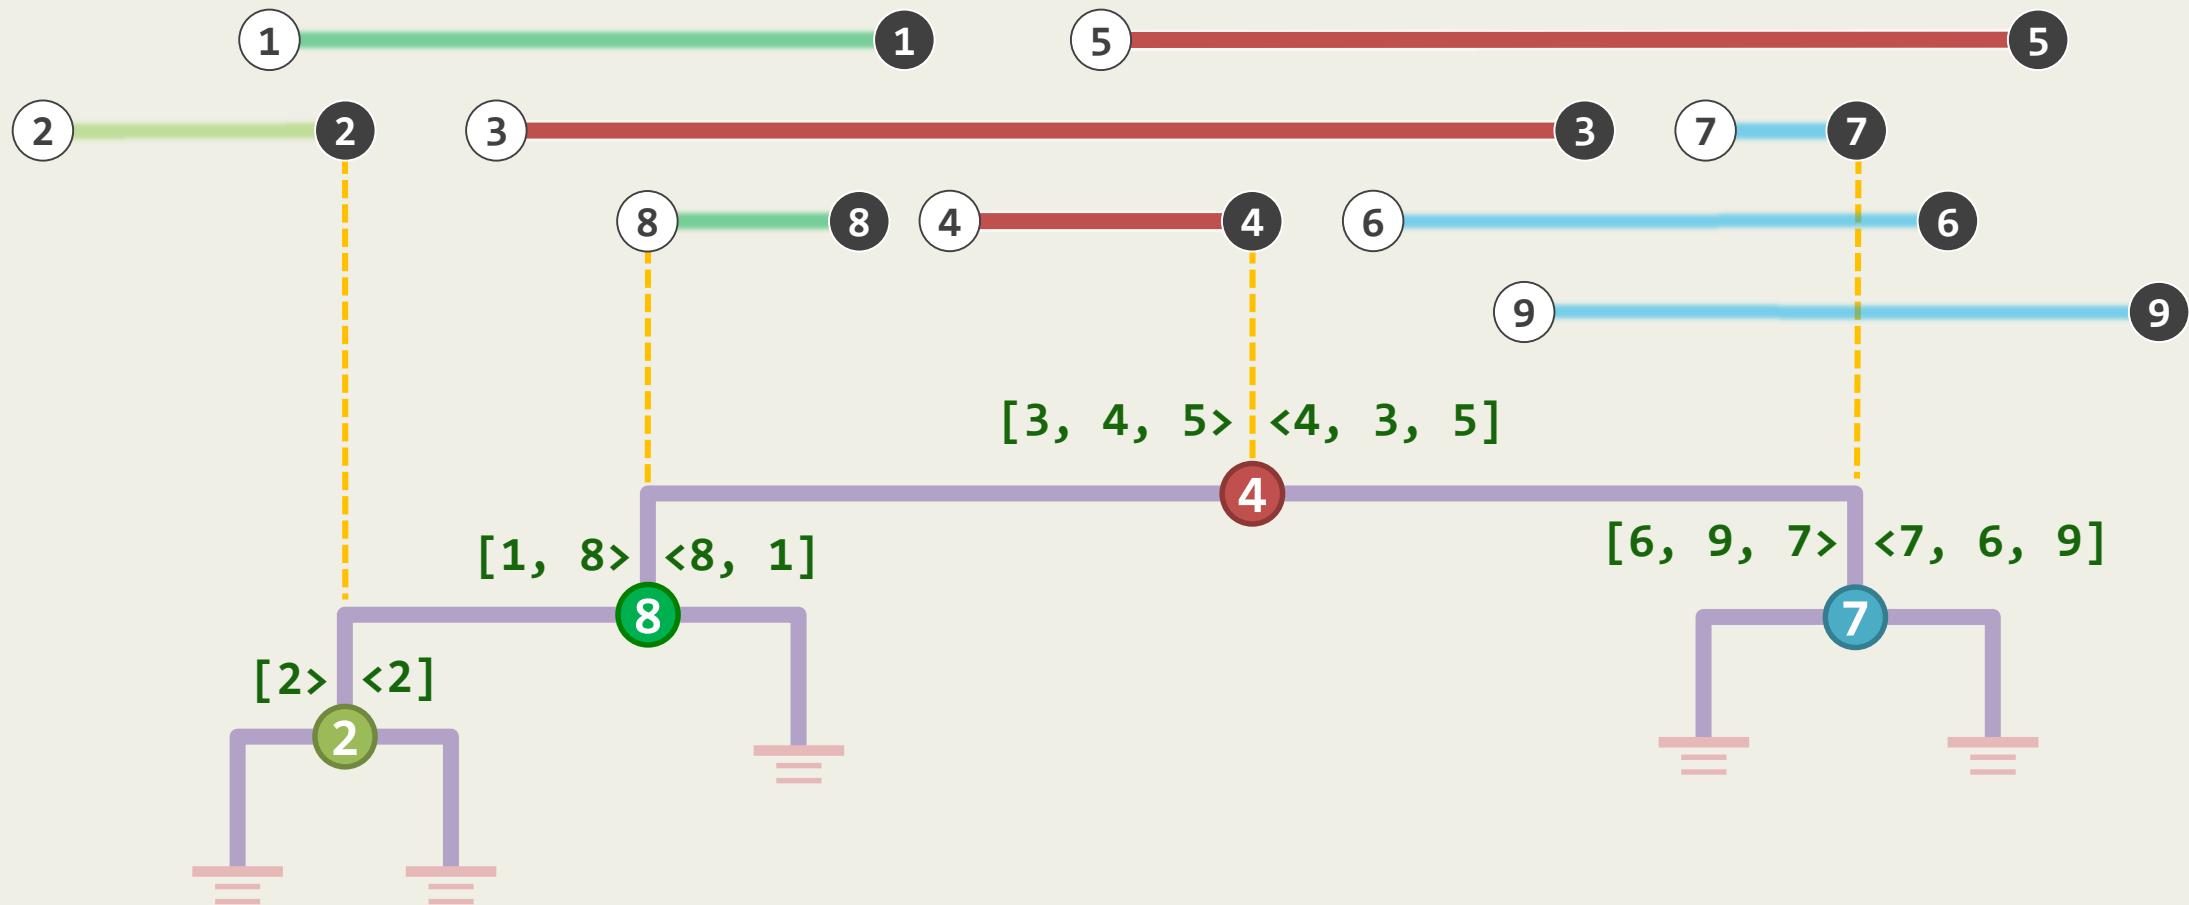
Best case: $|S_{mid}| = n$

Worst case: $|S_{mid}| = 1$



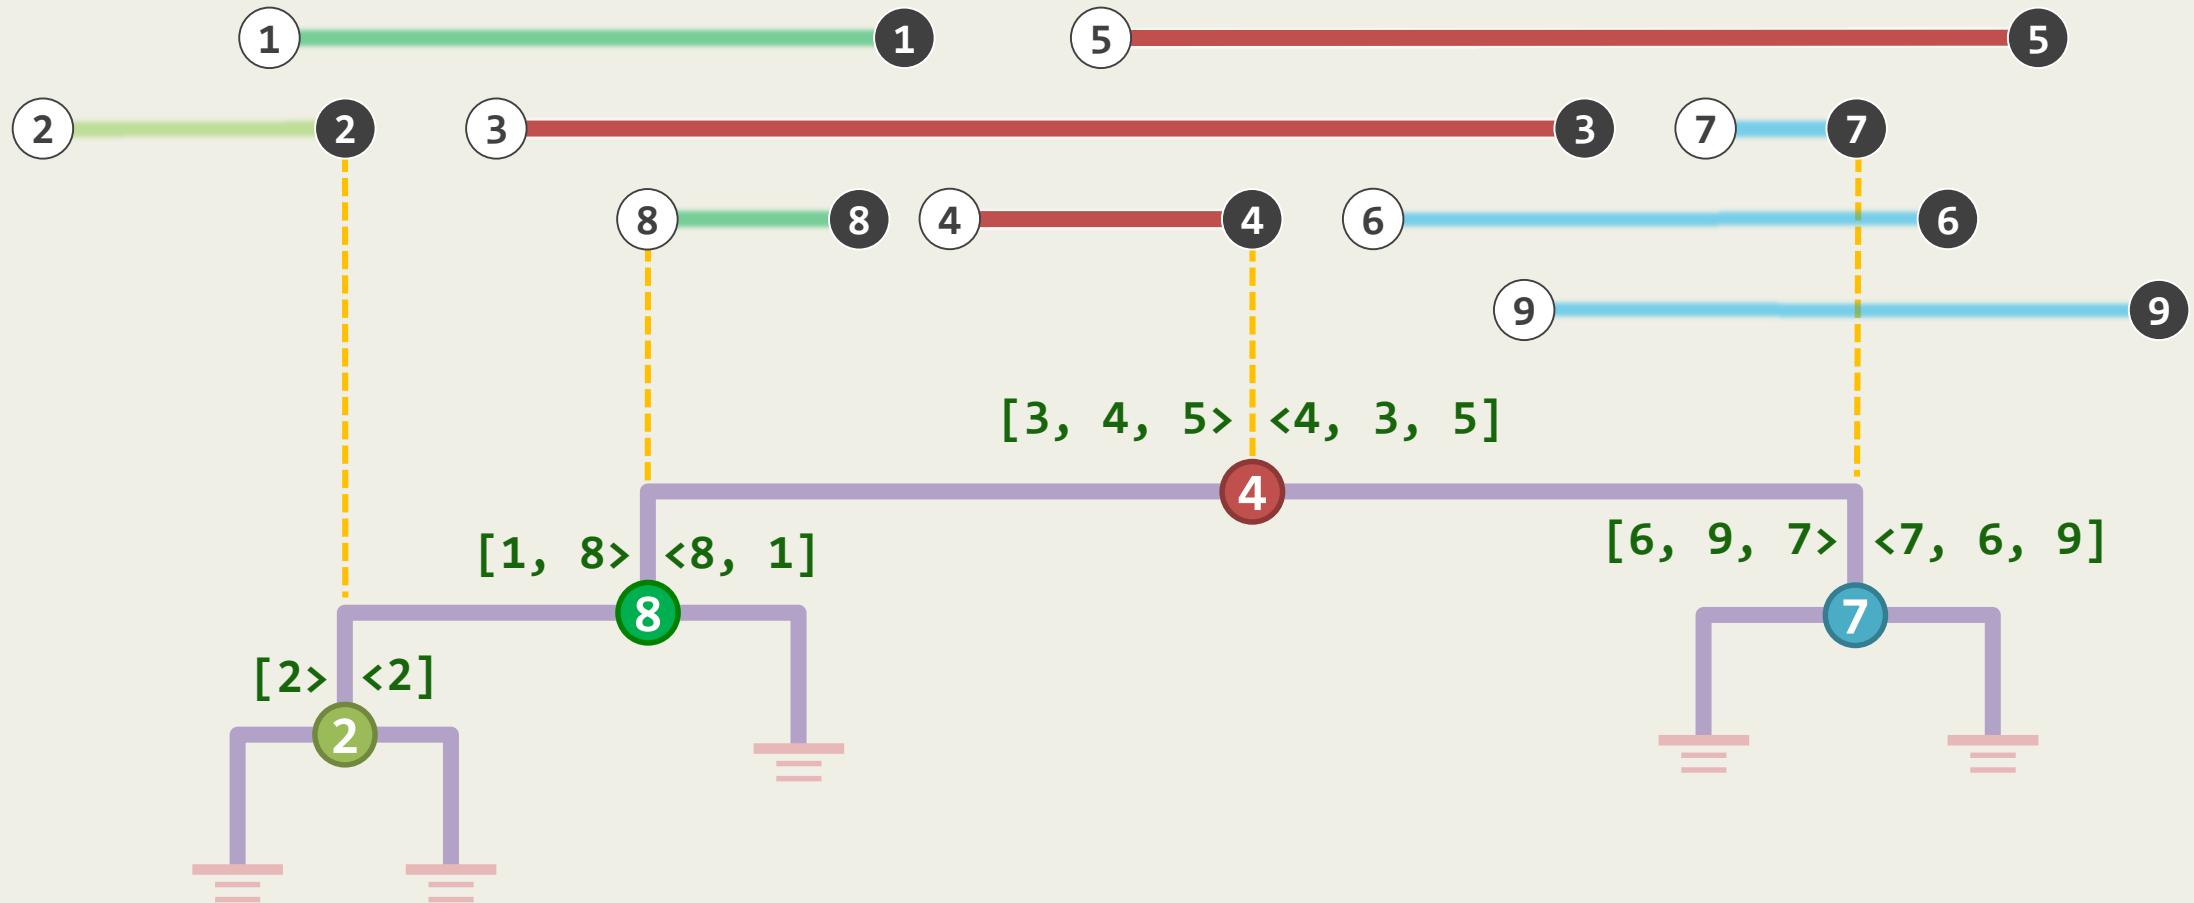
Associative Lists

❖ $L_{\text{left/right}} = \text{all intervals of } S_{\text{mid}}$ sorted by the **left/right** endpoints



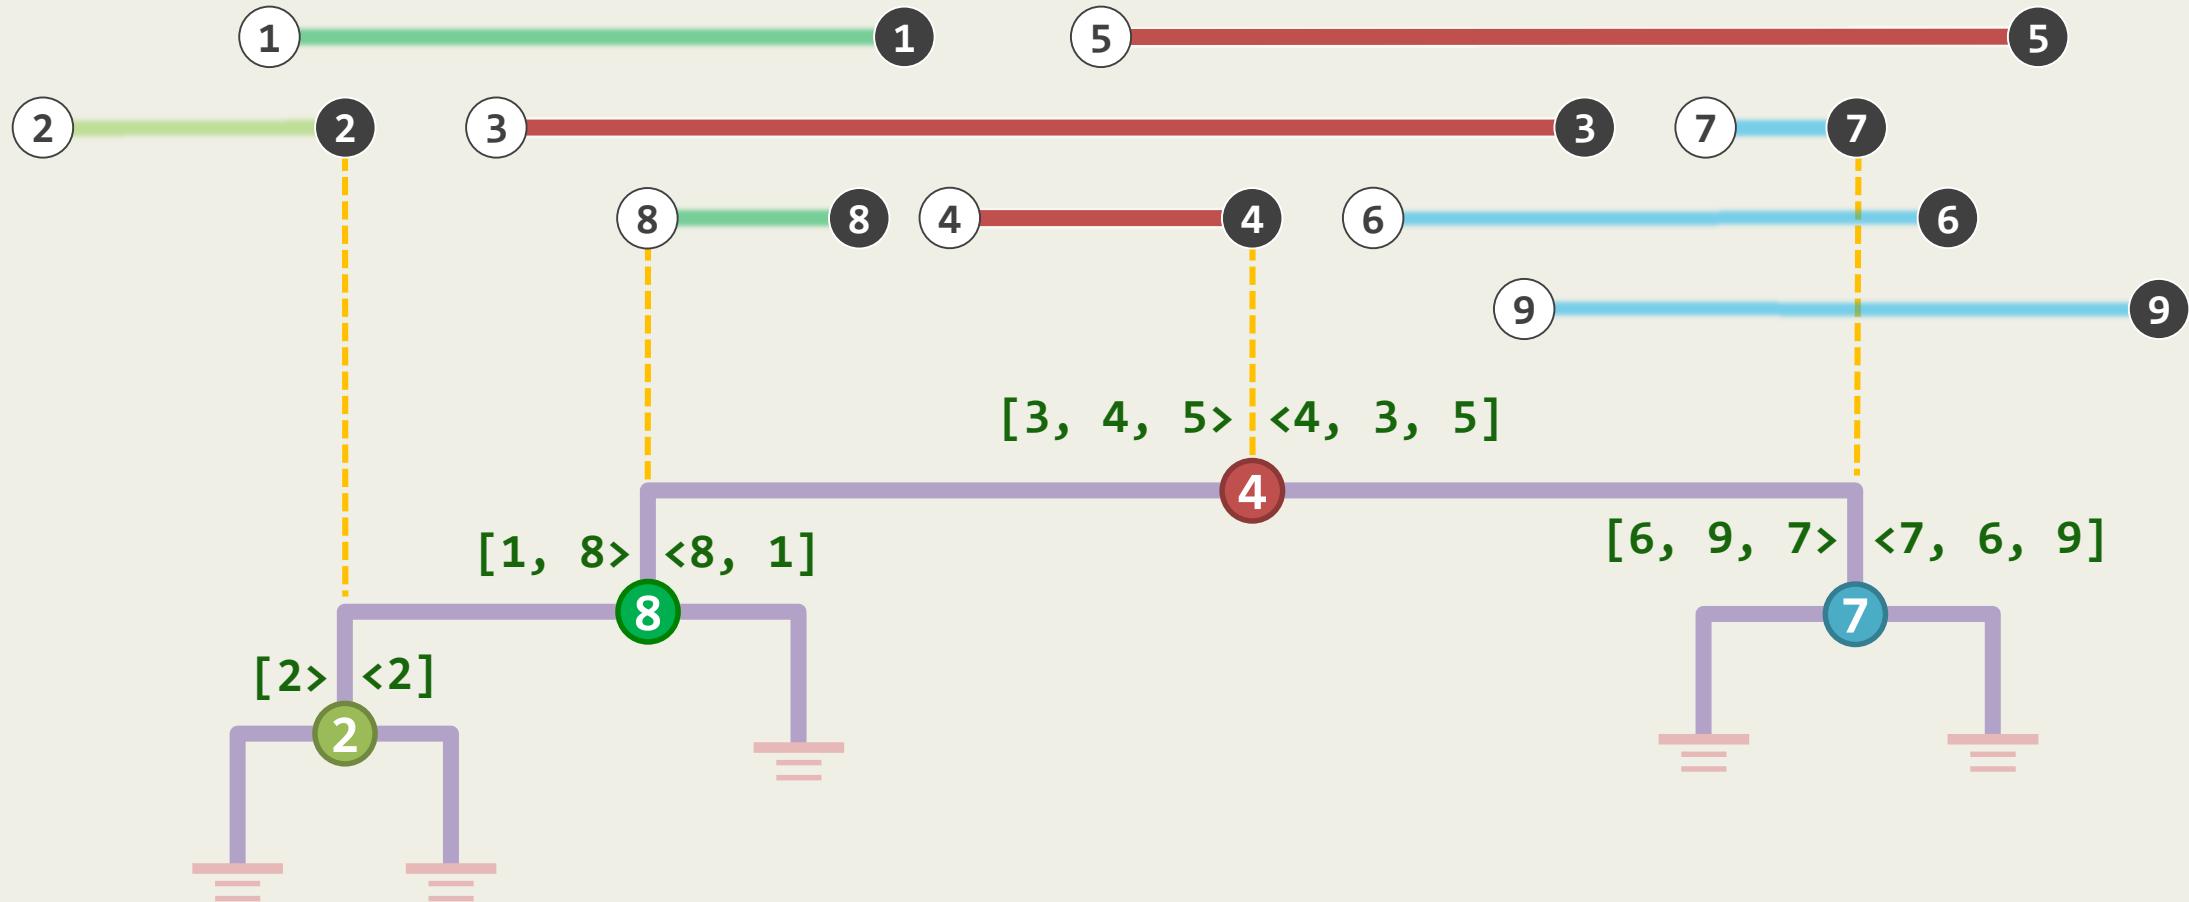
$\Theta(n)$ Size

❖ Each segment appears twice (one in each list)



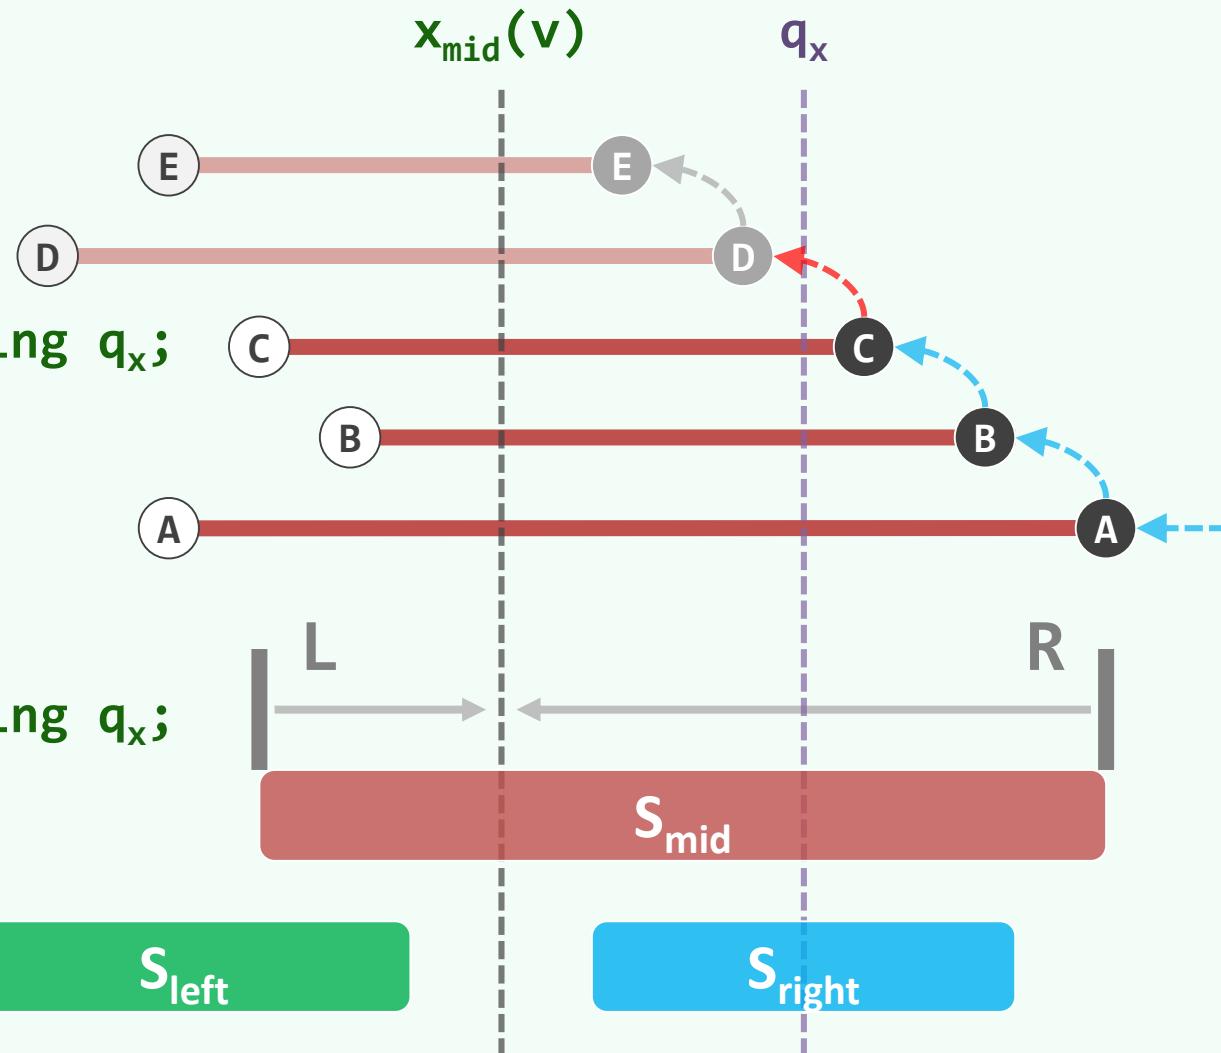
$\Theta(n \log n)$ Construction Time

❖ Hint: avoid repeatedly sorting



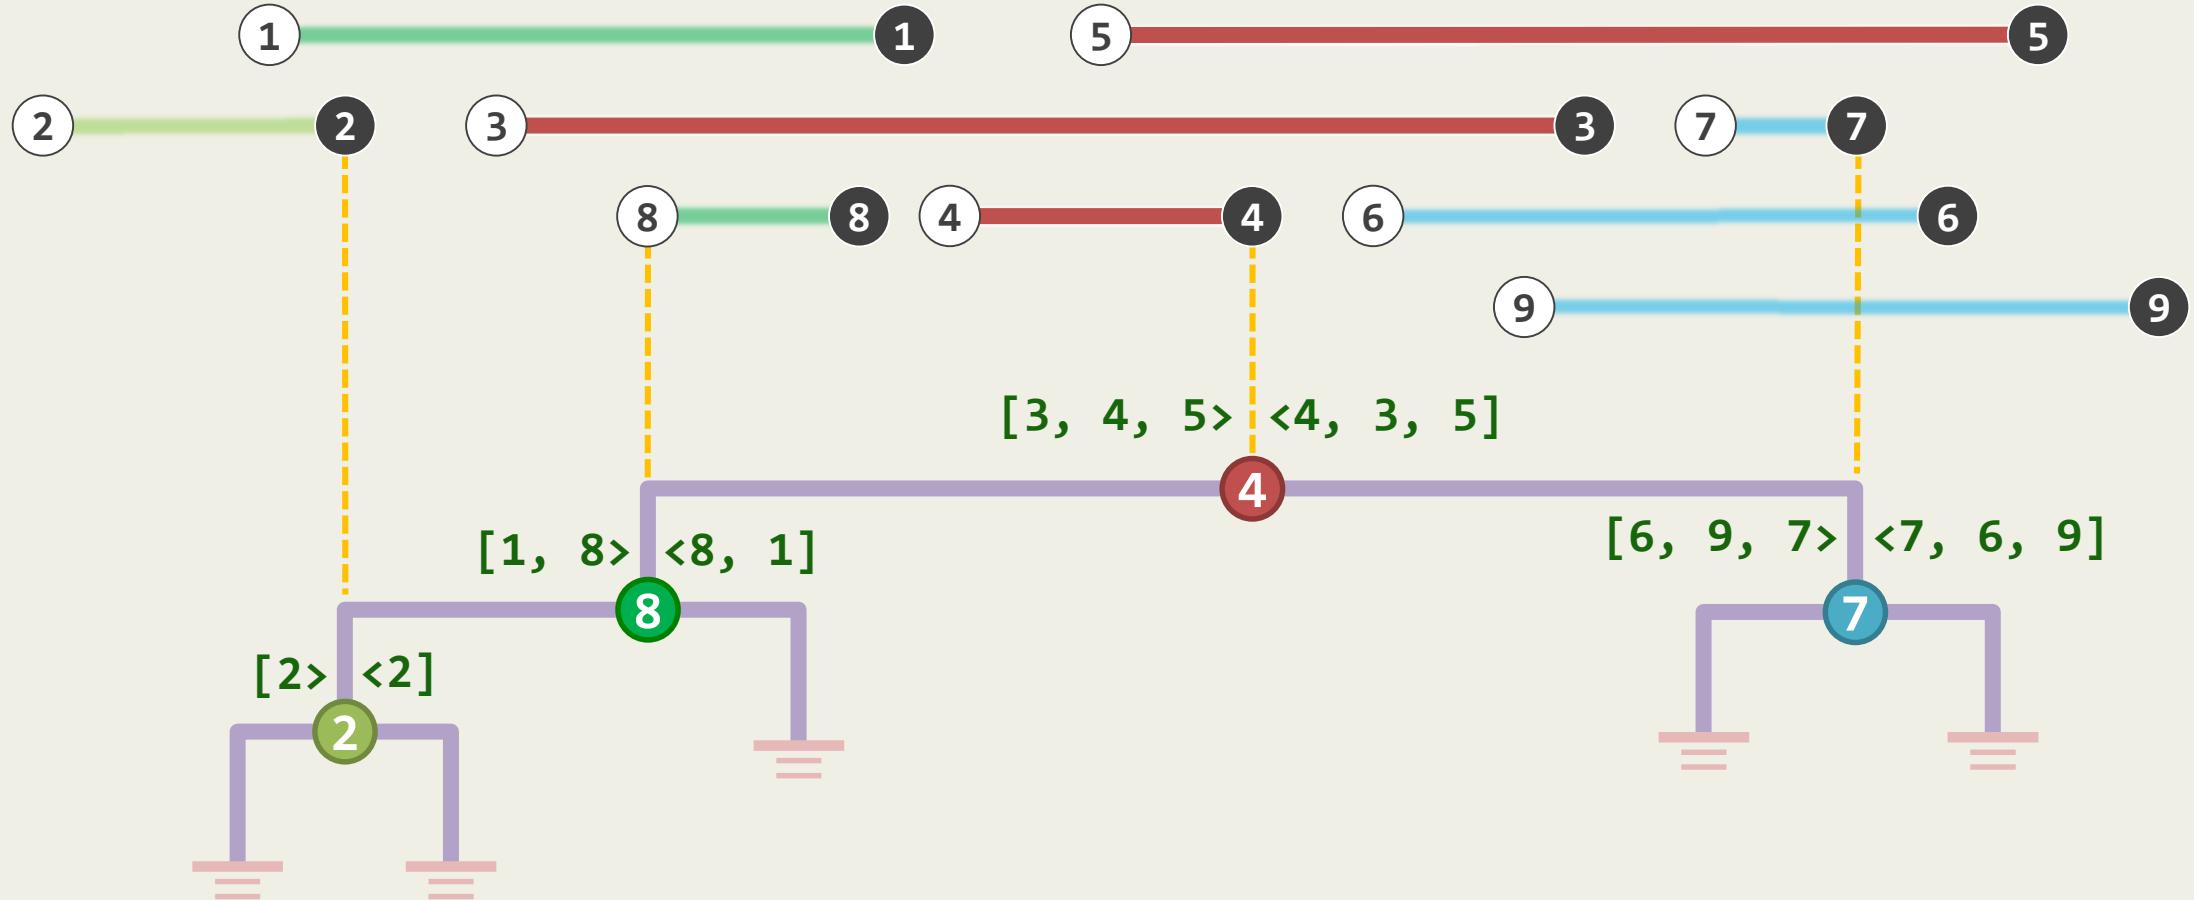
queryIntervalTree(v , q_x)

```
if ( ! v ) return; //base  
  
if (  $q_x < x_{\text{mid}}(v)$  )  
    report all segments of  $S_{\text{mid}}(v)$  containing  $q_x$ ;  
    queryIntervalTree( lc( $v$ ),  $q_x$  );  
  
else if (  $x_{\text{mid}}(v) < q_x$  )  
    report all segments of  $S_{\text{mid}}(v)$  containing  $q_x$ ;  
    queryIntervalTree( rc( $v$ ),  $q_x$  );  
  
else //with a probability  $\approx 0$   
    report all segments of  $S_{\text{mid}}(v)$ ; //both rc( $v$ ) & lc( $v$ ) can be ignored
```



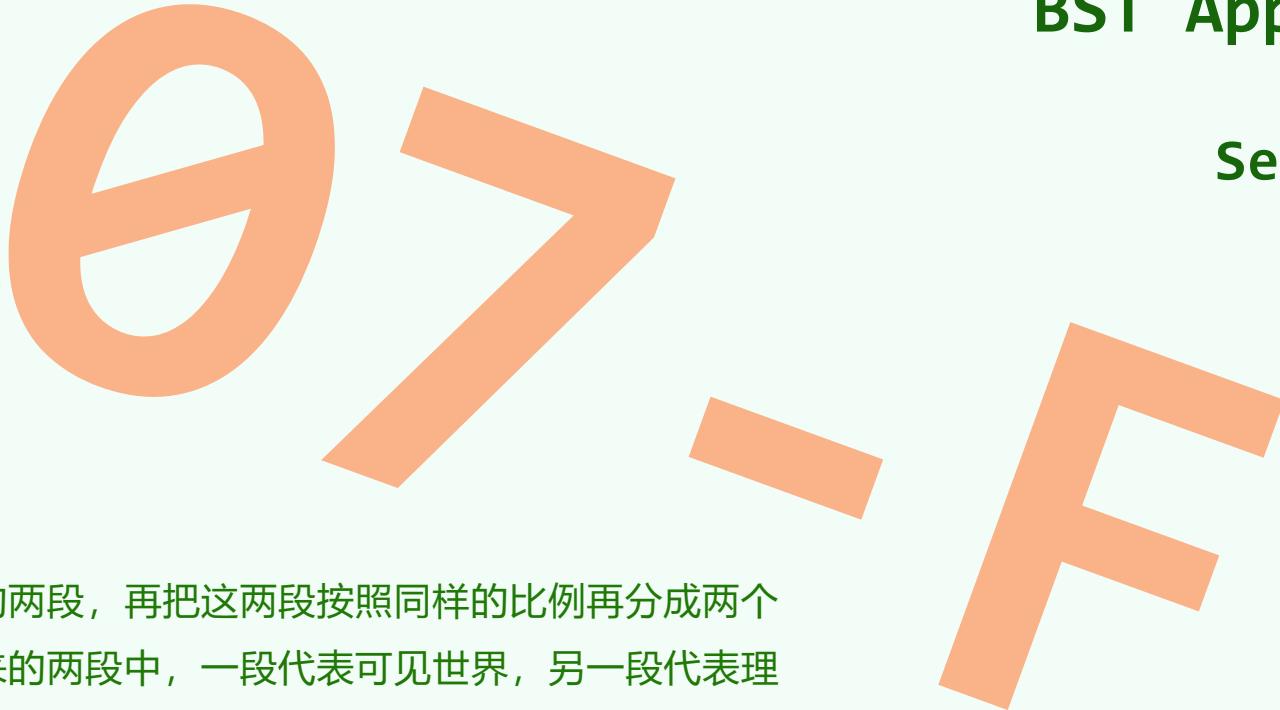
$\Theta(r + \log n)$ Query Time

❖ Each query visits $\Theta(\log n)$ nodes //LINEAR recursion



BST Application

Segment Tree



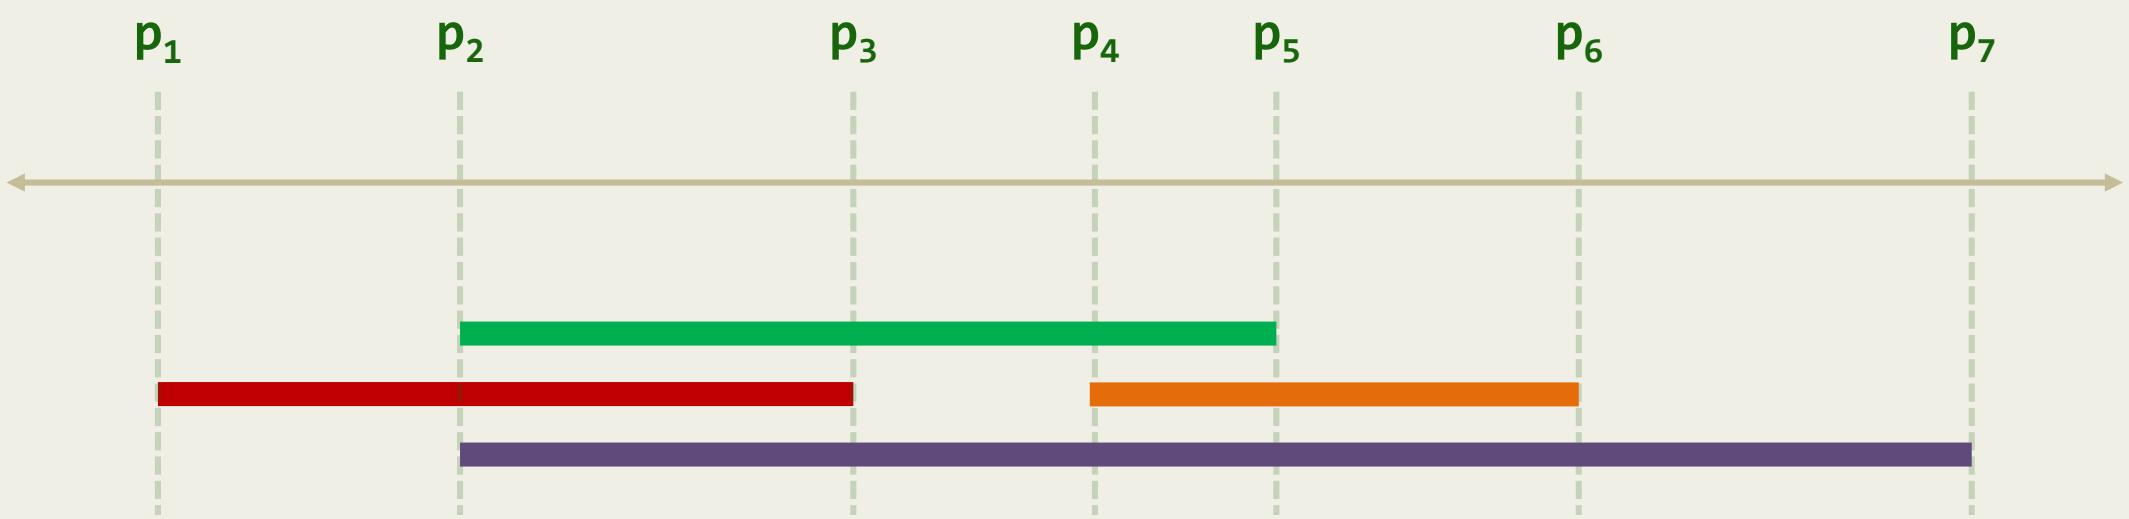
把一条线分割成不相等的两段，再把这两段按照同样的比例再分成两个部分。假设第一次分出来的两段中，一段代表可见世界，另一段代表理智世界。然后再看第二次分成的两段，他们分别代表清楚与不清楚的程度，你便会发现，可见世界那一段的第一部分是它的影像。

邓俊辉

deng@tsinghua.edu.cn

Elementary Intervals

- ❖ Let $I = \{ [x_i, x'_i] \mid i = 1, 2, 3, \dots, n \}$ be n intervals on the x-axis
- ❖ Sort all the endpoints into $\{ p_1, p_2, p_3, \dots, p_m \}$, $m \leq 2n$



- ❖ $m+1$ elementary intervals are hence defined as:

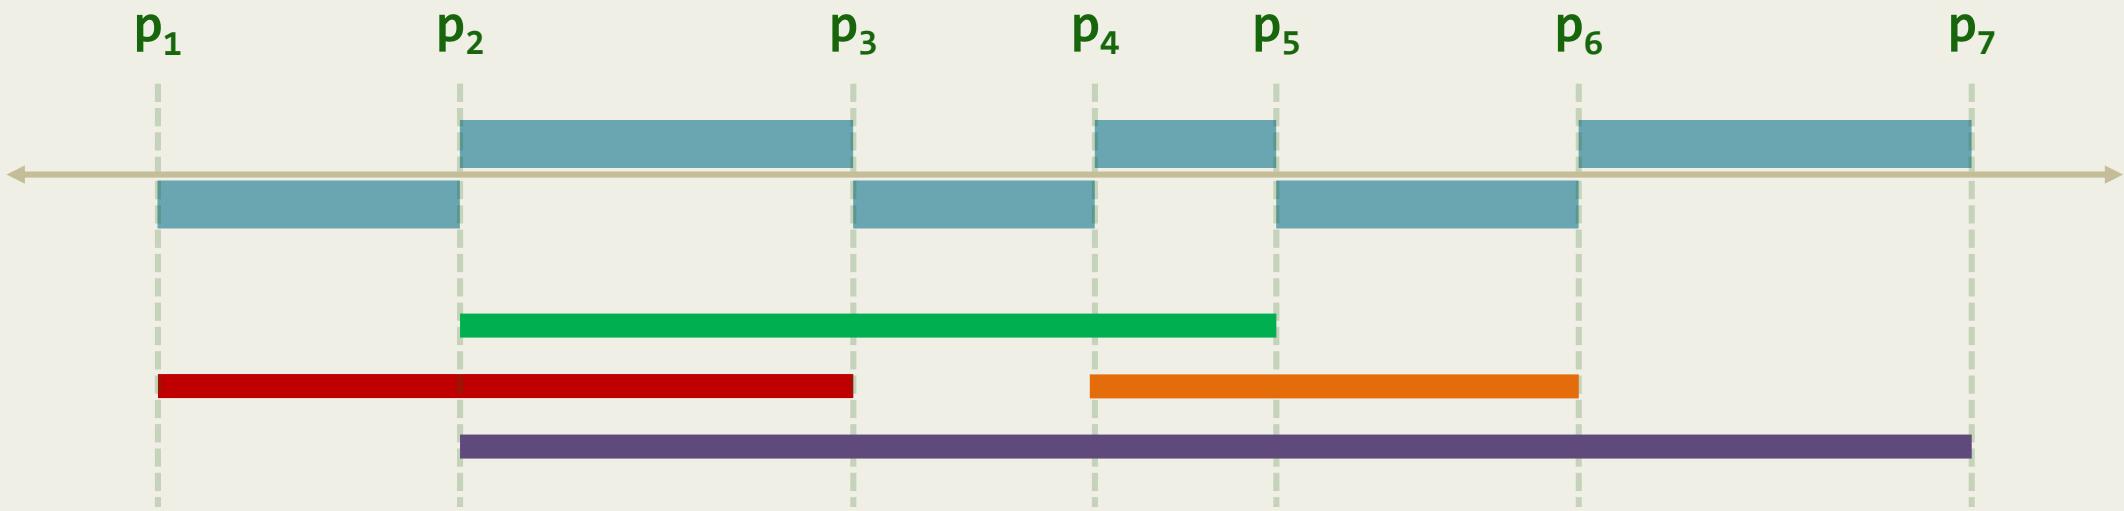
$$(-\infty, p_1], (p_1, p_2], (p_2, p_3], \dots, (p_{m-1}, p_m], (p_m, +\infty]$$

Discretization

⦿ Within each EI, all stabbing queries share a same output

∴ If we sort all EI's into a vector and

store the corresponding output with each EI, then ...



∴ Once a query position is determined, //by an $\mathcal{O}(\log n)$ time binary search
the output can then be returned directly // $\mathcal{O}(r)$

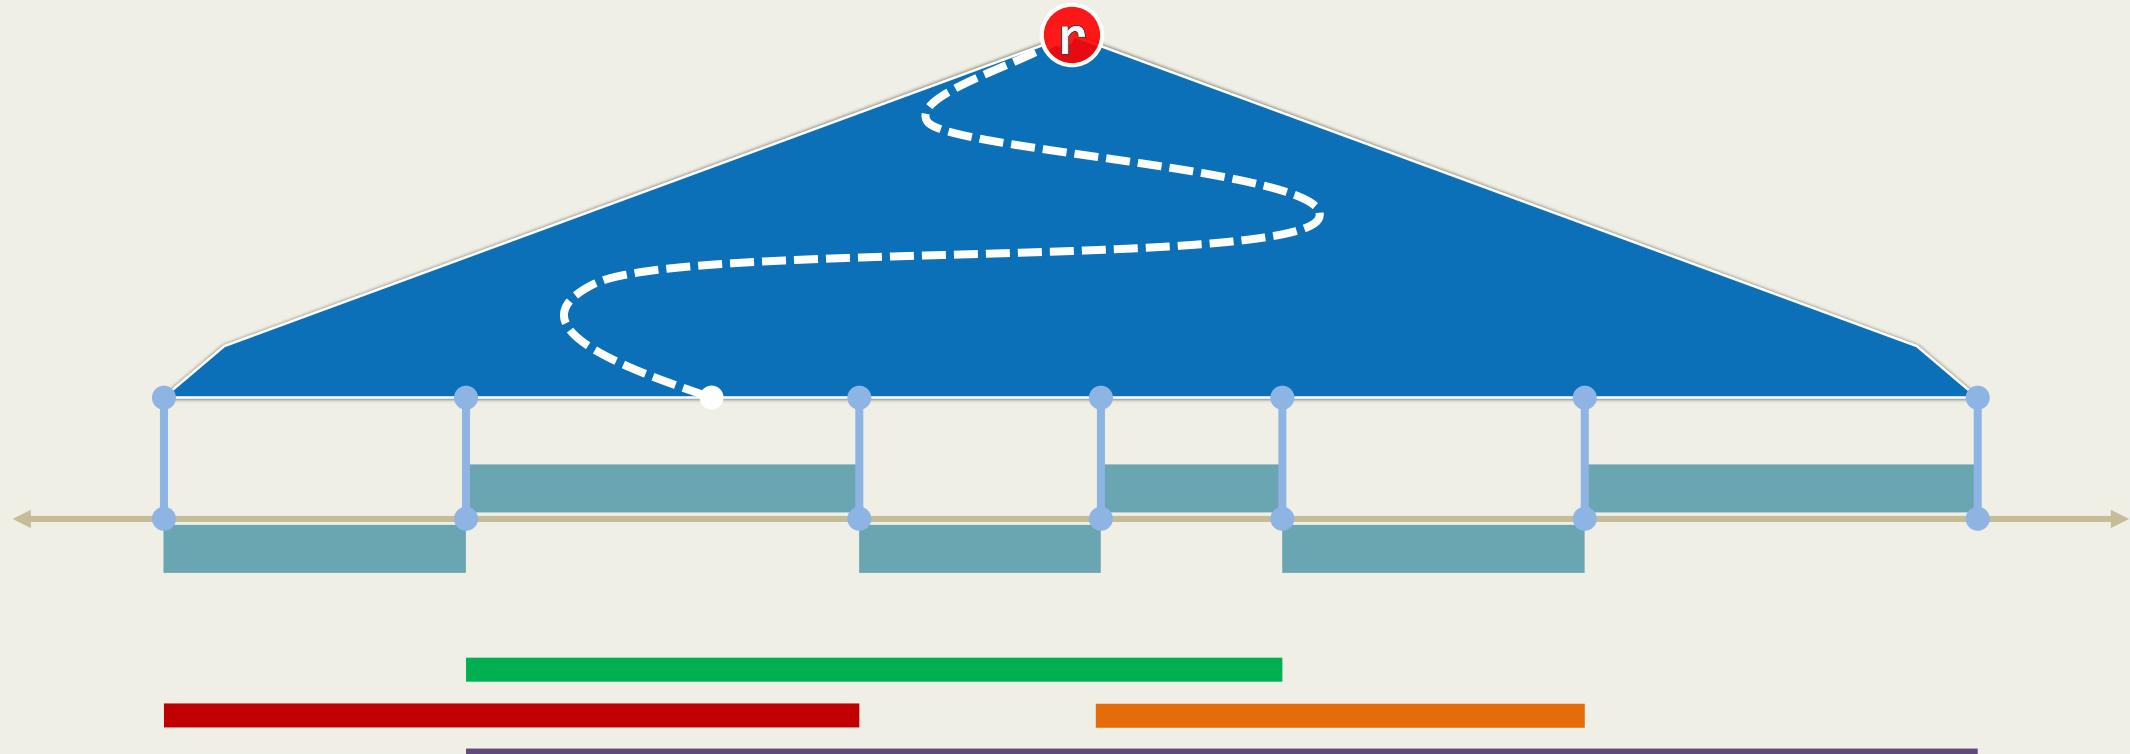
Worst Case

- ❖ Every interval spans $\Omega(n)$ EI's and a total space of $\Omega(n^2)$ is required



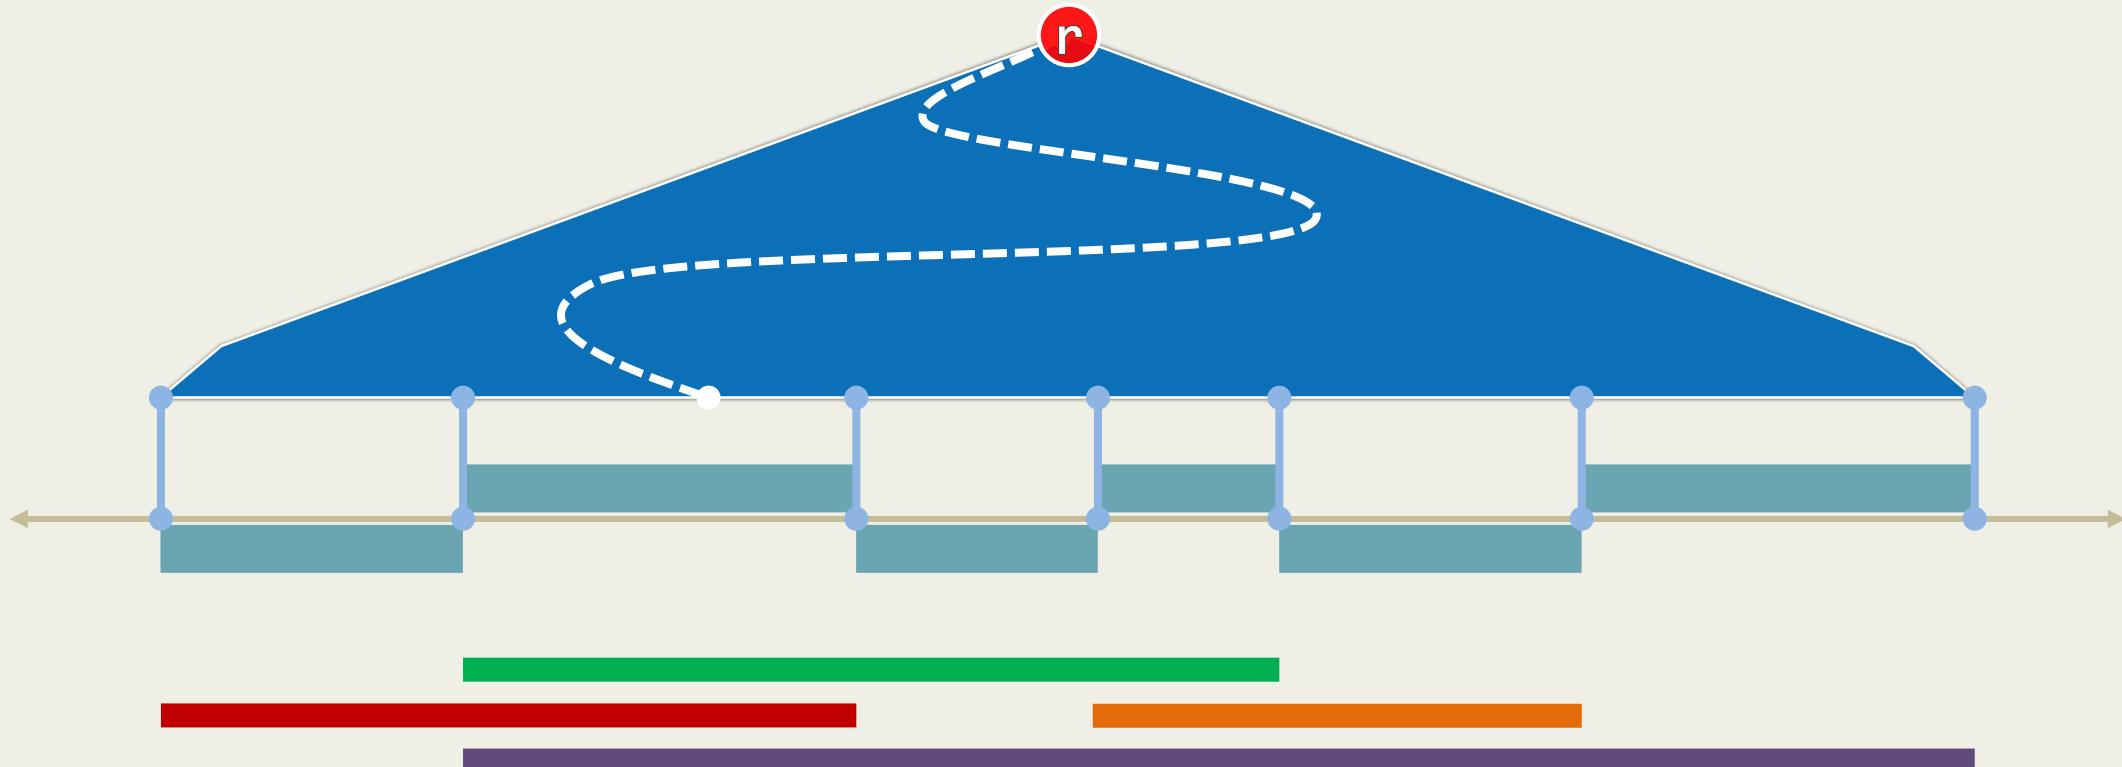
Sorted Vector \rightarrow BBST

- ❖ For each leaf v ,
denote the corresponding elementary interval as $R(v)$, //range of domination
denote the subset of intervals spanning $R(v)$ as $Int(v)$ and store $Int(v)$ at v

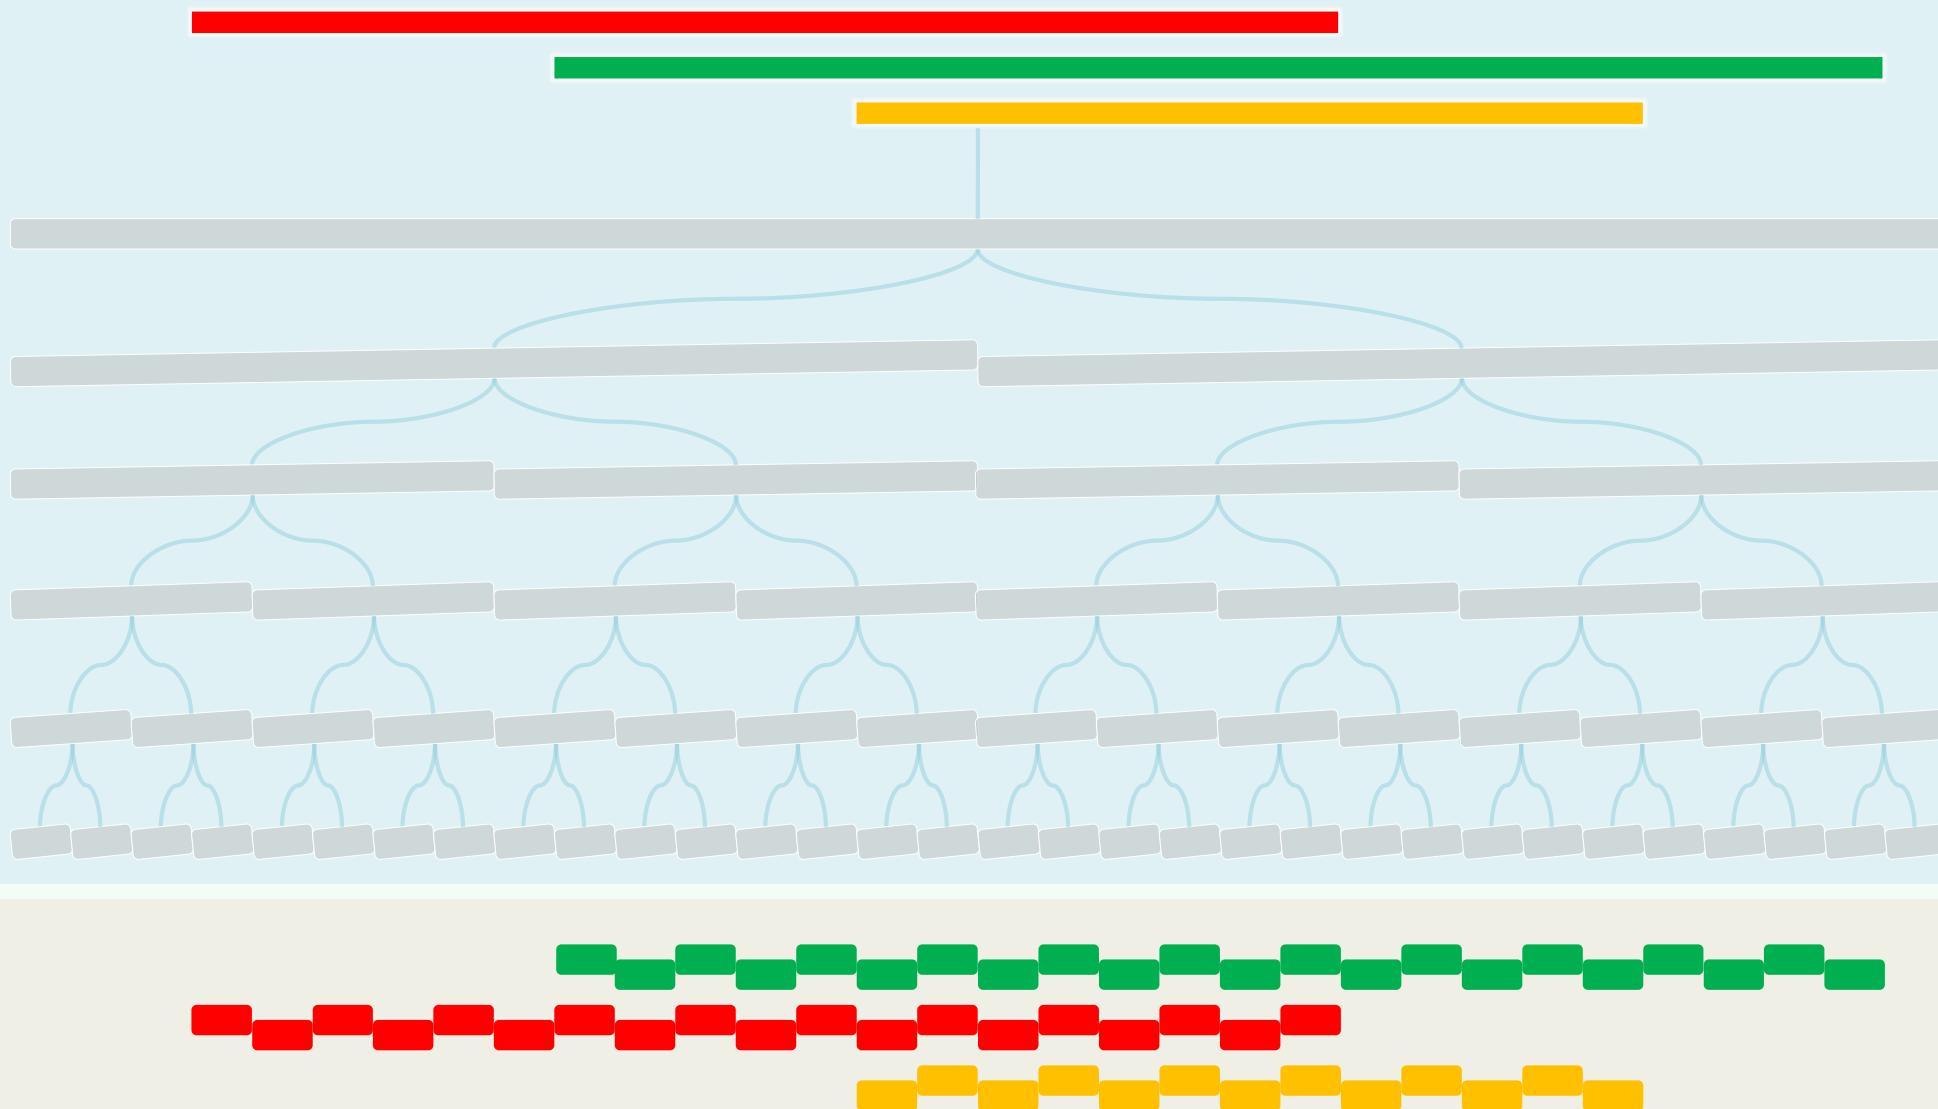


1D Stabbing Query with BBST

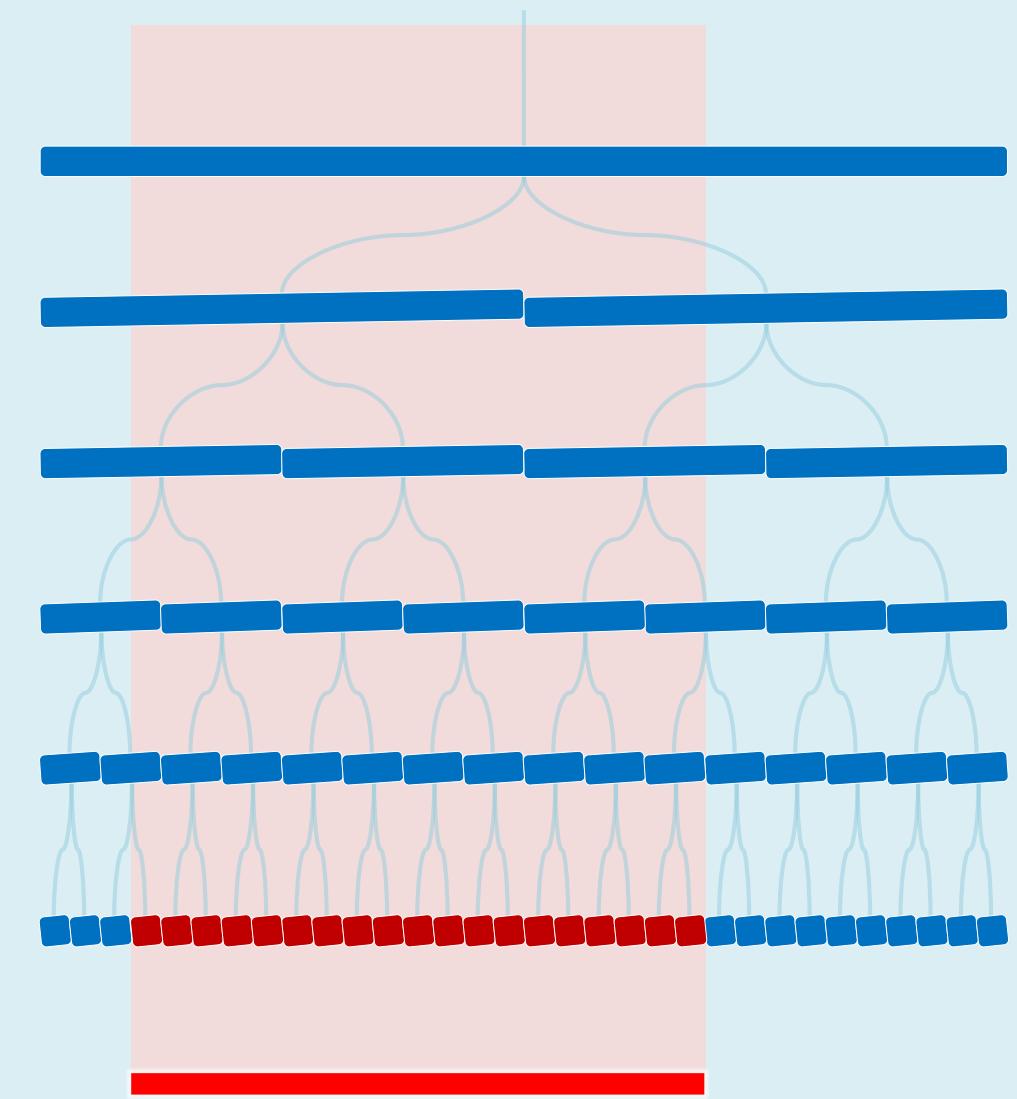
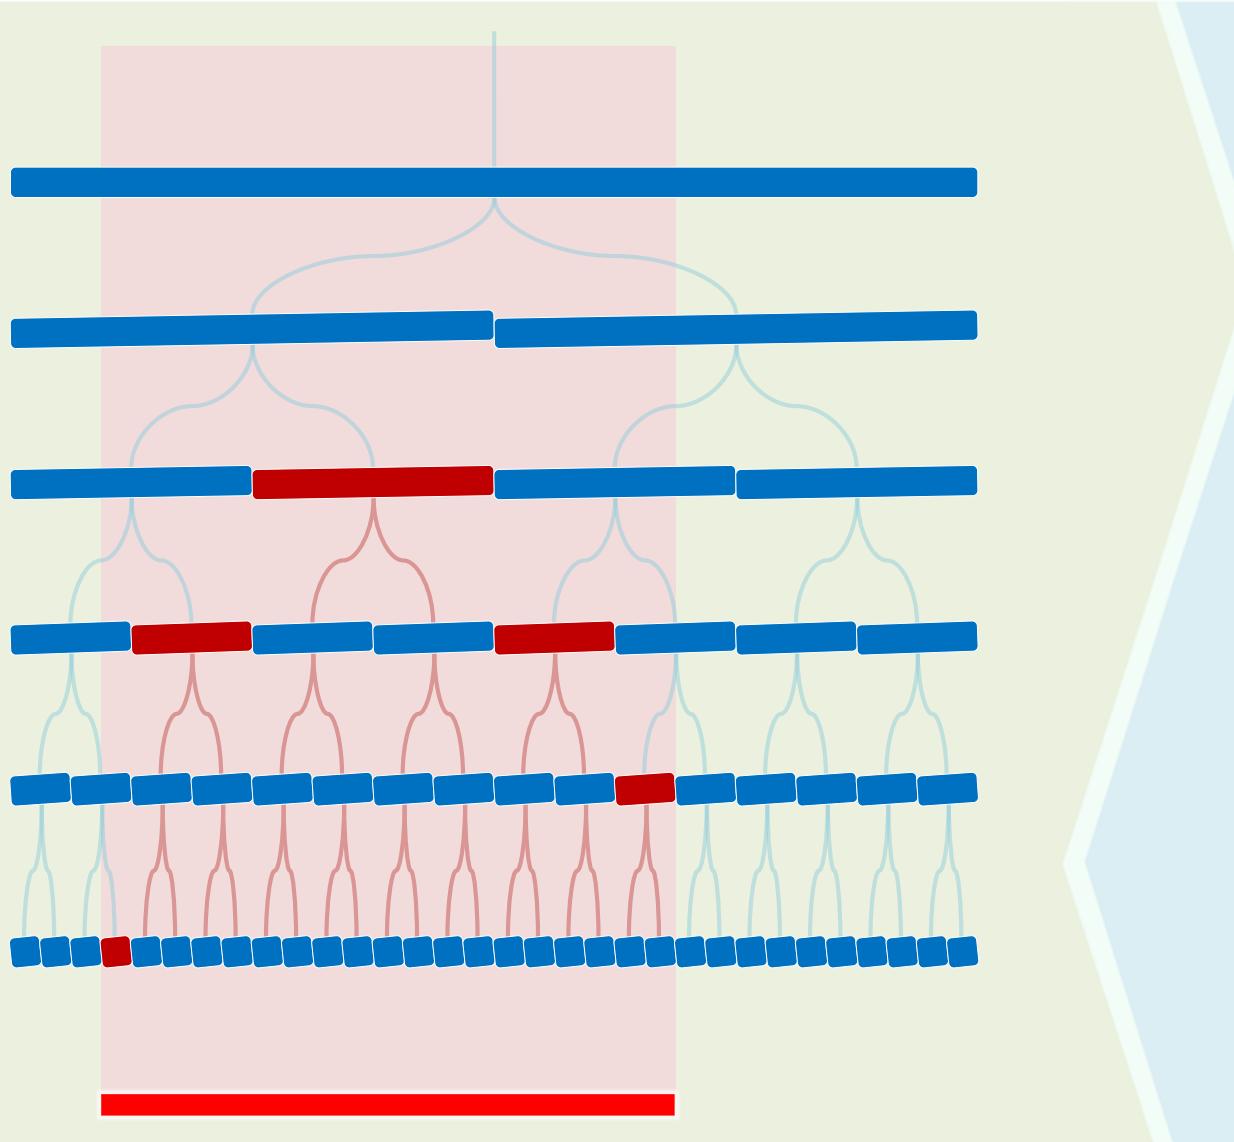
- ❖ To find all intervals containing q_x , we can
 - find the leaf v whose $R(v)$ contains q_x // $\mathcal{O}(\log n)$ time for a BBST
 - and then report $\text{Int}(v)$ // $\mathcal{O}(1 + r)$ time



$\Omega(n^2)$ Total Space In The Worst Cases



Store each interval at $\Theta(\log n)$ common ancestors by greedy merging



Canonical Subsets with $\Theta(n \log n)$ Space



BuildSegmentTree(I)

Sort all endpoints in I before

determining all the EI's // $\Theta(n \log n)$

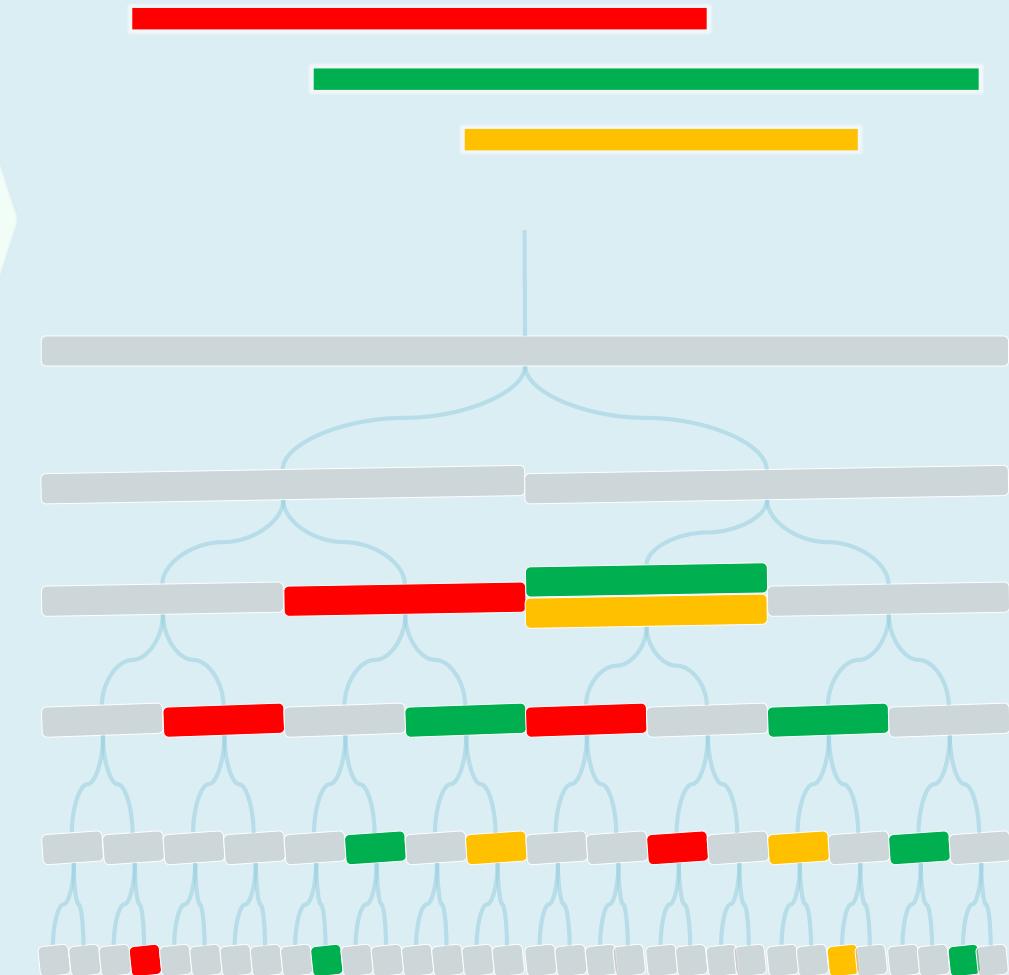
Create T a BBST on all the EI's // $\Theta(n)$

Determine $R(v)$ for each node v

// $\Theta(n)$ if done in a bottom-up manner

For each s of I

InsertSegment(T.root, s)

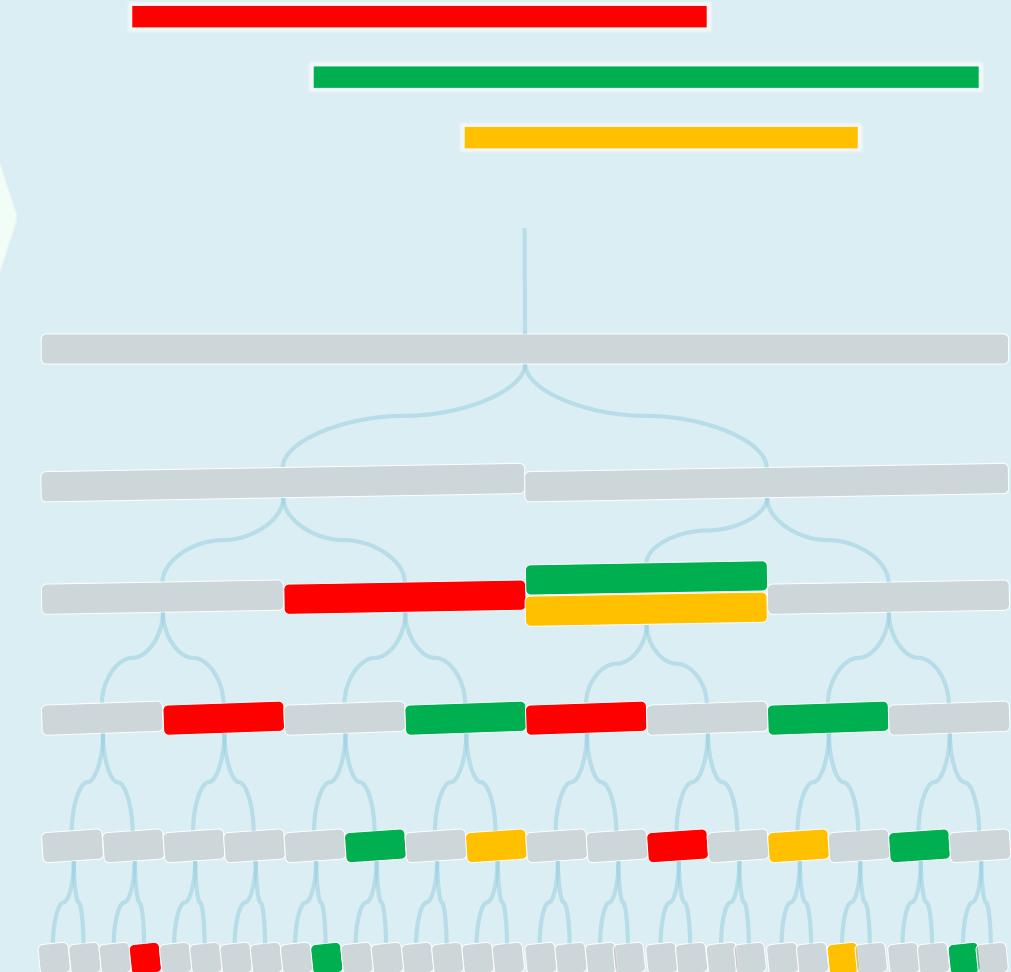


InsertSegment(v , s)

```
if (  $R(v) \subseteq s$  ) //greedy by top-down  
    store  $s$  at  $v$  and return;  
  
if (  $R( lc(v) ) \cap s \neq \emptyset$  ) //recurse  
    InsertSegment(  $lc(v)$ ,  $s$  );  
  
if (  $R( rc(v) ) \cap s \neq \emptyset$  ) //recurse  
    InsertSegment(  $rc(v)$ ,  $s$  );
```

- ⦿ At each level,
 ≤ 4 nodes are visited
(2 stores + 2 recursions)

∴ $O(\log n)$ time



Query(v , q_x)

report all the intervals in $\text{Int}(v)$

if (v is a leaf)

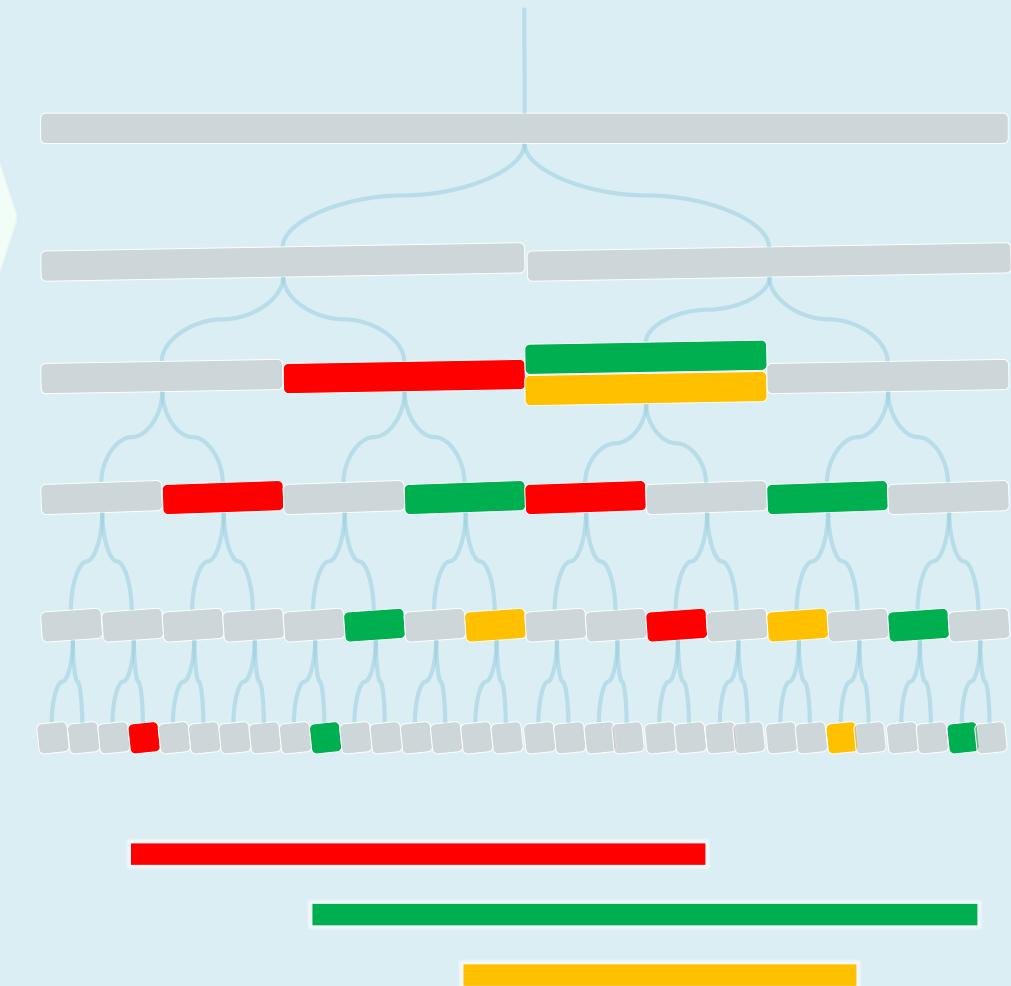
return

if ($q_x \in R(\text{lc}(v))$)

Query($\text{lc}(v)$, q_x)

else // $q_x \in R(\text{rc}(v))$

Query($\text{rc}(v)$, q_x)



$\mathcal{O}(r + \log n)$

⦿ Only one node is visited per level,

altogether $\mathcal{O}(\log n)$ nodes

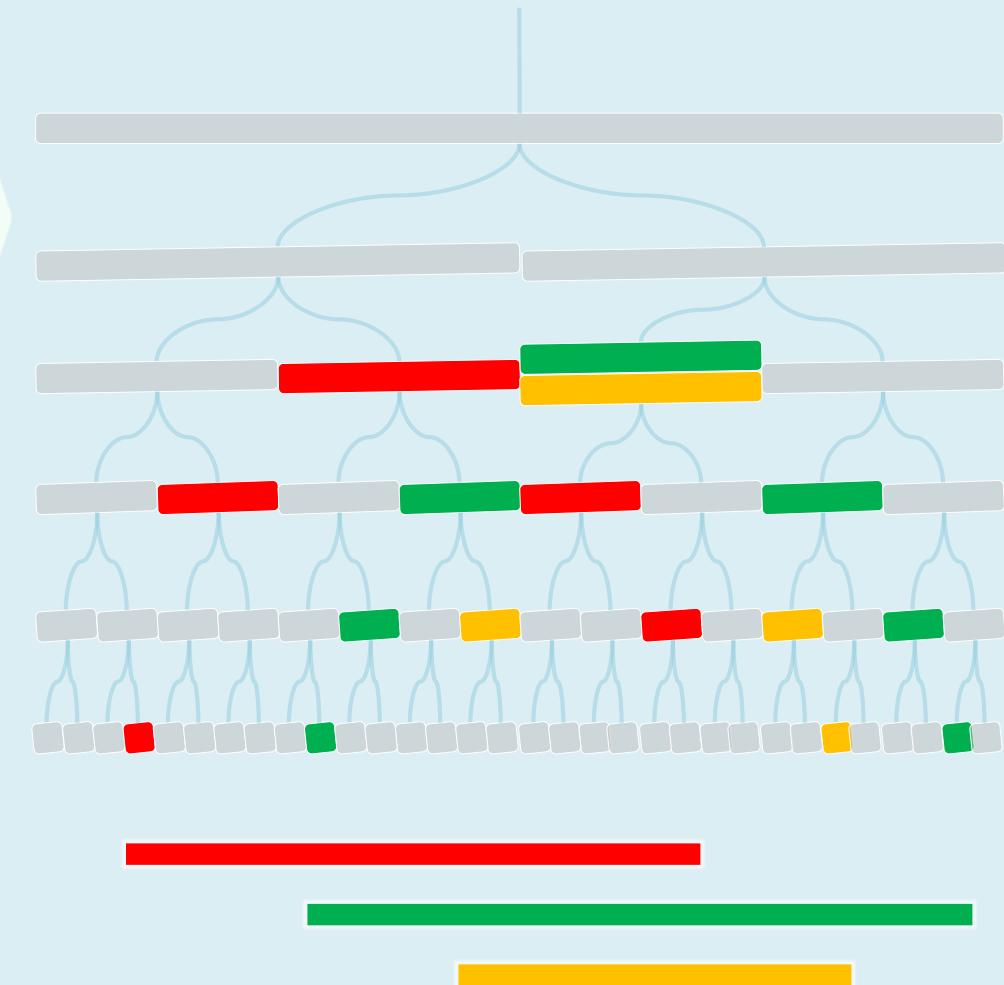
⦿ At each node v

- the CS $\text{Int}(v)$ is reported
- in time

$$1 + |\text{Int}(v)| = \mathcal{O}(1 + r_v)$$

∴ Reporting all the intervals

costs $\mathcal{O}(r)$ time



Conclusion

- ❖ For a set of n intervals,
 - a segment tree of size $\Theta(n \log n)$
 - can be built in $\Theta(n \log n)$ time
 - which reports all intervals containing a query point in $\Theta(r + \log n)$ time

